

3D Vision 1

Week 7

Image Formation
Camera Projection Matrix

Announcements

- Assignment 2 due in **one week** (11:59pm Friday 26 April)
 - This includes a one week extension that has already been applied
 - **Zero** marks if either report or code submitted late (unless extension)
 - Submit early; you can always resubmit an updated version later
 - Depending on your internet connection and load on the TurnItIn servers, uploading can sometimes be slow, so please factor this into your submission schedule
 - Submit your report (PDF) and code (ZIP file) **separately under the correct tab** in the submission box
 - Follow the instructions under Submission Requirements

Announcements

- Public holiday Thursday 25 April: Thursday lab rescheduled at **13:00-15:00 Tuesday Rm 109 CSIT Building**

Weekly Study Plan: Overview

Wk	Starting	Lecture	Lab	Assessment
1	19 Feb	Introduction	X	
2	26 Feb	Low-level Vision 1	1	
3	4 Mar	Low-level Vision 2	1	
		Mid-level Vision 1		
4	11 Mar	Mid-level Vision 2	1	CLab1 report due Friday
		High-level Vision 1		
5	18 Mar	High-level Vision 2	2	
6	25 Mar	High-level Vision 3 ¹	2	
	1 Apr	Teaching break	X	
	8 Apr	Teaching break	X	
7	15 Apr	3D Vision 1	2	CLab2 report due Friday
8	22 Apr	3D Vision 2	3	
9	29 Apr	3D Vision 3	3	
10	6 May	3D Vision 4	3	
		Mid-level Vision 3		
11	13 May	High-level Vision 4	X	CLab3 report due Friday
12	20 May	Course Review	X	



Weekly Study Plan: Part B

Wk	Starting	Lecture	By
7	15 Apr	3D vision: introduction, camera model, single-view geometry	Dylan
8	22 Apr	3D vision: camera calibration, two-view geometry (homography)	Dylan
9	29 Apr	3D vision: two-view geometry (epipolar geometry, triangulation, stereo)	Dylan
10	6 May	3D vision: multiple-view geometry	Weijian
		Mid-level vision: optical flow, shape-from-X	Dylan
11	13 May	High-level vision: self-supervised learning, detection, segmentation	Dylan
12	20 May	Course review	Dylan

Outline

1. Introduction to 3D vision
2. Model fitting (line fitting)
 1. Least squares
 2. M-estimation
 3. RANSAC
 4. Hough transform
3. Image formation (review): pinhole camera model
4. Camera projection matrix and single view geometry
5. Camera calibration
6. Resectioning and camera pose

Hough Transform

Fitting Multiple Lines Using a Hough Transform

- Given a binary edge image, find the lines (or curves) that explain the data points best in the parameter space
- This parameter space is called a **Hough space**

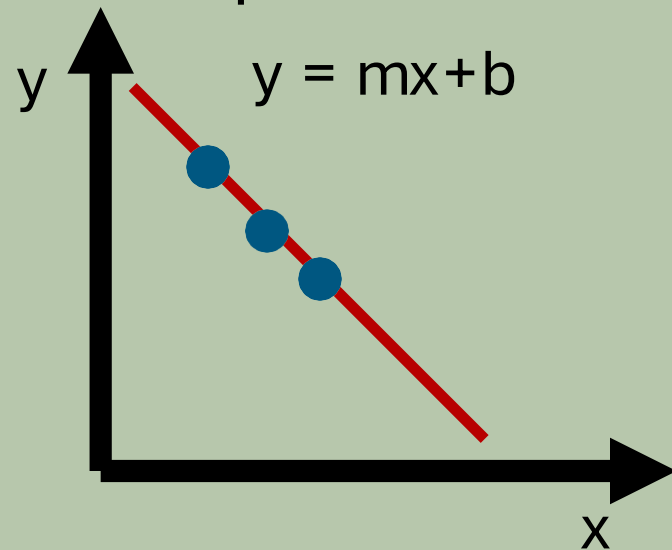
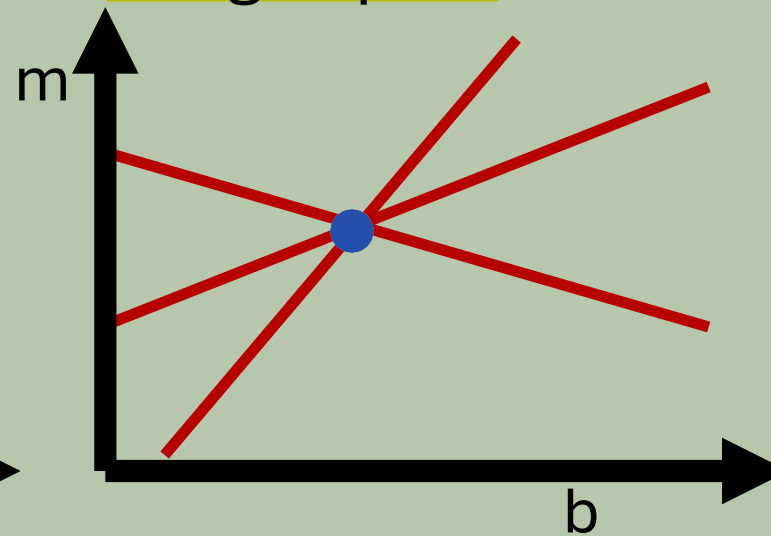


Image space

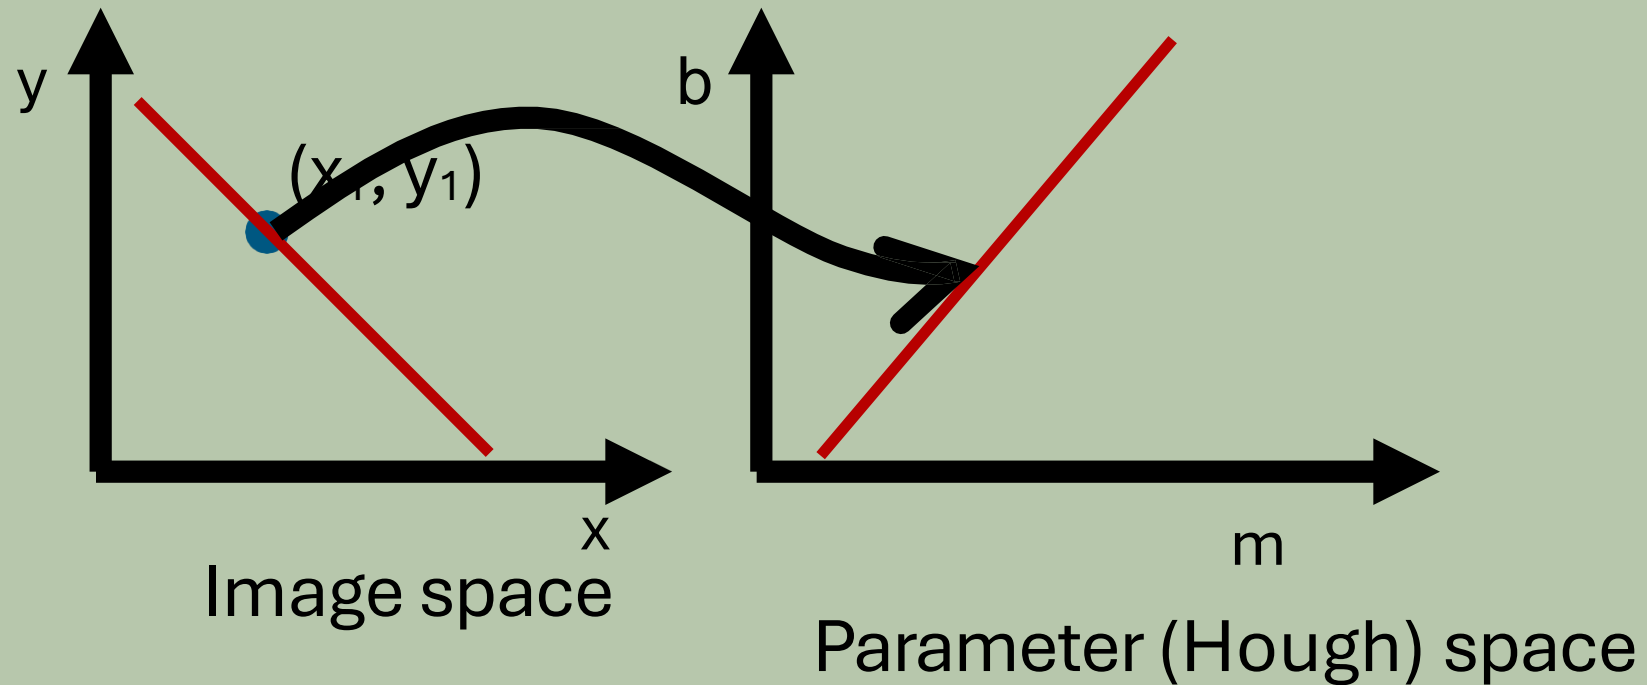


Parameter (Hough) space

A point (x_1, y_1) is mapped to a line in Hough space

$$y_1 = m x_1 + b$$

$$b = -x_1 m + y_1 \quad (\text{rewrite the equation in terms of } m \text{ and } b)$$

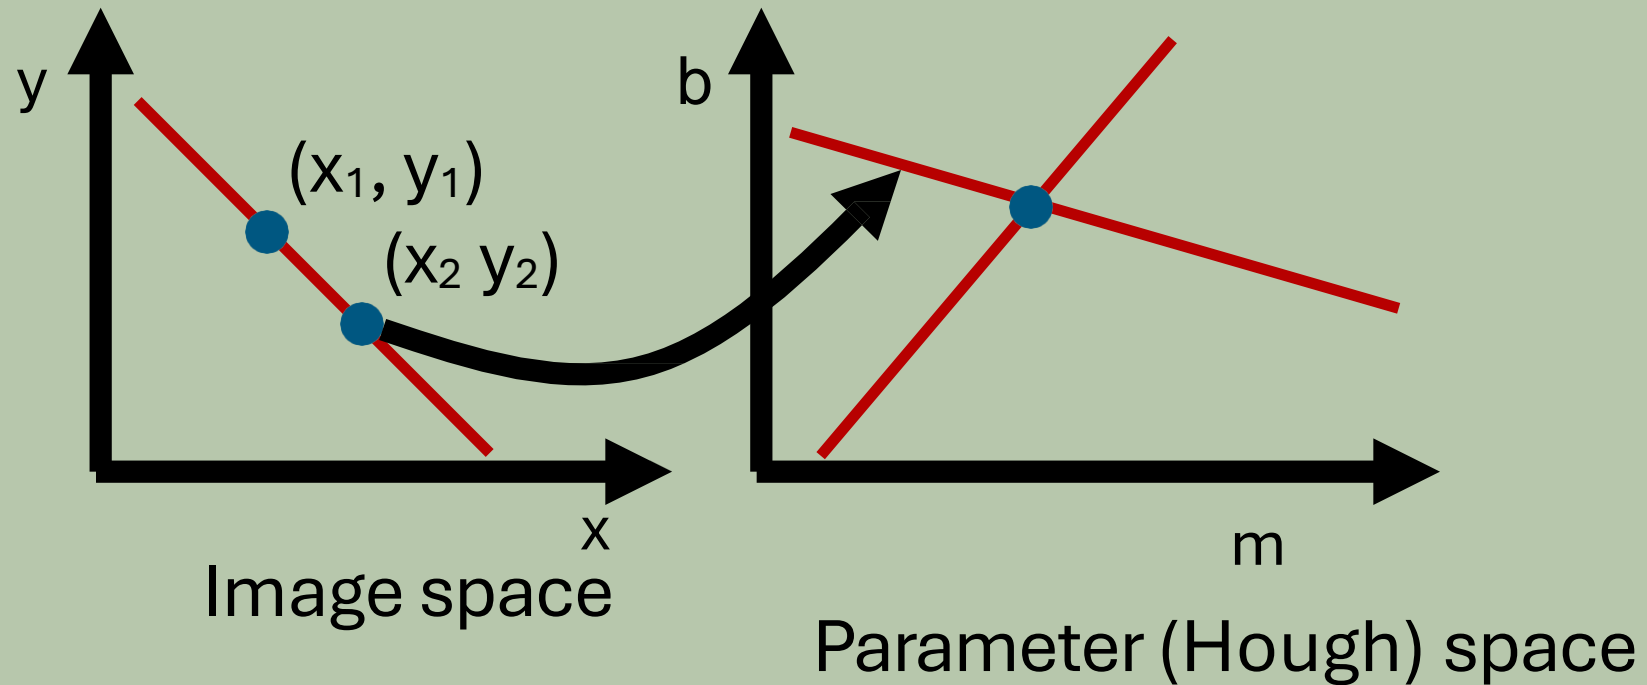


A point (x_2, y_2) is mapped to a line in Hough space


$$y_2 = m x_2 + b$$

$$b = -x_2 m + y_2$$

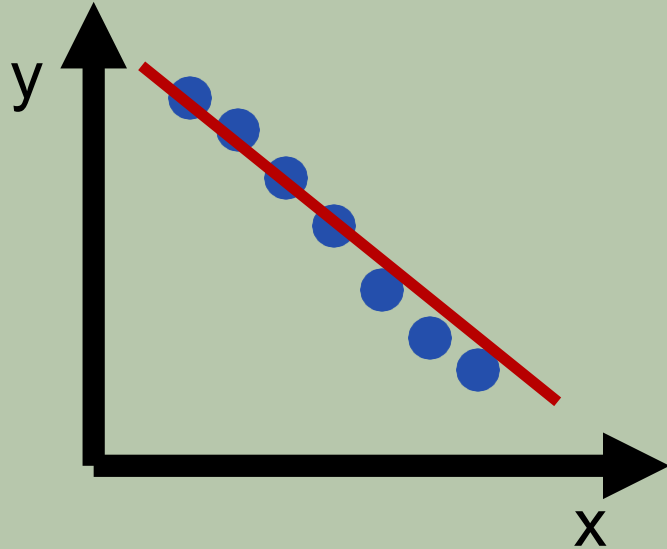
The intersection point satisfies both (x_1, y_1) and (x_2, y_2) , thus the line





A point (x_3, y_3) is mapped to a line in Hough space



Hough Transform: Accumulator



- For each pixel, draw a line in the discretised Hough space, assigning a value of one to all the discretised positions it passes through
- Process all image pixels

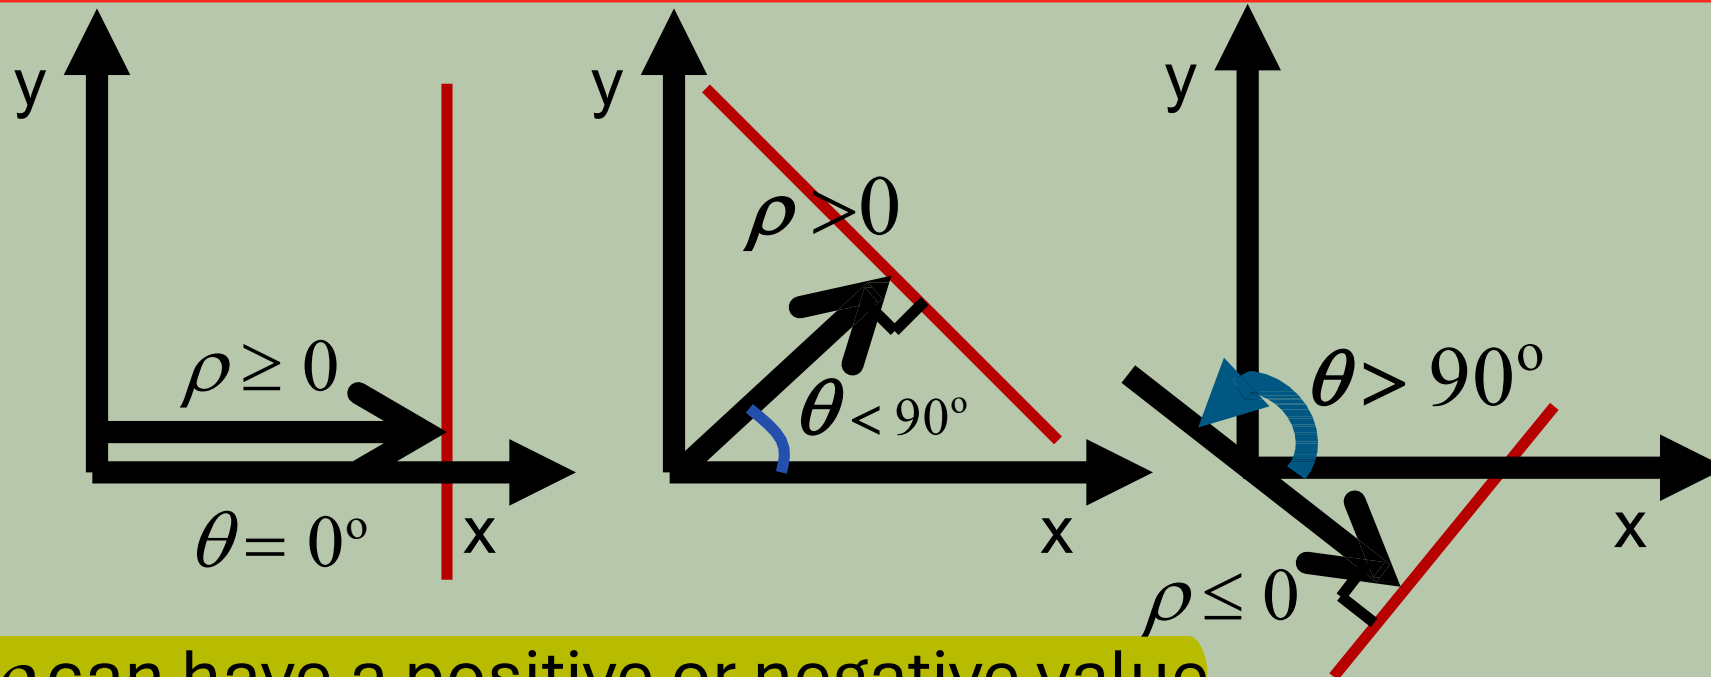


Hough Transform

- Can detect multiple lines in an image
 - from multiple local maxima
- Can easily be extended for circles and ellipses
- Computationally efficient
- Problem: (m, b) are unbounded
 - E.g., the slope parameter m can have an infinite value

Hough Transform: Polar Form

- Use a polar representation for the parameter space
- $x \cos \theta + y \sin \theta = \rho \rightarrow y = -\frac{\cos \theta}{\sin \theta} x + \frac{\rho}{\sin \theta} \quad (0 \leq \theta \leq 180^\circ)$



- Note: ρ can have a positive or negative value

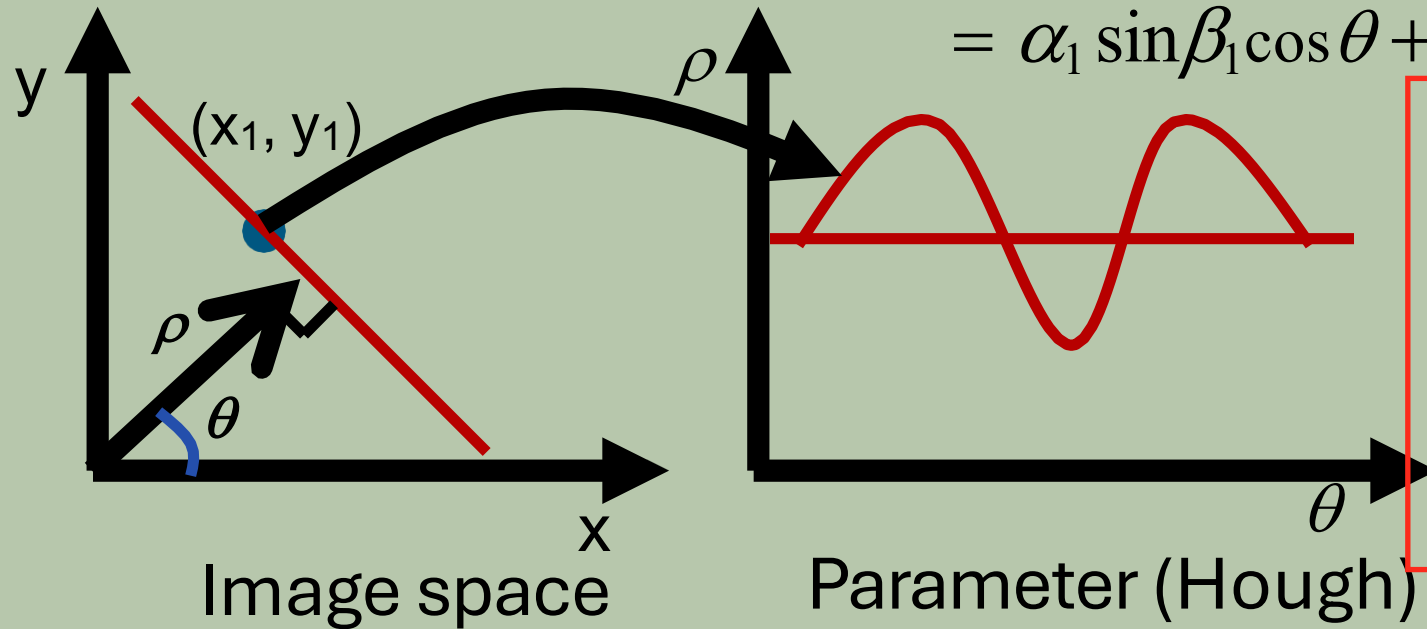
A point (x_1, y_1) is mapped to a **sinusoid** in the polar parameter space

$$x_1 \cos \theta + y_1 \sin \theta = \rho$$

$$\rho = x_1 \cos \theta + y_1 \sin \theta$$

$$= \alpha_1 \sin(\theta + \beta_1)$$

$$= \alpha_1 \sin \beta_1 \cos \theta + \alpha_1 \cos \beta_1 \sin \theta$$



$$\alpha_1 = \sqrt{x_1^2 + y_1^2}$$

$$\sin \beta_1 = \frac{x_1}{\sqrt{x_1^2 + y_1^2}}$$

$$\cos \beta_1 = \frac{y_1}{\sqrt{x_1^2 + y_1^2}}$$

$$\beta_1 = \tan^{-1} \frac{x_1}{y_1}$$

Parameter (Hough) space


$$0 \leq \theta \leq 180^\circ$$

How to map: just divide the right-hand side by $\sqrt{x_1^2 + y_1^2}$ and multiply by the same

A point (x_2, y_2) is mapped to a sinusoid in the polar parameter space

$$x_2 \cos \theta + y_2 \sin \theta = \rho$$

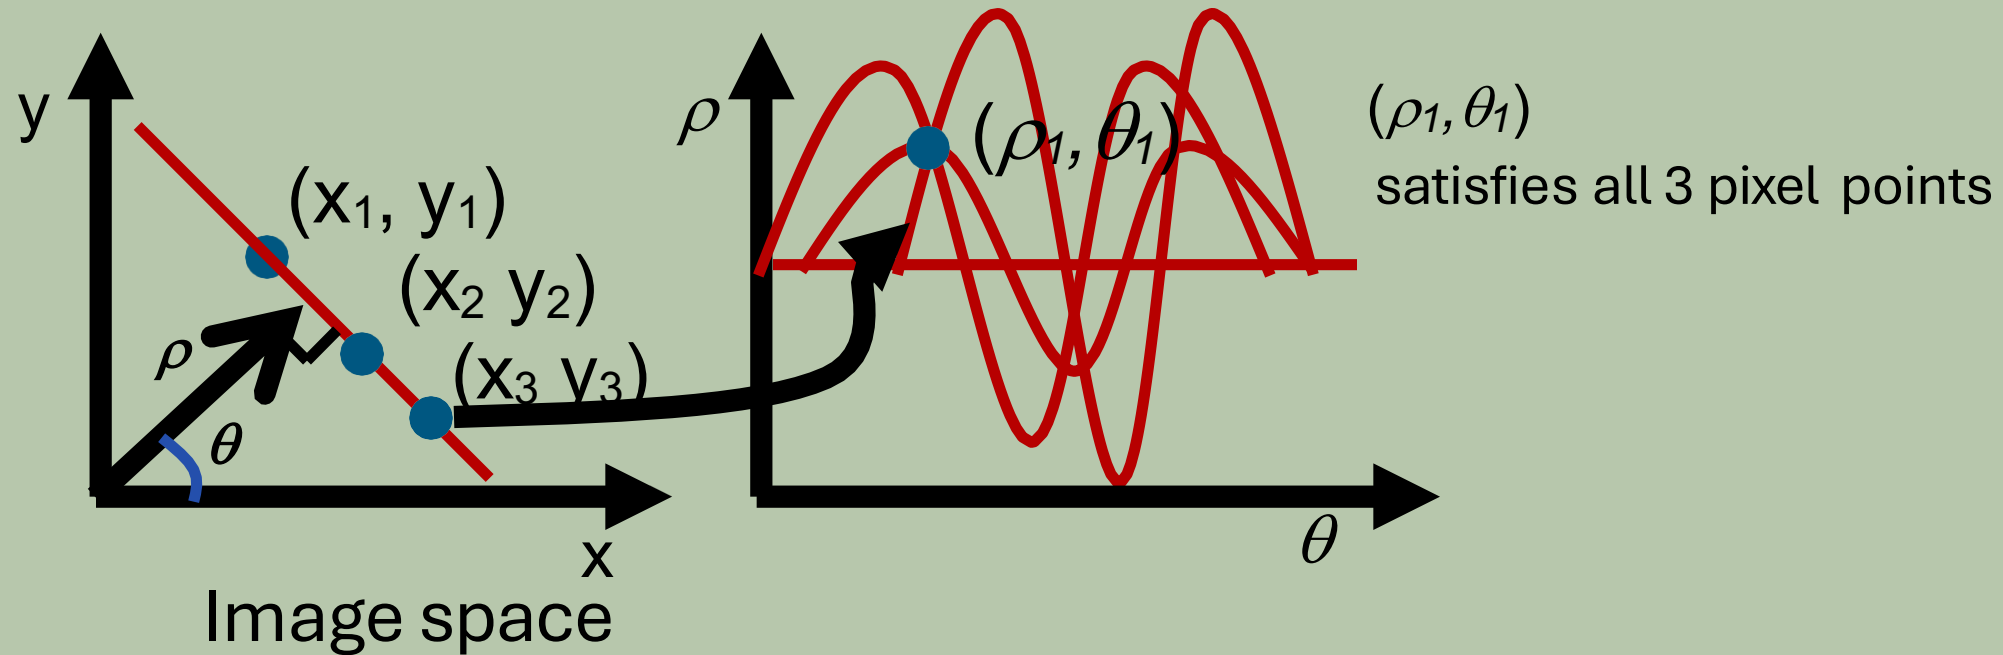
$$\begin{aligned} \rho &= x_2 \cos \theta + y_2 \sin \theta \\ &= \alpha_2 \sin(\theta + \beta_2) \end{aligned}$$



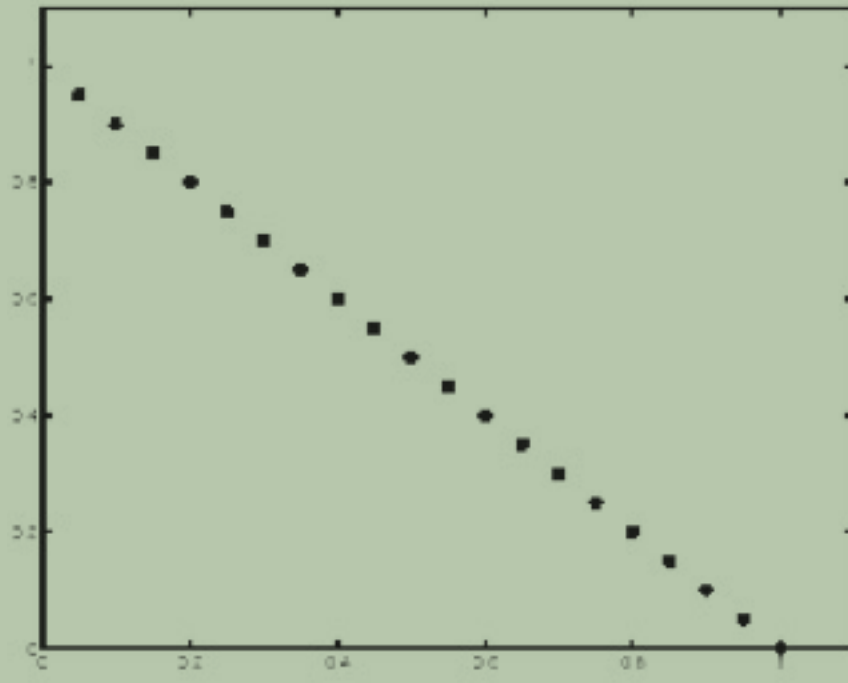
A point (x_3, y_3) is mapped to a sinusoid in the polar parameter space

$$x_3 \cos \theta + y_3 \sin \theta = \rho$$

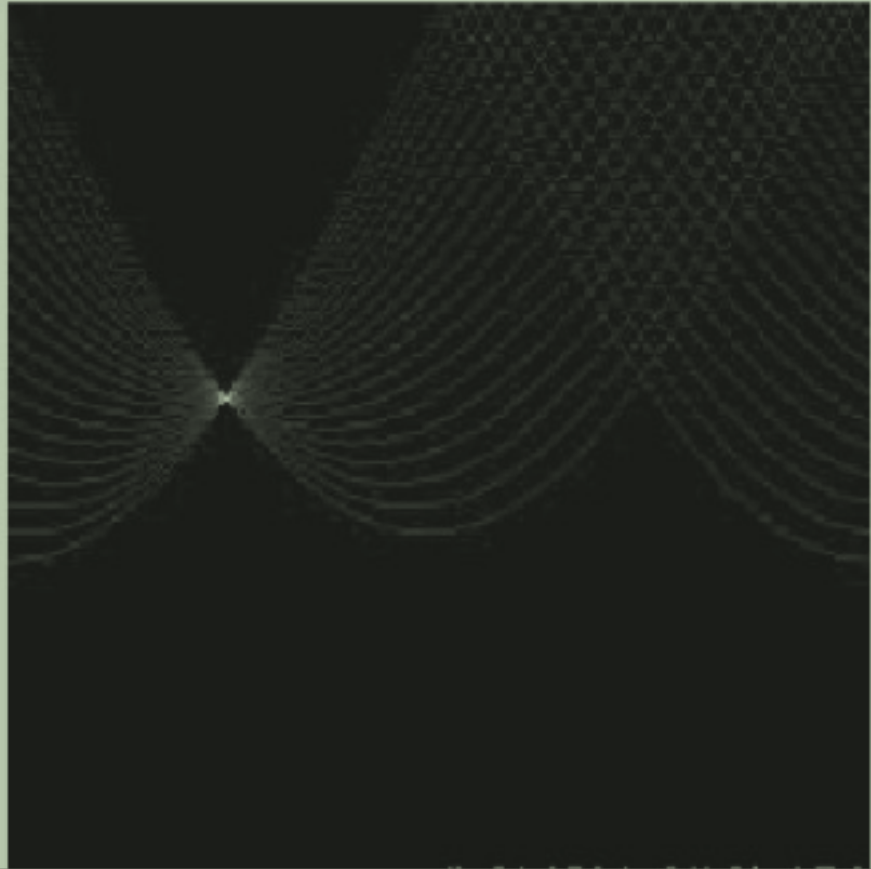
$$\begin{aligned} \rho &= x_3 \cos \theta + y_3 \sin \theta \\ &= \alpha_3 \sin(\theta + \beta_3) \end{aligned}$$



Hough Transform: Example

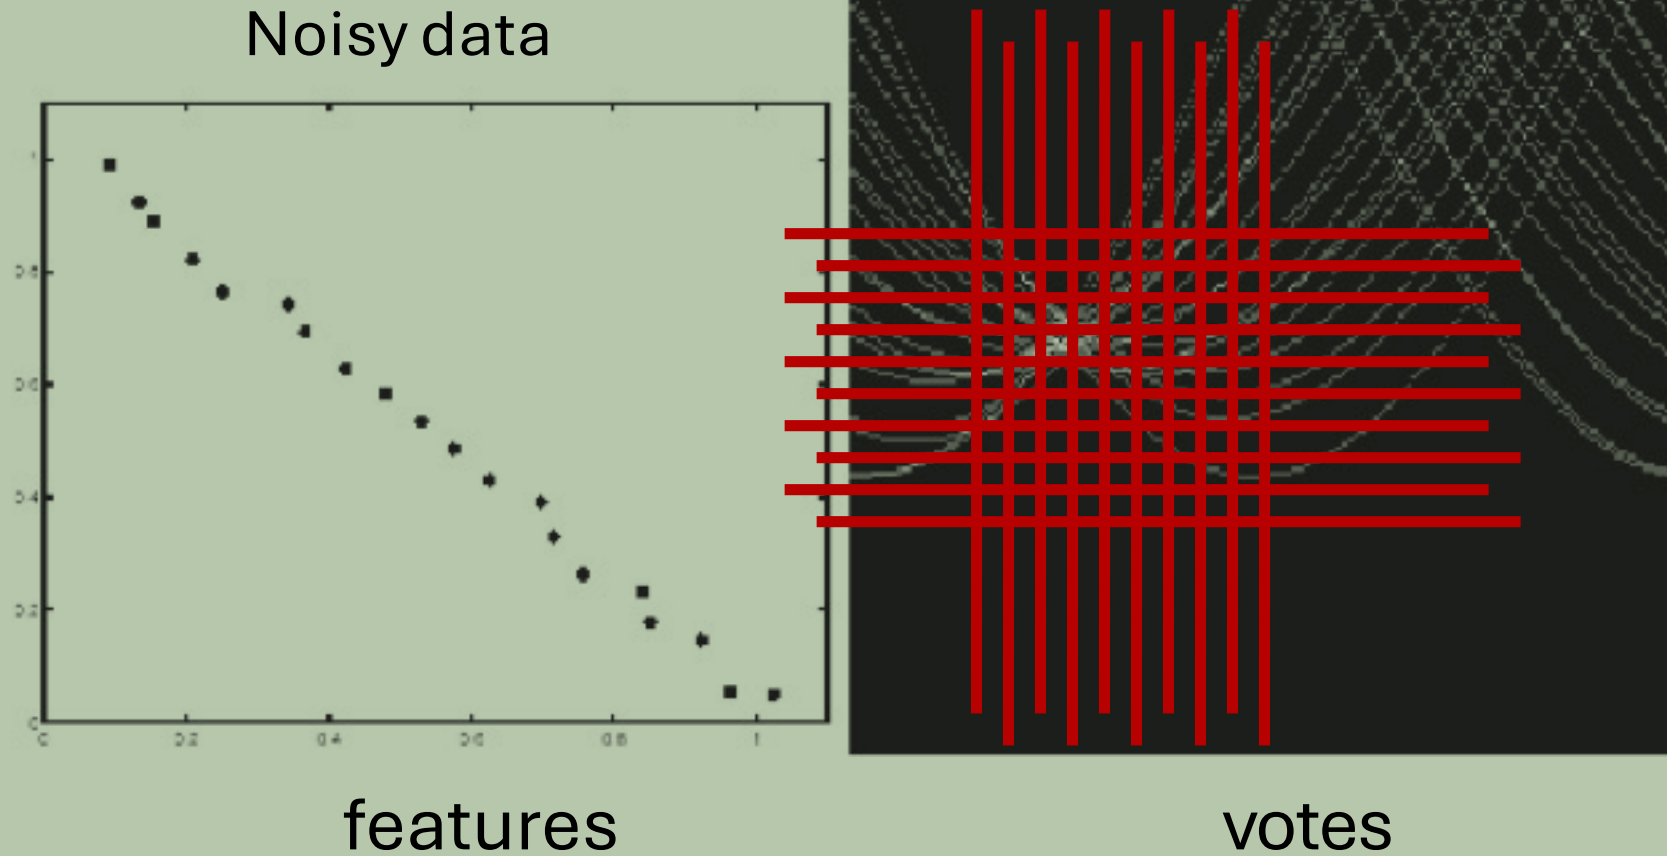


features



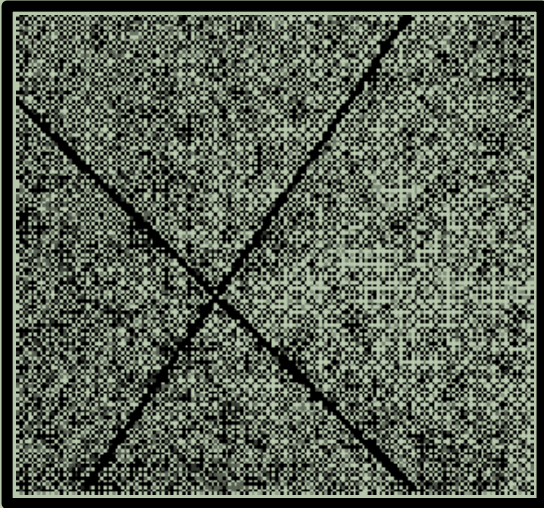
votes

Hough Transform: Example

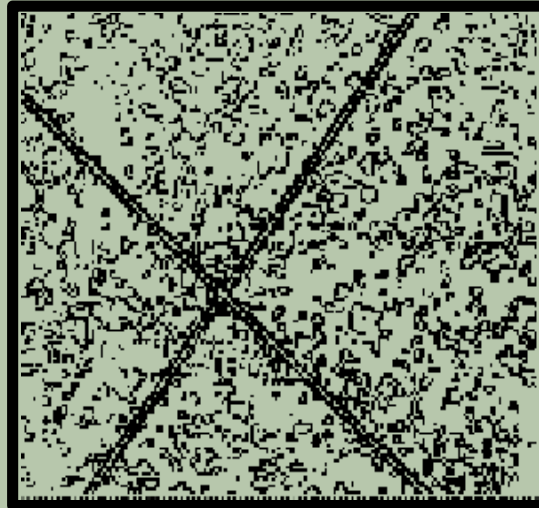


Issue: the grid size needs to be adjusted...

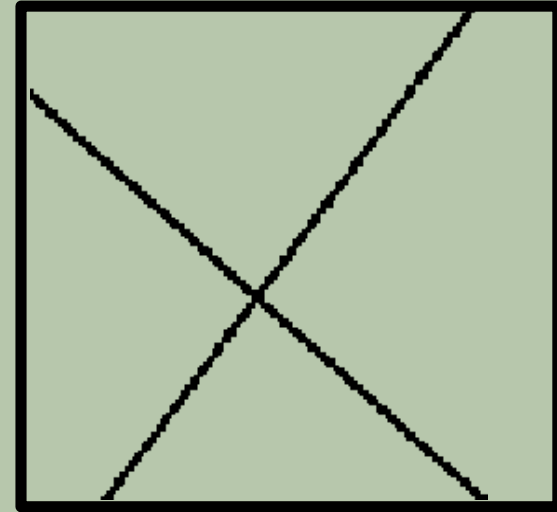
Hough Transform: Example



Image



Edge Detection

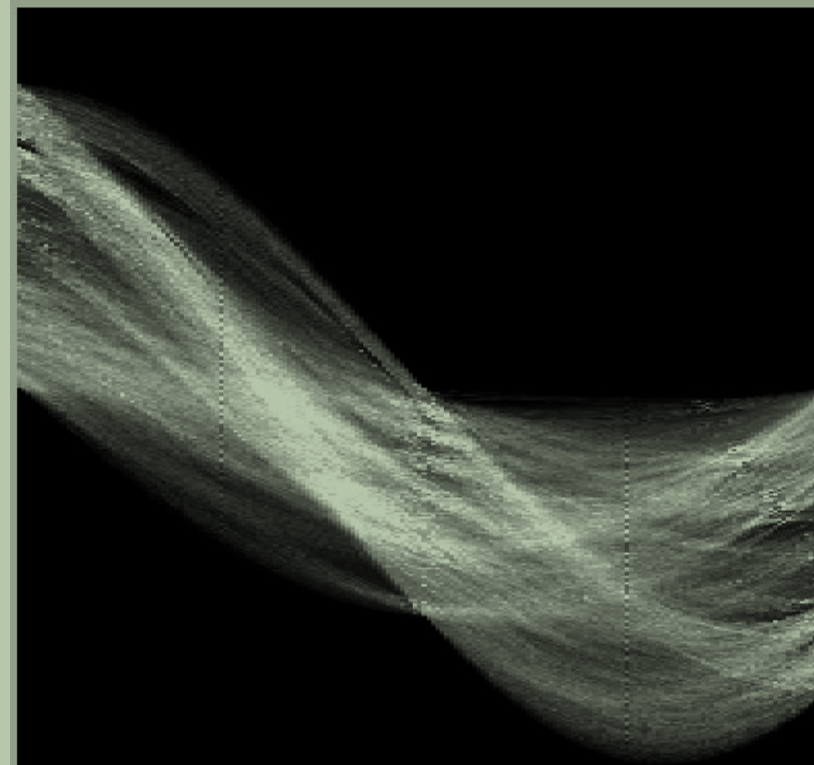


Hough Transform

Hough Transform: 1) Image \rightarrow Canny Edges



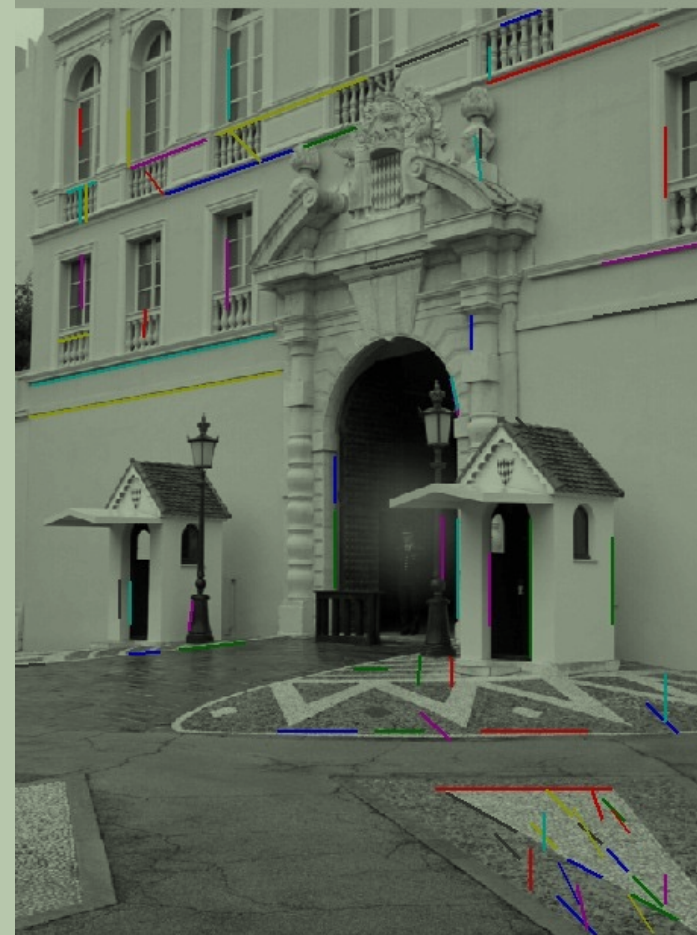
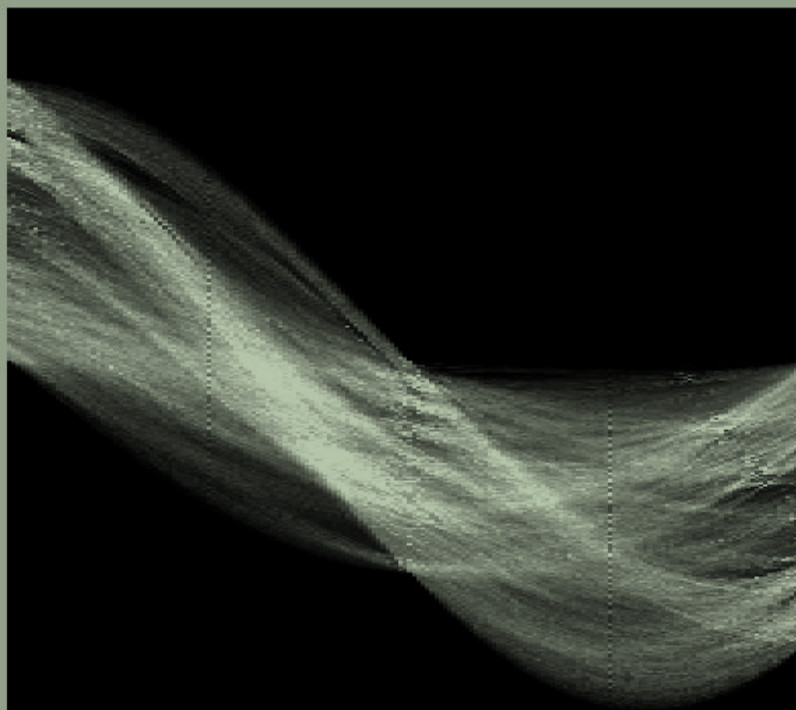
Hough Transform: 2) Canny Edges \rightarrow Votes



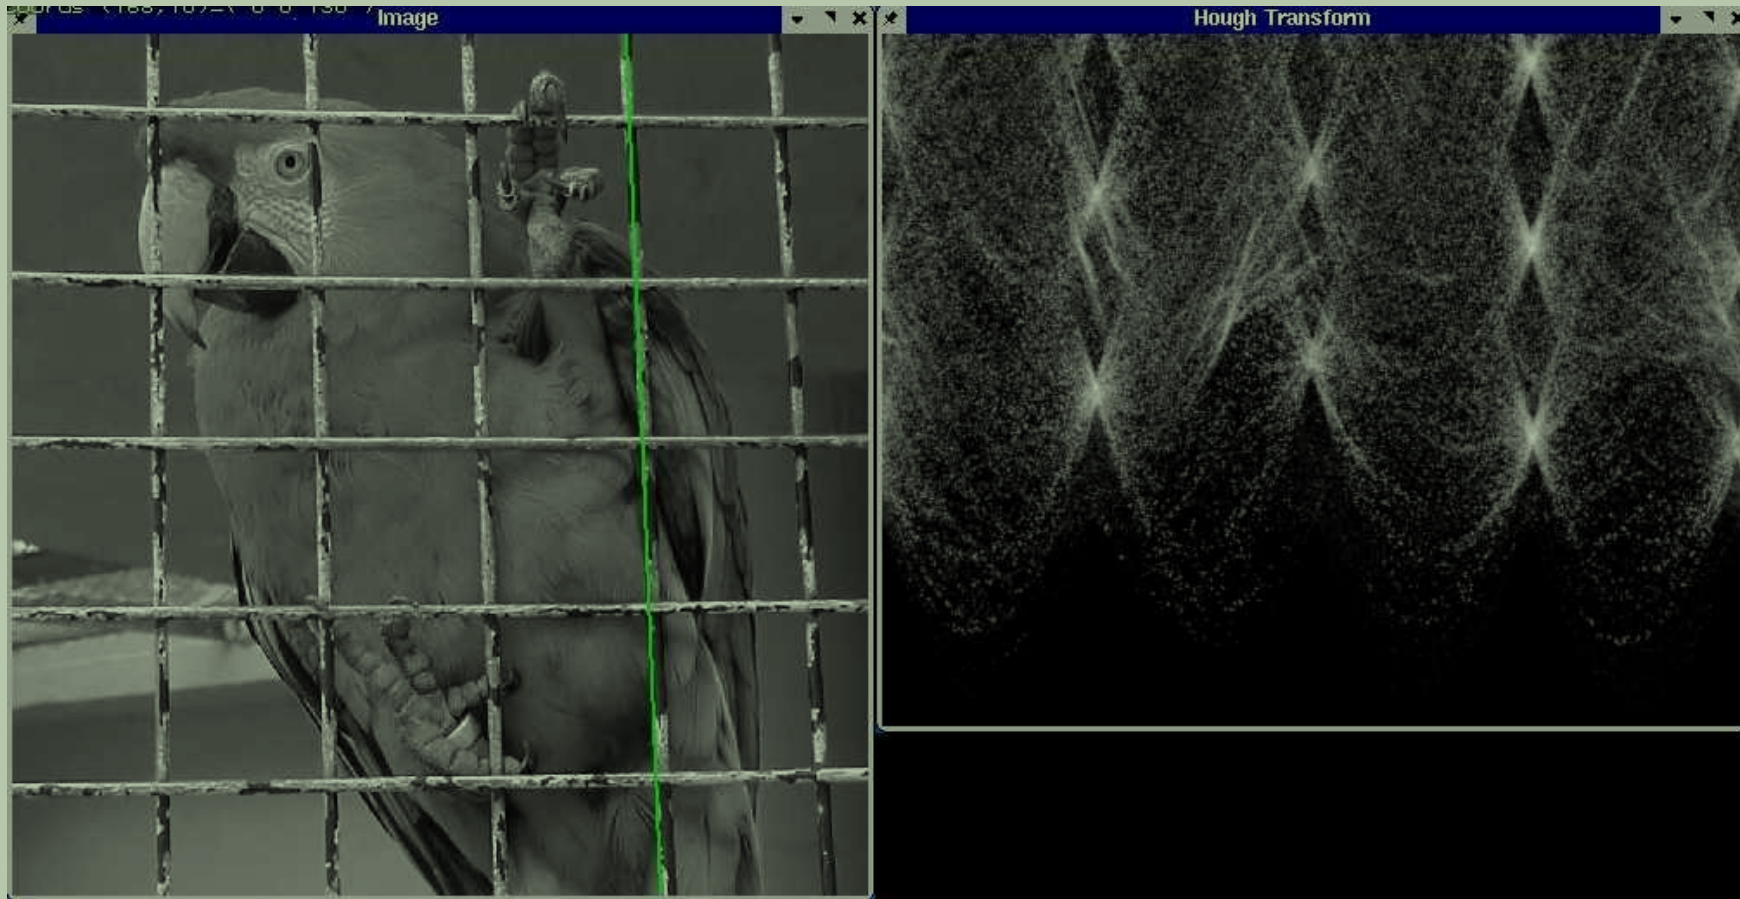
Hough Transform: 3) Votes \rightarrow Lines

threshold + non-maximum suppression

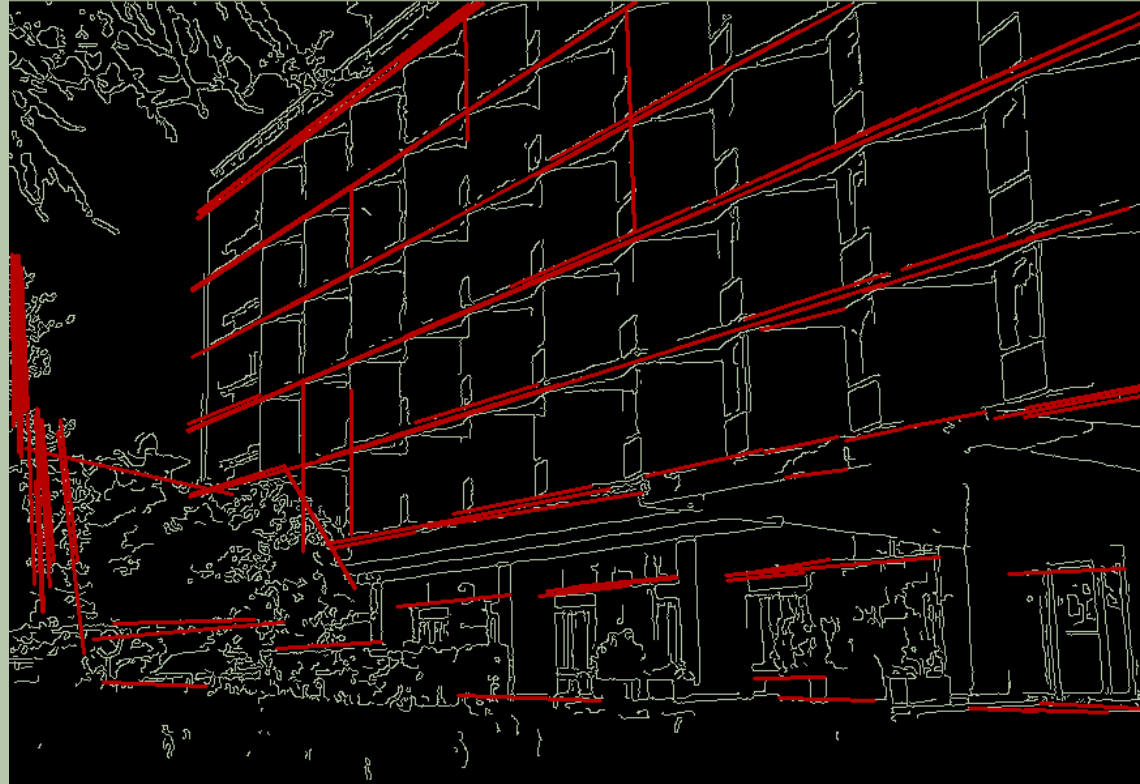
- Find peaks and post-process



Hough Transform: Example



Hough Transform: Example



Hough Transform: Pros and Cons

Pros:

- All points processed **independently**, so can **cope with occlusion**
- Some **robustness to noise**: noisy points are unlikely to contribute consistently to any single bin
- Can detect **multiple instances** of line/circle/ellipse in a **single pass**

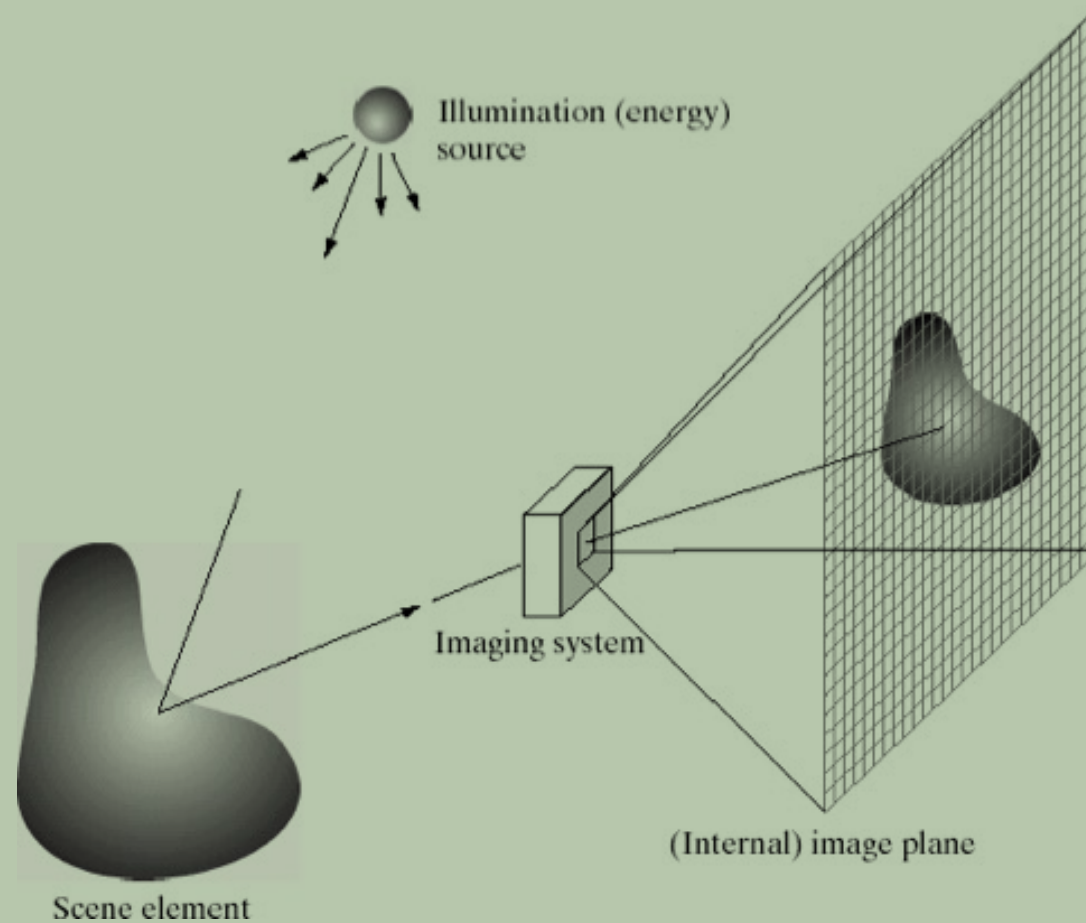
Cons:

- **Complexity of search time** increases exponentially with the **number of model parameters**
- **Non-target shapes** can produce **spurious peaks** in parameter space
- Quantization: hard to pick a **grid size**

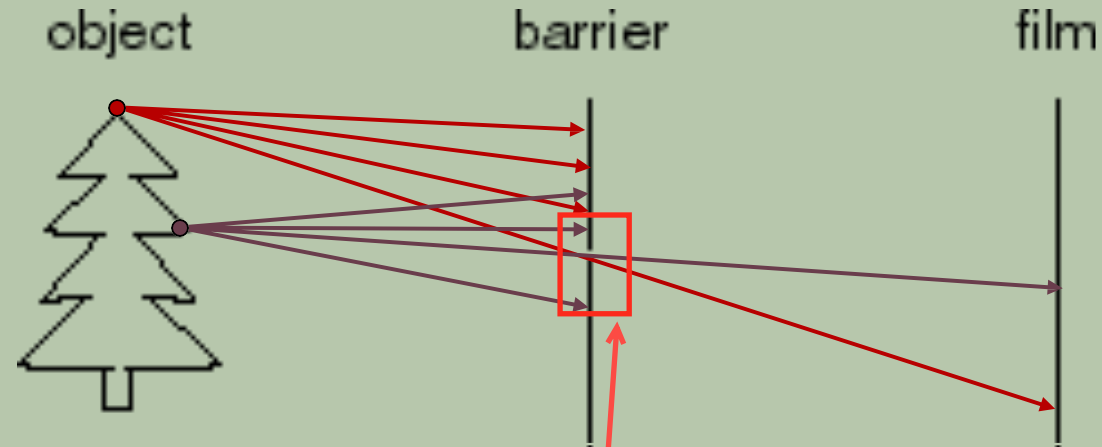
Image Formation (Review)

Pinhole Camera Model

Photometric Image Formation

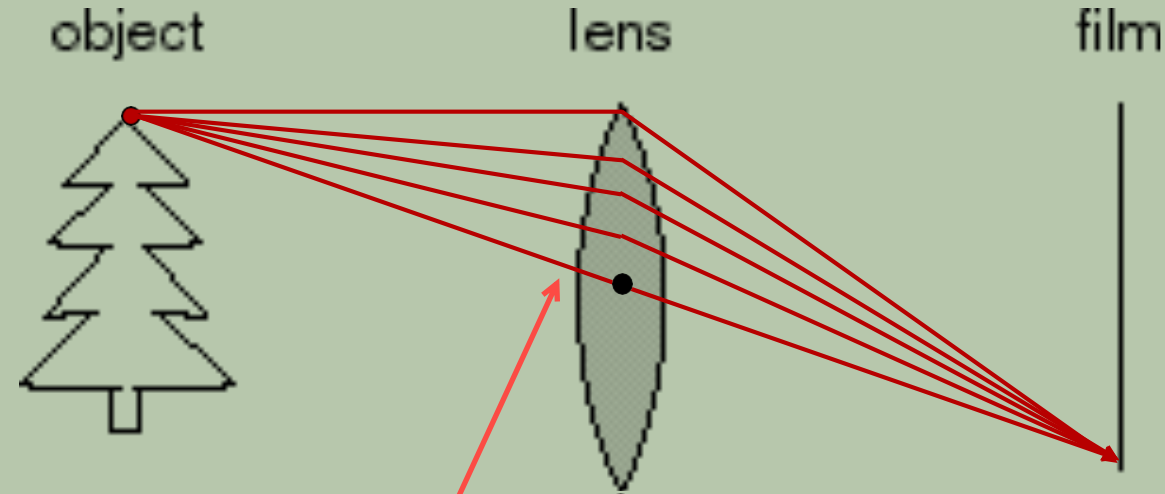


Pinhole Camera



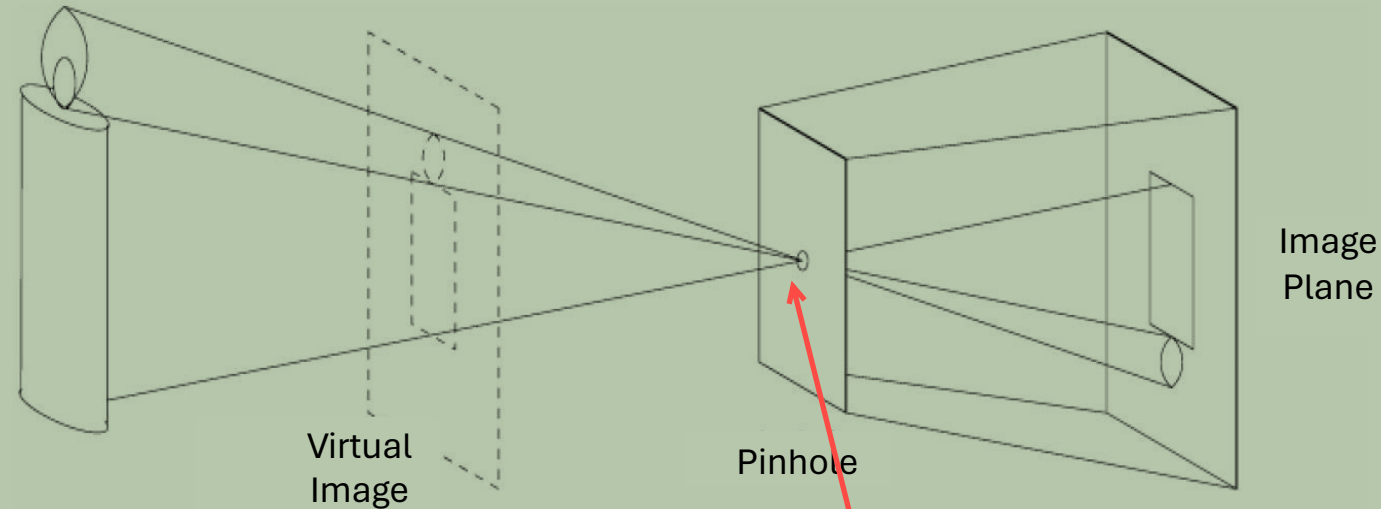
- Add a barrier to block off most of the rays
- This reduces blurring
- The opening is known as the aperture

Lensed Camera



- A lens focuses light onto the film: captures more light
- Rays passing through the centre do not deviate

Pinhole Camera Model



- Pinhole camera is an abstract model to approximate the imaging process: **perspective projection**
- If we treat the pinhole as a point, only one ray from any given point can enter the camera

Camera Projection Matrix and Single View Geometry

Homogeneous Coordinates

from Cartesian to homogeneous

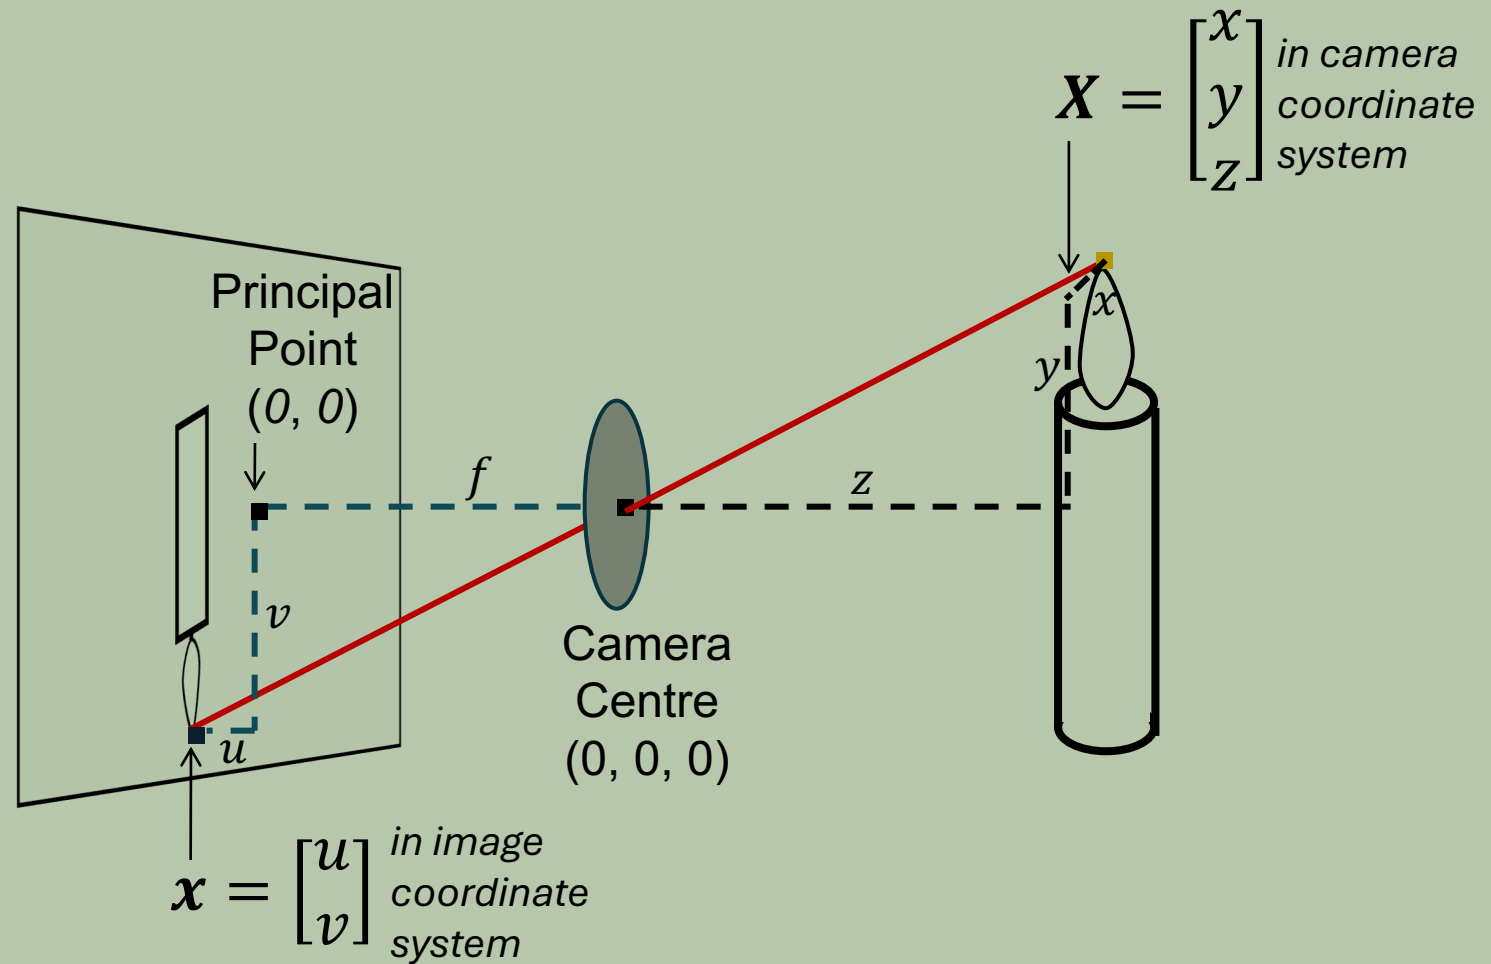
- To homogeneous: $(x, y) \rightarrow (x, y, 1)$ $(x, y, z) \rightarrow (x, y, z, 1)$
- From homogeneous: $(x, y, w) \rightarrow (x/w, y/w)$ $(x, y, z, w) \rightarrow (x/w, y/w, z/w)$

from homogeneous to Cartesian

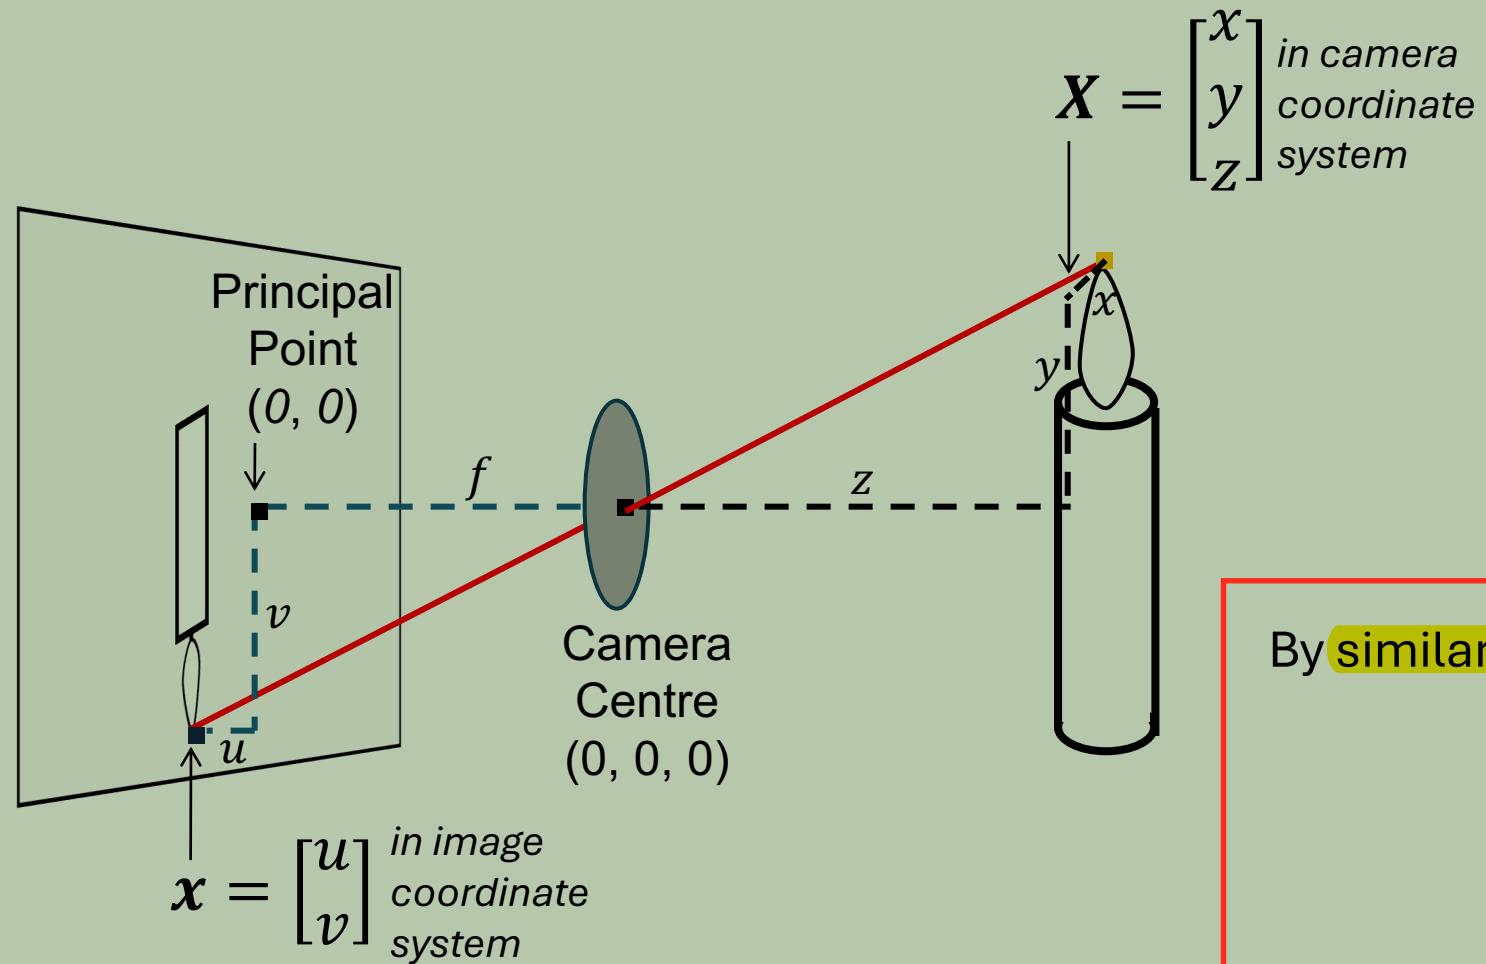
- Invariant to scaling: $k \begin{bmatrix} x \\ y \\ w \end{bmatrix} = \begin{bmatrix} kx \\ ky \\ kw \end{bmatrix} \Rightarrow \begin{bmatrix} \frac{kx}{kw} \\ \frac{ky}{kw} \\ \frac{kw}{kw} \end{bmatrix} = \begin{bmatrix} \frac{x}{w} \\ \frac{y}{w} \\ 1 \end{bmatrix}$
Homogeneous Cartesian

- A **point** in Cartesian coordinates is a **ray** in homogeneous coords
k in (-inf, +inf)

Perspective Projection



Perspective Projection



By similar triangles:

$$u = f \frac{x}{z}$$

$$v = f \frac{y}{z}$$

Perspective Projection

- $(x, y, z) \mapsto \left(\frac{fx}{z}, \frac{fy}{z}\right)$ In Cartesian coord: 3D \rightarrow 2D

- Using **homogeneous** coordinates:

$$2D: \left(\frac{fx}{z}, \frac{fy}{z}\right) \mapsto \left(\frac{fx}{z}, \frac{fy}{z}, 1\right) = (fx, fy, z)$$

$$3D: (x, y, z) \mapsto (x, y, z, 1)$$

Cart.

Homo.

- Linear projection in homogeneous coordinates!

In homogeneous coord:
3D \rightarrow 2D

$$\begin{matrix} 3D \\ \begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix} \end{matrix} \mapsto \begin{matrix} 2D \\ \boxed{\begin{pmatrix} fx \\ fy \\ z \end{pmatrix}} \end{matrix} = \begin{bmatrix} f & 0 & 0 \\ 0 & f & 0 \\ 0 & 0 & 1 \end{bmatrix} \begin{matrix} 3D \\ \boxed{\begin{pmatrix} x \\ y \\ z \\ 1 \end{pmatrix}} \end{matrix}$$

2D projected coord. in homo. coord.

3D point coord. in homo. coord.

Perspective Projection

Cart. to Homo. for 2D

2D in homo. coord.

$[fx/z, fy/z].T$

$$\begin{pmatrix} fx \\ fy \\ z \end{pmatrix} = \begin{bmatrix} f & & \\ & f & \\ & & 1 \end{bmatrix} \begin{bmatrix} 1 & & 0 \\ & 1 & 0 \\ & & 1 & 0 \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$

Homo. to Cart. for 3D

3D in homo. coord.

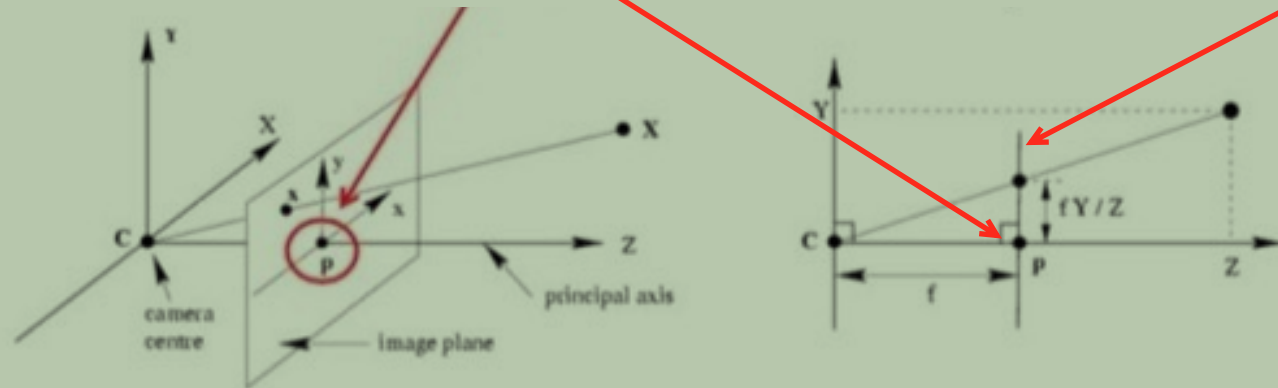
$[x, y, z].T$

$$\Leftrightarrow \mathbf{x} = \mathbf{P}\mathbf{X} \text{ with } \mathbf{P} = \text{diag}(f, f, 1) [\mathbf{I} \mid 0]$$

= 3×4 homogeneous camera projection matrix

Principal Point

- The point where the principal axis intersects with the image plane



Principal Point Offset


<https://au.mathworks.com/help/vision/ug/camera-calibration.html>

- So far, we have assumed that the **origin of points** in the image plane is at **principal point**

- However, origin is often elsewhere (e.g., at image corner)

- ~~Inhomogeneous:~~ ^{Cartesian.} ^{3D} $(X, Y, Z) \mapsto (fX/Z + p_x, fY/Z + p_y)$ ^{2D}

- **Homogeneous:** ^{3D} $\begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} \mapsto \begin{pmatrix} fX + Zp_x \\ fY + Zp_y \\ Z \end{pmatrix} = \begin{bmatrix} f & p_x & 0 \\ f & p_y & 0 \\ 1 & 0 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$ ^{proj. mat.} ^{3D}



Camera (Intrinsic) Calibration Matrix

x_{img} (homo)

x_{world} (homo)

$$\begin{pmatrix} fX + Zp_x \\ fY + Zp_y \\ Z \end{pmatrix} = \begin{bmatrix} f & p_x & 0 \\ & f & p_y \\ & & 1 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

calibration matrix

$$= \begin{bmatrix} f & p_x \\ & f & p_y \\ & & 1 \end{bmatrix} \begin{bmatrix} 1 & 0 & 0 \\ & 1 & 0 \\ & & 1 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix}$$

$$= \mathbf{K}[I \mid 0]\mathbf{X}$$

$$\mathbf{K} = \begin{bmatrix} f & p_x \\ & f & p_y \\ & & 1 \end{bmatrix} =$$

camera
calibration matrix

Rectangular Pixels: CCD Cameras


- Charge-Coupled Device (CCD)
- From image plane to pixel coordinates ^{measured by # pixels}



$$\mathbf{K} = \begin{bmatrix} \alpha_x & x_0 \\ & \alpha_y & y_0 \\ & & 1 \end{bmatrix} = \begin{bmatrix} m_x & & \\ & m_y & \\ & & 1 \end{bmatrix} \begin{bmatrix} f & p_x \\ & f & p_y \\ & & 1 \end{bmatrix}$$

$$\begin{cases} x_0 = m_x p_x \\ y_0 = m_y p_y \\ \alpha_x = m_x f \\ \alpha_y = m_y f \end{cases}, \text{ with } \begin{cases} m_x = \# \text{ pixels / unit distance along } x \\ m_y = \# \text{ pixels / unit distance along } y \end{cases} \begin{matrix} \text{offsets} \times \# \text{ pixels per offset} \\ \text{on the image plane} \end{matrix}$$

Camera Parameters



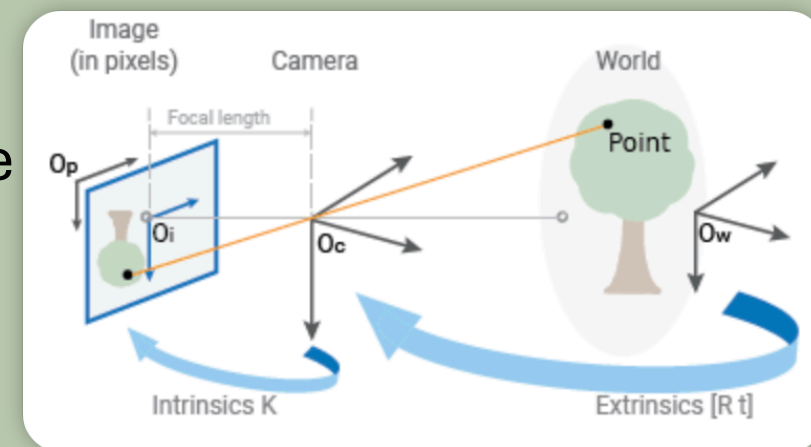
- So far: **intrinsic** camera parameters

- How to map spatial directions to pixel coordinates ^{cam → img}
- ^f Focal length, ^{p_x, p_y (offsets)} principal point, ^{m_x, m_y} pixel width/height

$$\mathbf{K} = \begin{bmatrix} \alpha_x & & x_0 \\ & \alpha_y & y_0 \\ & & 1 \end{bmatrix} = \begin{bmatrix} m_x & & \\ & m_y & \\ & & 1 \end{bmatrix} \begin{bmatrix} f & & p_x \\ & f & p_y \\ & & 1 \end{bmatrix}$$

- What else? **Extrinsic** camera parameters

- How to transform a 3D point into the camera frame ^{world → cam}
- Depends on the camera rotation and translation



Extrinsic Parameters

The extrinsic parameters consist of a **rotation, R** , and a **translation, t** . The origin of the camera's coordinate system is at its optical center and its x - and y -axis define the image plane.

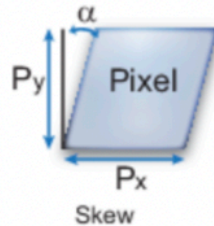


Intrinsic Parameters

The intrinsic parameters include the **focal length**, the **optical center**, also known as the *principal point*, and the **skew coefficient**. The camera intrinsic matrix, K , is defined as:

$$\begin{bmatrix} f_x & s & c_x \\ 0 & f_y & c_y \\ 0 & 0 & 1 \end{bmatrix}$$

The pixel skew is defined as:



$\begin{bmatrix} c_x & c_y \end{bmatrix}$ – Optical center (the principal point), in pixels.

(f_x, f_y) – Focal length in pixels.

$$f_x = F/p_x$$

$$f_y = F/p_y$$

F – Focal length in world units, typically expressed in millimeters.


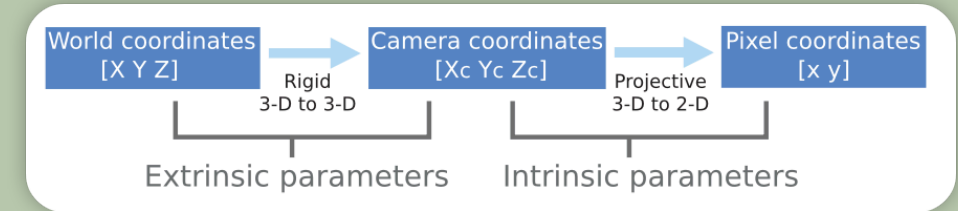
(p_x, p_y) – Size of the pixel in world units.

s – Skew coefficient, which is non-zero if the image axes are not perpendicular.

$$s = f_x \tan \alpha$$

Extrinsic Camera Parameters

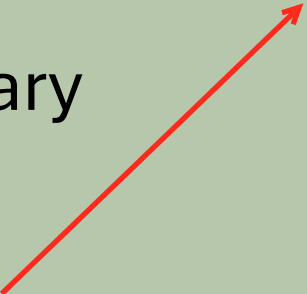
- Camera rotation and translation
- **R: Rotation matrix**
 - Orthogonal + unit determinant
 - $SO(3)$
- **C: Camera centre (vector)** C_{world}
 - Location of the camera in the world coordinate system
($X_{world} - C_{world}$)
- $X_{cam} = R(X - C)$ (in camera coordinates)



$$X_{cam} = \begin{bmatrix} R & -RC \\ 0 & 1 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} = \begin{bmatrix} R & -RC \\ 0 & 1 \end{bmatrix} X_{world} \quad (\text{in homogeneous coordinates})$$

Rotation About Coordinate Axes in 3D

- Express 3D rotation as series of rotations around coordinate axes by angles α, β, γ
- The overall rotation is the product of these elementary rotations:
- $\mathbf{R} = \mathbf{R}_x \mathbf{R}_y \mathbf{R}_z$ rotation matrix
- They describe clockwise rotations


$$\begin{aligned}\mathbf{R}_x(\alpha) &= \begin{bmatrix} 1 & 0 & 0 \\ 0 & \cos \alpha & -\sin \alpha \\ 0 & \sin \alpha & \cos \alpha \end{bmatrix} \\ \mathbf{R}_y(\beta) &= \begin{bmatrix} \cos \beta & 0 & \sin \beta \\ 0 & 1 & 0 \\ -\sin \beta & 0 & \cos \beta \end{bmatrix} \\ \mathbf{R}_z(\gamma) &= \begin{bmatrix} \cos \gamma & -\sin \gamma & 0 \\ \sin \gamma & \cos \gamma & 0 \\ 0 & 0 & 1 \end{bmatrix}\end{aligned}$$

Complete Camera Matrix

3*1 (homo)
X_img

4*1 (homo)
X_cam

$$\bullet \mathbf{X} = \mathbf{K} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \mathbf{X}_{\text{cam}}$$

3*3 K (intrinsic)

$$= \begin{bmatrix} \alpha_x & & x_0 \\ & \alpha_y & y_0 \\ & & 1 \end{bmatrix} \begin{bmatrix} \mathbf{I} & \mathbf{0} \end{bmatrix} \begin{bmatrix} \mathbf{R} & -\mathbf{RC} \\ \mathbf{0} & 1 \end{bmatrix} \mathbf{X}$$

4*1 (homo)
X_world

$$= \underbrace{\mathbf{K}}_{\substack{\text{Camera} \\ \text{intrinsics} \\ 3*3}} \underbrace{\mathbf{R} \begin{bmatrix} \mathbf{I} & -\mathbf{C} \end{bmatrix}}_{\substack{\text{Camera} \\ \text{extrinsics} \\ 3*4}} \mathbf{X}$$

4*1 (homo)
X_world

$$= \mathbf{P} \mathbf{X}$$

4*1 (homo)
X_cam

↑
3x4 projective camera matrix

$$\begin{cases} x_0 = m_x p_x \\ y_0 = m_y p_y \\ \alpha_x = m_x f \\ \alpha_y = m_y f \end{cases}, \text{ with } \begin{cases} m_x = \# \text{ pixels / unit distance along } x \\ m_y = \# \text{ pixels / unit distance along } y \end{cases} \text{ on the image plane}$$

offsets × #pixels per offset

Camera (Intrinsic) Calibration Matrix

- \mathbf{K} is a 3×3 upper triangular matrix, the “camera calibration matrix”

$$\mathbf{K} = \begin{bmatrix} \alpha_x & & x_0 \\ & \alpha_y & y_0 \\ & & 1 \end{bmatrix} \text{ measured by \# pixels}$$

- Four parameters:
 - The scaling in the image x and y directions, α_x and α_y
 - The principal point (x_0, y_0) , which is the point where the optical axis intersects the image plane

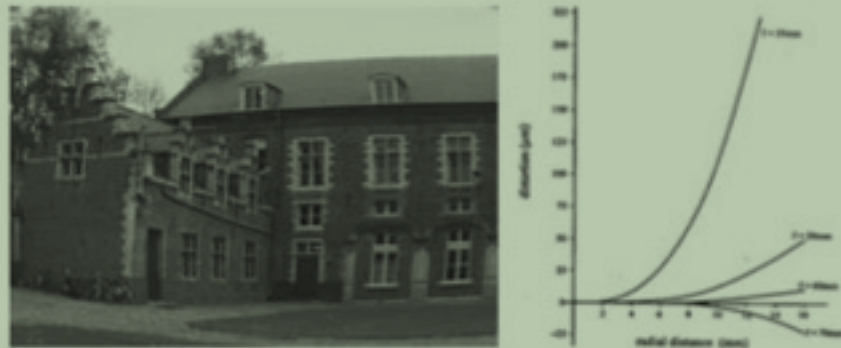
- The aspect ratio is α_y / α_x

$$\begin{cases} x_0 = m_x p_x \text{ offsets} \times \text{\#pixels per offset} \\ y_0 = m_y p_y \\ \alpha_x = m_x f \\ \alpha_y = m_y f \end{cases}, \text{ with } \begin{cases} m_x = \# \text{ pixels / unit distance along } x \\ m_y = \# \text{ pixels / unit distance along } y \end{cases} \text{ on the image plane}$$

Radial Lens Distortion

Radial distortion occurs when light rays bend more near the edges of a lens than they do at its optical center. The smaller the lens, the greater the distortion.

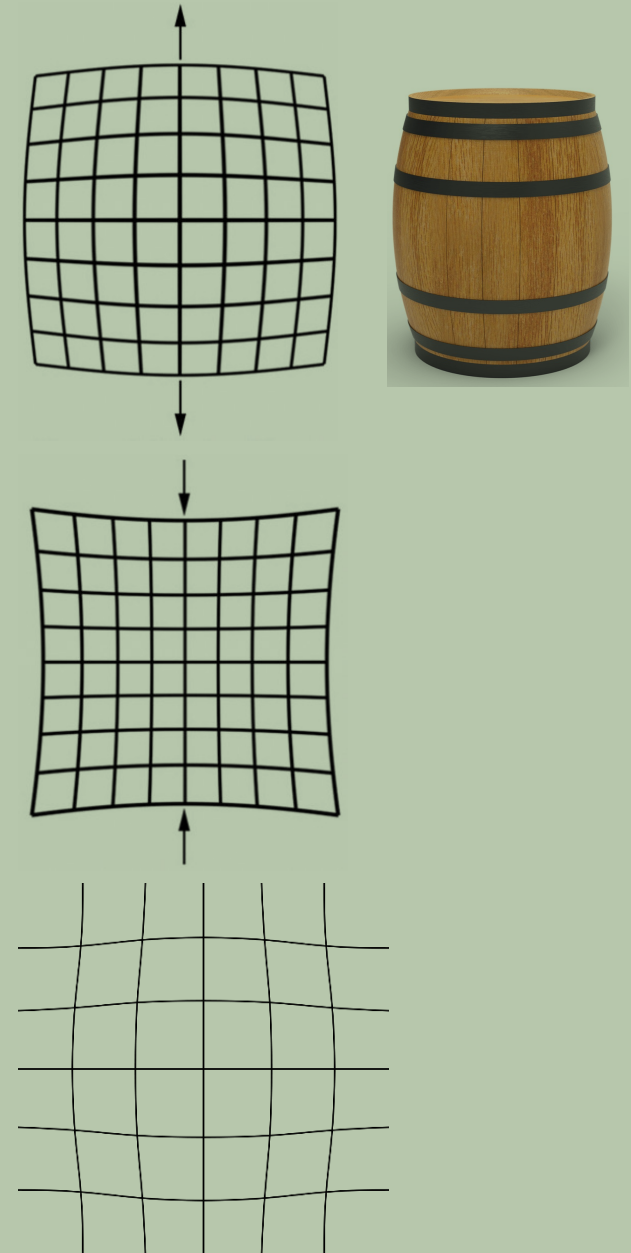
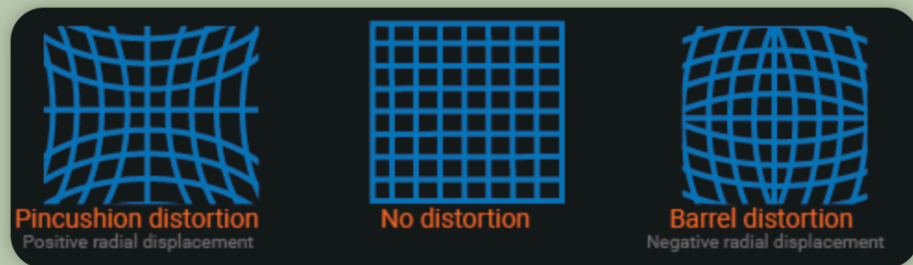
- There is no such thing as a perfect lens
- Straight lines are no longer straight!



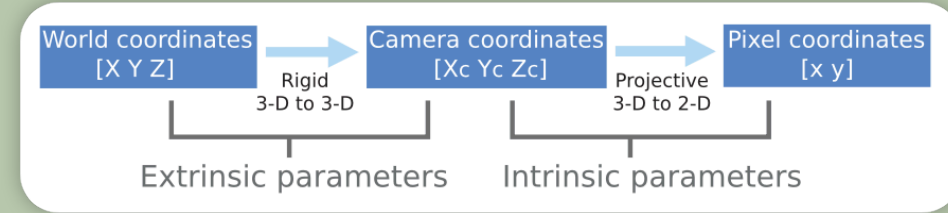
http://foto.hut.fi/opetus/260/luennot/11/kinson_6-11_radial_distortion_zoom_lenses.jpg

Radial Lens Distortion

- Due to spherical lenses (cheaper)
 - **Barrel distortion**
 - Image magnification **decreases** with distance from **optical axis**
 - **Pincushion distortion**
 - Image magnification **increases** with distance from **optical axis**
 - **Mustache distortion**
 - A mixture of both types



Radial Lens Distortion



- Model for radial distortion:

- Change based on distance of point on image plane from principal point

- If $\mathbf{x} = \mathbf{P}\mathbf{X}_{\text{world}} = \mathbf{K}\mathbf{R}[\mathbf{I} \mid -\mathbf{C}]\mathbf{X}_{\text{world}} = \mathbf{K}[\mathbf{I} \mid 0]\mathbf{X}_{\text{cam}}$

$$= \mathbf{K} \begin{bmatrix} X/Z \\ Y/Z \\ 1 \end{bmatrix} = \mathbf{K} \begin{bmatrix} x_{\text{cam}} \\ y_{\text{cam}} \\ 1 \end{bmatrix}$$

$$\mathbf{X}_{\text{cam}} = \begin{bmatrix} \mathbf{R} & -\mathbf{RC} \\ 0 & 1 \end{bmatrix} \begin{pmatrix} X \\ Y \\ Z \\ 1 \end{pmatrix} = \begin{bmatrix} \mathbf{R} & -\mathbf{RC} \\ 0 & 1 \end{bmatrix} \mathbf{X}_{\text{world}} \quad (\text{in homogeneous coordinates})$$

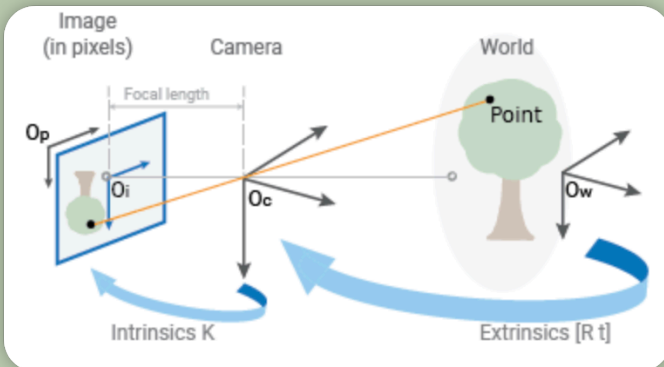
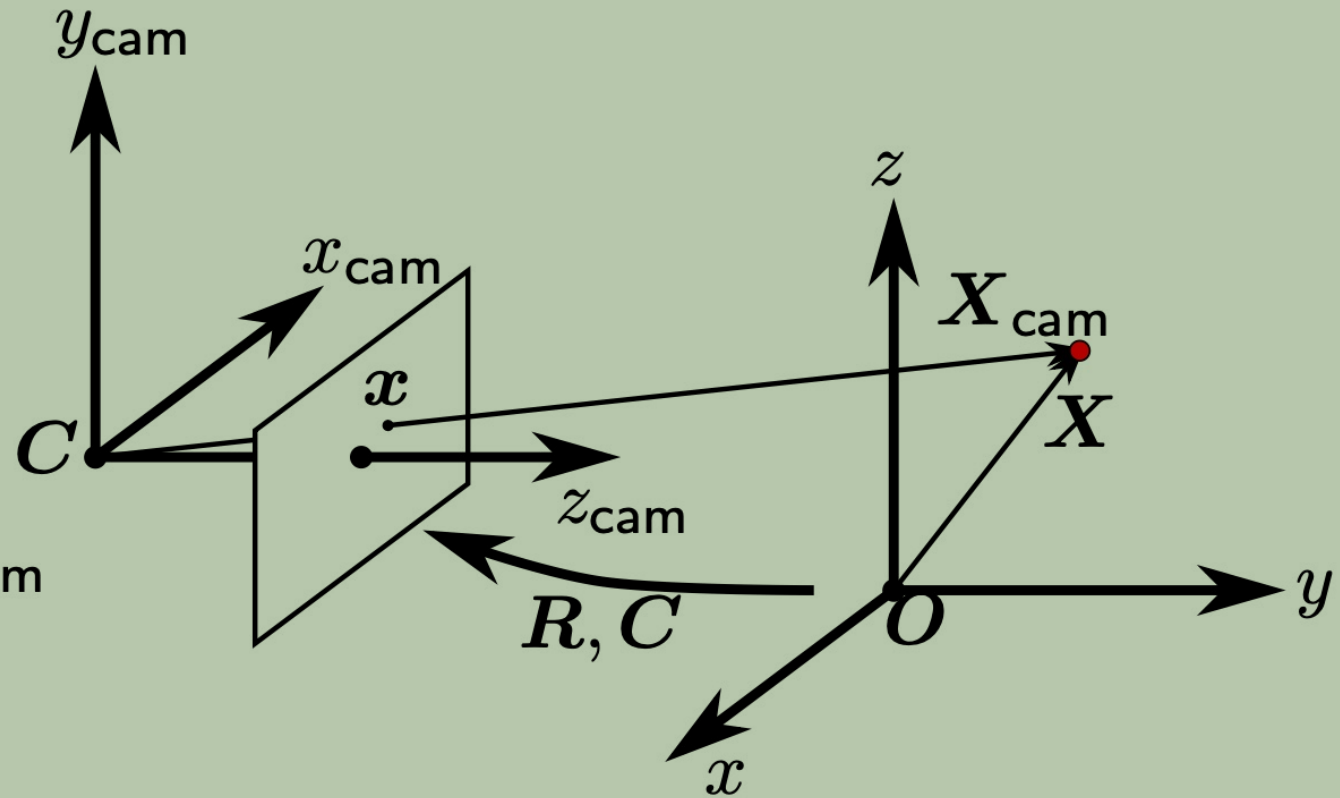
we change to $\mathbf{x} = \mathbf{K} \begin{bmatrix} r \\ r \\ 1 \end{bmatrix} [\mathbf{I} \mid 0]\mathbf{X}_{\text{cam}}$

with $r = 1 + k_1(x_{\text{cam}}^2 + y_{\text{cam}}^2) + k_2(x_{\text{cam}}^2 + y_{\text{cam}}^2)^2$

x_{cam} and y_{cam} are from homo \mathbf{X}_{cam} when $Z=1$ (X, Y divided by Z)


Summary: Camera Projection Matrix

- ▶ image plane
- ▶ camera centre C
- ▶ principal axis z_{cam}
- ▶ image coordinates x
- ▶ world coordinates X
- ▶ camera coordinates X_{cam}



Summary: Camera Projection Matrix

- ▶ image plane
- ▶ camera centre C
- ▶ principal axis z_{cam}
- ▶ image coordinates x
- ▶ world coordinates X
- ▶ camera coordinates X_{cam}



$$\begin{aligned}
 \text{X_img } \hat{x} &= P \hat{X} \\
 &= \underbrace{K}_{\substack{\text{intrinsic} \\ \text{measured by \#pixels}}} \underbrace{R \begin{bmatrix} I & -C \end{bmatrix}}_{\substack{\text{extrinsic} \\ \text{measured by \#pixels}}} \hat{X} = K \begin{bmatrix} R & t \end{bmatrix} \hat{X} \quad \text{X_world} \\
 &= \begin{bmatrix} \alpha & \gamma & u_0 \\ 0 & \beta & v_0 \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \hat{X} \quad \text{X_world} \\
 &= \begin{bmatrix} m_x f_x & \gamma & m_x p_x \\ 0 & m_y f_y & m_y p_y \\ 0 & 0 & 1 \end{bmatrix} \begin{bmatrix} r_{11} & r_{12} & r_{13} & t_x \\ r_{21} & r_{22} & r_{23} & t_y \\ r_{31} & r_{32} & r_{33} & t_z \end{bmatrix} \hat{X} \quad \text{X_world}
 \end{aligned}$$

Next Week

- How to **calibrate** a perspective camera
- How to find the **P** matrix
- How to estimate camera focal length, etc.
- The DLT algorithm

$$\begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \propto \underbrace{\begin{bmatrix} p_{11} & p_{12} & p_{13} & p_{14} \\ p_{21} & p_{22} & p_{23} & p_{24} \\ p_{31} & p_{32} & p_{33} & p_{34} \end{bmatrix}}_{P_{\{3 \times 4\}}} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix}$$