

# Assignment 3

Yifan Luo, u7351505

## Task 1: Model Fitting with the Hough Circle Transformation

### 1. Read RGB images and convert them into greyscale

```
Assignment_3 - task1.ipynb

1 # read RGB/RGBA images with the format (H, W, C)
2 smarties_rgb = cv2.imread('Task1/images/smarties.png') # (356, 413, 3)
3 coins_rgba = cv2.imread('Task1/images/coins.png') # (246, 300, 4)
4
5 # convert to grayscale images
6 smarties = cv2.cvtColor(smarties_rgb, cv2.COLOR_RGB2GRAY) # (356, 413)
7 coins = cv2.cvtColor(coins_rgba, cv2.COLOR_RGB2GRAY) # (246, 300)
8
9 # show grayscale images
10 fig, ax = plt.subplots(1, 2, figsize=(12, 8))
11 ax[0].imshow(smarties, cmap='gray')
12 ax[0].set_title(f"smarties\n{n{smarties.shape}}")
13 ax[0].set_axis_off()
14 ax[1].imshow(coins, cmap='gray')
15 ax[1].set_title(f"coins\n{n{coins.shape}}")
16 ax[1].set_axis_off()
```

Figure 1 Read and convert images

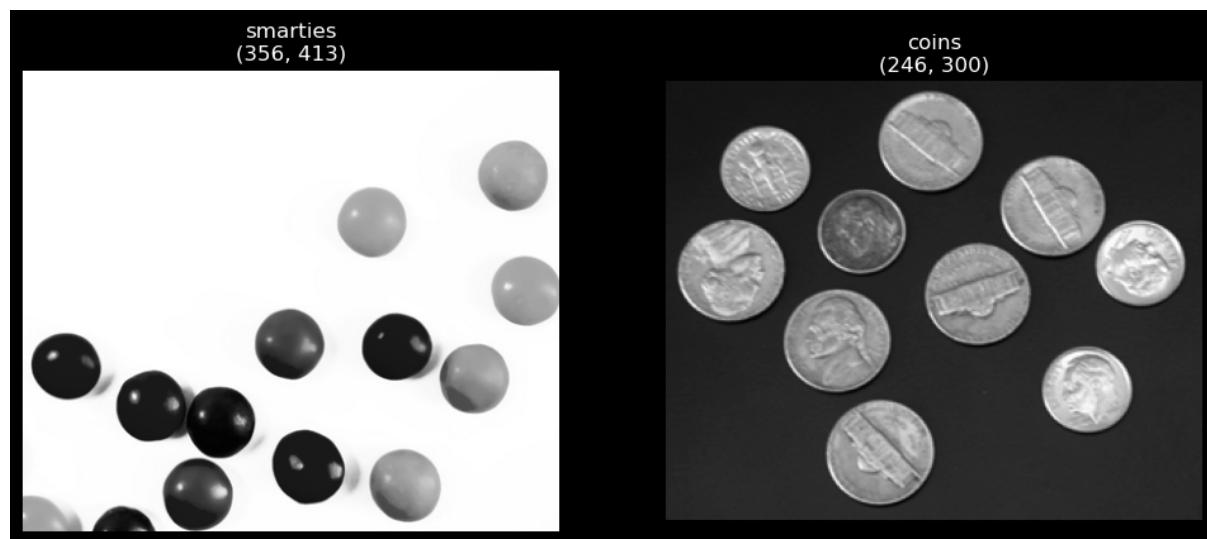


Figure 2 Grayscale images for Task 1

## 2. Apply the Canny edge detector

```
Assignment_3 - task1.ipynb

1 # apply Canny edge detectors on grayscale images
2 dge_smarties = cv2.Canny(smarties, threshold1=230, threshold2=320)
3 edge_coins = cv2.Canny(coins, threshold1=300, threshold2=600)
4
5 fig, ax = plt.subplots(1, 2, figsize=(12, 8))
6 ax[0].imshow(edge_smarties, cmap='gray')
7 ax[0].set_title(f"edges on smarties\n{edge_smarties.shape}")
8 ax[0].set_axis_off()
9 ax[1].imshow(edge_coins, cmap='gray')
10 ax[1].set_title(f"edges on coins\n{edge_coins.shape}")
11 ax[1].set_axis_off()
```

Figure 3 Canny edge detection

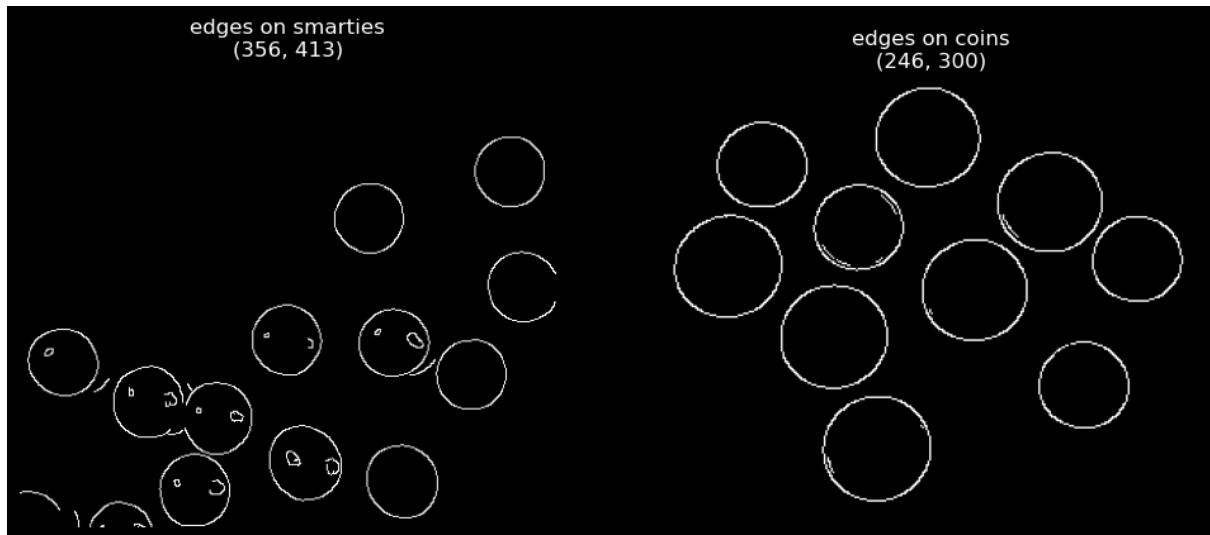


Figure 4 Results from Canny edge detectors

The screenshot above shows the “best” Canny edge detection results after trying multiple threshold combinations.

### 3. Implement the Hough transform function

```
Assignment_3 - task1.ipynb

 1 def hough_circle_transform(
 2     edge,
 3     threshold_ratio=0.8,
 4     radius_increment=1,
 5     circle_centre_increment=[1, 1],
 6     theta_increment=(2 * np.pi / 180),
 7     radius_range=[5, 35],
 8     non_max_suppression=True,
 9     *args,
10     **kwargs
11 ):
12     """Implementation of the Hough transform function with non-maximum suppression.
13
14     Args:
15         edge (np.ndarray): The Canny edge detection result.
16         threshold_ratio (float, optional): Filter small votes. Defaults to 0.8.
17         radius_increment (int, optional): The increment of radius during. Defaults to 1.
18         circle_centre_increment (list, optional): The increment of the circle centre. Defaults to [1, 1].
19         theta_increment (tuple, optional): The increment of theta during. Defaults to (2 * np.pi / 180).
20         radius_range (list, optional): The preset radius range. Defaults to [5, 35].
21         non_max_suppression (bool, optional): Whether to apply non-maximum suppression. Defaults to True.
22
23     Returns:
24         List[int, int, int]: A list of votes [a, b, r] in descending order.
25     """
26
27     # initialise a matrix for recording votes
28     a_range = int(edge.shape[1] + 2 * radius_range[1] // circle_centre_increment[0]) # padding
29     b_range = int(edge.shape[0] + 2 * radius_range[1]// circle_centre_increment[1]) # padding
30     r_range = int((radius_range[1] - radius_range[0]) // radius_increment)
31     vote = np.zeros((a_range, b_range, r_range))
32
33     # iterate over the edge image and accumulate votes
34     for x in range(0, edge.shape[1], circle_centre_increment[0]):
35         for y in range(0, edge.shape[0], circle_centre_increment[1]):
36             # skip if no edge detected
37             if edgely, x] == 0:
38                 continue
39             # iterate over all the possible theta
40             for theta in np.arange(0, 2 * np.pi, theta_increment):
41                 # iterate over all the possible radius within the range
42                 for radius in np.arange(radius_range[0], radius_range[1], radius_increment):
43                     # a = x - r * cos(theta)
44                     candidate_a = x - radius * np.cos(theta)
45                     # b = y - r * sin(theta)
46                     candidate_b = y - radius * np.sin(theta)
47                     # calculate the indices in the vote matrix
48                     candidate_a_idx = int((candidate_a // circle_centre_increment[0]) + radius_range[1])
49                     candidate_b_idx = int((candidate_b // circle_centre_increment[1]) + radius_range[1])
50                     candidate_r_idx = int((radius - radius_range[0]) // radius_increment)
51                     # vote if the index is valid
52                     if (
53                         (0 <= candidate_a_idx < vote.shape[0])
54                         and (0 <= candidate_b_idx < vote.shape[1])
55                         and (0 <= candidate_r_idx < vote.shape[2])
56                     ):
57                         vote[candidate_a_idx, candidate_b_idx, candidate_r_idx] += 1
58
59     print(f"after voting: {np.count_nonzero(vote)}")
60     # set votes that are lower than the threshold to zeros
61     vote = np.where(vote < (threshold_ratio * np.max(vote)), 0, vote)
62     print(f"after implying threshold: {np.count_nonzero(vote)}")
63
64     # implement non-max suppression
65     if non_max_suppression:
66         vote = non_maximum_suppression(vote,
67                                         kwargs['threshold_a'],
68                                         kwargs['threshold_b'],
69                                         kwargs['threshold_r'])
70
71     print(f"after non-max suppression: {np.count_nonzero(vote)}")
72
73     # return a list of votes [a, b, r] in a descending order
74     asc_idx = np.unravel_index(np.argsort(vote, axis=None), vote.shape)
75     res = []
76     for a_idx, b_idx, r_idx in zip(asc_idx[0][::-1][:np.count_nonzero(vote)],
77                                   asc_idx[1][::-1][:np.count_nonzero(vote)],
78                                   asc_idx[2][::-1][:np.count_nonzero(vote)]):
79         res.append([(a_idx - radius_range[1]) * circle_centre_increment[0],
80                     (b_idx - radius_range[1]) * circle_centre_increment[1],
81                     radius_range[0] + r_idx * radius_increment])
82
83     return np.unique(res, axis=0).tolist()
```

Figure 5 Implementation of the Hough transform function

```

Assignment_3 - task1.ipynb

1 def non_maximum_suppression(vote, threshold_a=20, threshold_b=20, threshold_r=20):
2     """Implementation of non-maximum suppression.
3
4     Args:
5         vote (np.ndarray): The vote matrix.
6         threshold_a (int, optional): The half window size in the vertical direction. Defaults to 20.
7         threshold_b (int, optional): The half window size in the horizontal direction. Defaults to 20.
8         threshold_r (int, optional): The half window size for radius. Defaults to 20.
9
10    Returns:
11        _type_: _description_
12    """
13    a_len, b_len, r_len = vote.shape
14    res = np.zeros_like(vote)
15    # zero padding
16    vote_pad = np.zeros(
17        (a_len + 2 * threshold_a, b_len + 2 * threshold_b, r_len + 2 * threshold_r)
18    )
19    vote_pad[threshold_a : threshold_a + a_len,
20             threshold_b : threshold_b + b_len,
21             threshold_r : threshold_r + r_len] = vote
22    # non-max suppression
23    for a in range(threshold_a, threshold_a + a_len):
24        for b in range(threshold_b, threshold_b + b_len):
25            for r in range(threshold_r, threshold_r + r_len):
26                if vote_pad[a, b, r] == 0:
27                    continue
28                # the neighbours around [a, b, r]
29                # controlled by the three thresholds
30                neighbours = vote_pad[
31                    a - threshold_a : a + threshold_a + 1,
32                    b - threshold_b : b + threshold_b + 1,
33                    r - threshold_r : r + threshold_r + 1,
34                ]
35                # suppression if the vote is not the maximum in its neighbourhood
36                if vote_pad[a, b, r] == np.max(neighbours):
37                    res[a - threshold_a, b - threshold_b, r - threshold_r] = vote_pad[a, b, r]
38
return res

```

Figure 6 The non-maximum suppression in Hough transformation

Figure 5 shows the implementation of the Hough transformation based on the required interface. During the voting stage, in the search space, we only record the corresponding index instead of the real value. Notice that each iteration is controlled by a searching range and an increment.

After applying a threshold and non-maximum suppression, a list of circle candidates is filtered. Before returning the result in descending order, we need to transform the index from the voting matrix into the real  $[a, b, r]$  value.

#### 4. Design and implement a non-maximum suppression algorithm

The screenshot is shown in Figure 6 in the last task.

The function takes three thresholds to control the window size of its searching area. Especially, `threshold_a` (units of pixels) controls the neighbourhood in the vertical direction; `threshold_b` (units of pixels) controls the neighbourhood in the horizontal

direction threshold\_r (units of pixels) controls the searching range of radius. Non-maximum vote will be suppressed to 0 after comparing with its neighbourhood. Only local maxima of the designed neighbourhood are kept.

The function is called inside the hough\_circle\_transform function shown in Figure 5.

## 5. Detect circles

```
Assignment_3 - task1.ipynb

1 # detect coins
2 circles_coins = hough_circle_transform(
3     edge_coins,
4     threshold_ratio=0.7,
5     radius_increment=1,
6     circle_centre_increment=[1, 1],
7     theta_increment=(2 * np.pi / 180),
8     radius_range=[5, 35],
9     non_max_suppression=True,
10    threshold_a=20,
11    threshold_b=20,
12    threshold_r=20
13 )
14
15 fig, ax = plt.subplots()
16 ax.imshow(coins_rgb)
17 for x, y, r in circles_coins:
18     circle = plt.Circle((x, y), r, color='red', linewidth=2, fill=False)
19     ax.add_patch(circle)
20     ax.text(x - r + np.random.randint(0, 8),
21             y + np.random.randint(-5, 5),
22             f"({x}, {y}), {r}",
23             color='red', fontsize=10)
24 ax.set_axis_off()
25 ax.set_title("coins")
```

Figure 7 Detect circles on coins

after voting: 1528133  
after implying threshold: 46  
after non-max suppression: 11

Figure 8 Circles during each stage

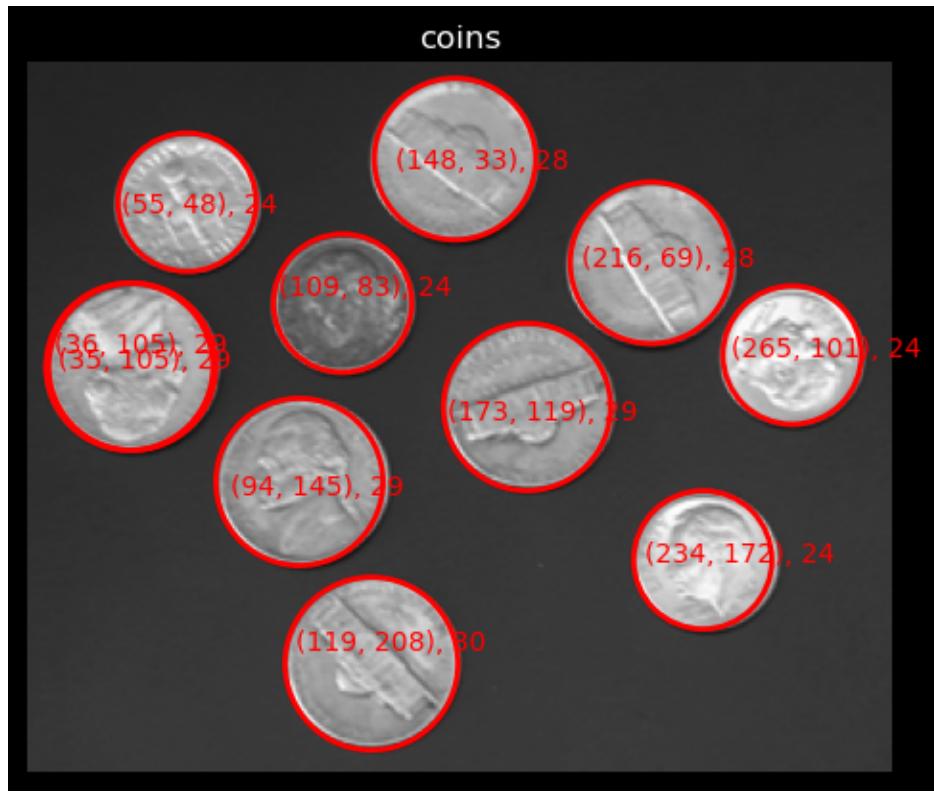


Figure 9 Circles on coins

There are 11 circles detected on the coins. The text over the coins shows the [a, b, r] value for each detection. Two circles are duplicated but are not been suppressed due to they have the same high votes.

```
Assignment_3 - task1.ipynb

1 # detect smarties
2 circles_smarties = hough_circle_transform(
3     edge_smarties,
4     threshold_ratio=0.7,
5     radius_increment=1,
6     circle_centre_increment=[1, 1],
7     theta_increment=(2 * np.pi / 180),
8     radius_range=[10, 35],
9     non_max_suppression=True,
10    threshold_a=20,
11    threshold_b=20,
12    threshold_r=20
13 )
14
15 fig, ax = plt.subplots()
16 ax.imshow(smarties_rgb)
17 for x, y, r in circles_smarties:
18     circle = plt.Circle((x, y), r, color='red', linewidth=2, fill=False)
19     ax.add_patch(circle)
20     ax.text(x - r + np.random.randint(0, 8),
21             y + np.random.randint(-5, 5),
22             f"({x}, {y}), {r}",
23             color='red', fontsize=8)
24 ax.set_title('circles on smarties')
25 ax.set_axis_off()
```

Figure 10 Detect circles on smarties

after voting: 2070295  
after implying threshold: 24  
after non-max suppression: 12

Figure 11 Detecting results during each stage

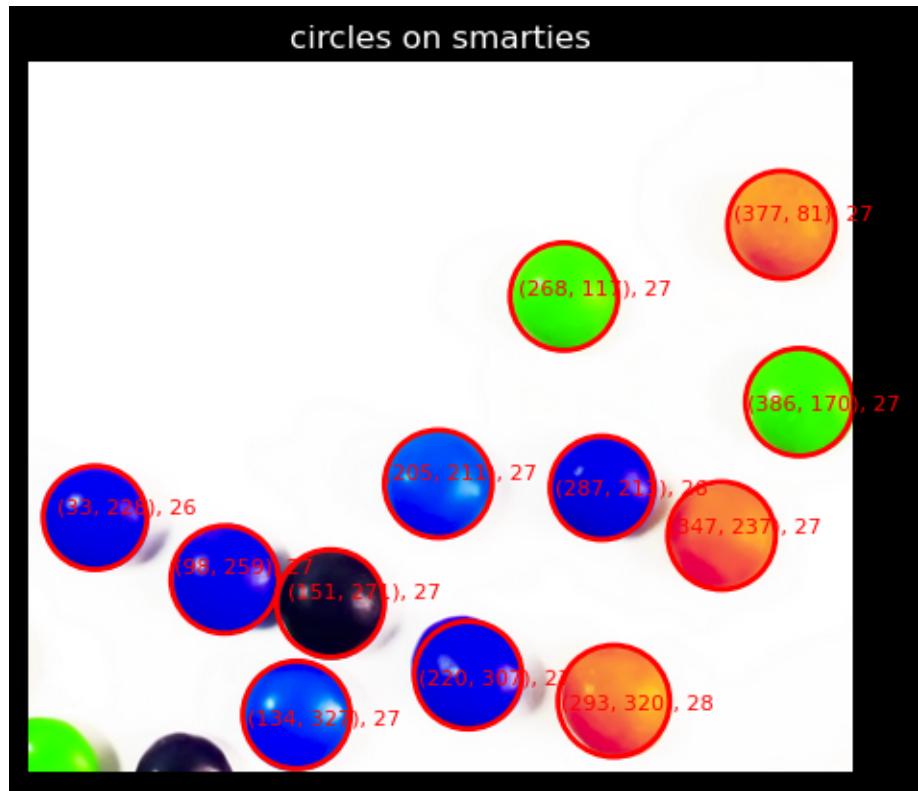


Figure 12 Circles on smarties

Similarly, the above three pictures show the results of smarties. Notice that the bottom-left smarties are not detected due to they have fewer votes and are filtered later by applying the threshold.

## 6. Analysis the results

For coins:

There are 10 coins in the image as the ground truth. There are 11 circles detected at the end. All ground truths are detected, however, there is a duplicated circle overlayed on the same coin. Hence,  $FP = 1$  (the duplicated circle),  $FNR = 0$  (all circles are detected),  $ACC = 10/11$ . The reason for the duplicates is that the votes of both circles are the same (89 and 89 shown in the below figure), and they cannot be filtered by applying the threshold or non-maximum suppression.

[118. 117. 109. 109. 108. 105. 104. 101. 97. 89. 89.]

Figure 13 Top 11 votes in coins

For smarties:

There are 12 full smarties and 2 partial smarties (bottom-left corner) in total as ground truth. The 12 full smarties are all detected, while the 2 partial smarties are not detected. Suppose the 2 partial smarties are also considered as ground truth circles, then  $FP = 0$ .

(all detected circles are true positive), FNR = 2 / 14 (the bottom-left corner circles are not detected). ACC = 1 (all detected circles are TP).

When using a smaller threshold (more detections) threshold\_ratio = 0.1, threshold\_a = threshold\_b = threshold\_r = 5, we can see the bottom-left corner circles can be detected:

```
after voting: 2070295  
after implying threshold: 149014  
after non-max suppression: 779
```

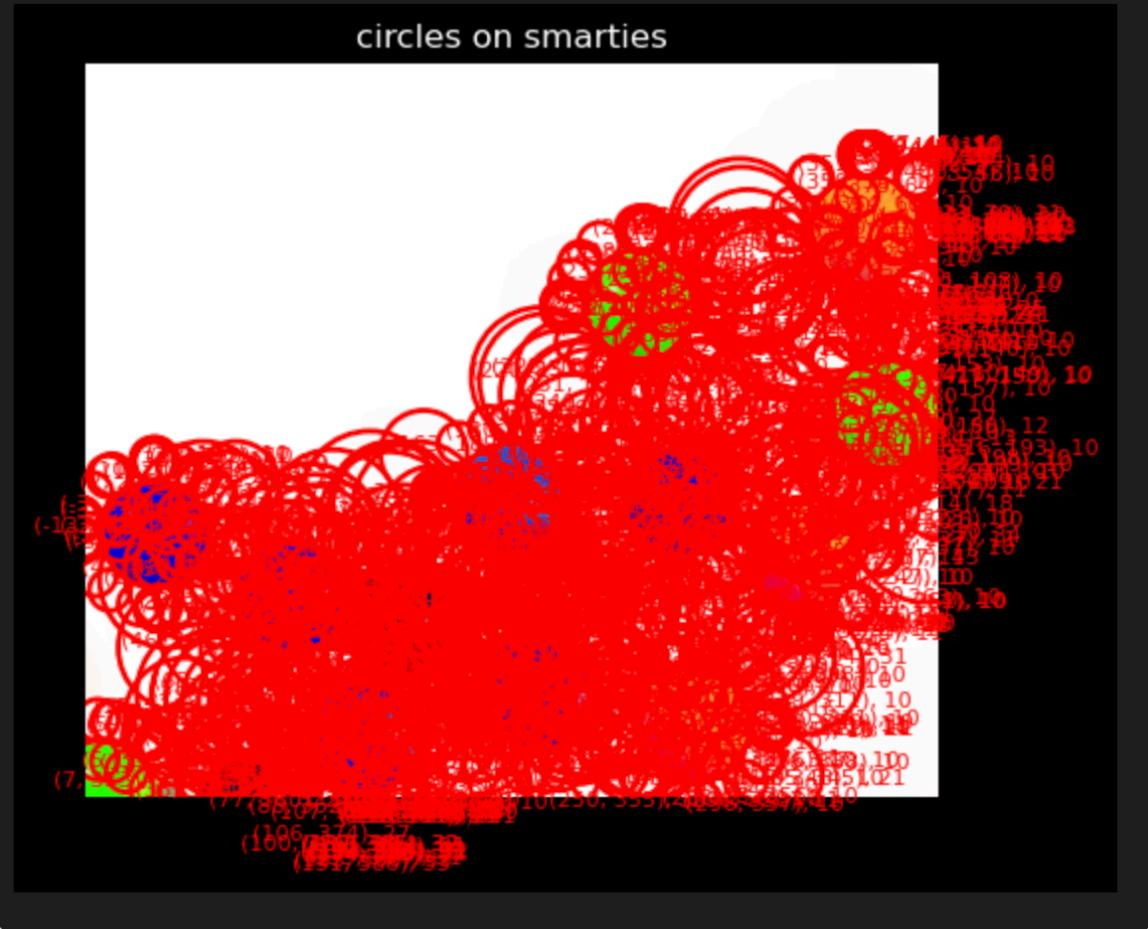


Figure 14 smarties with smaller threshold

After raising threshold\_ratio to 0.3 while keeping all other parameters the same, the result shows that only one bottom-left circle is detected, while many false positive circles around other 12 smarlies:

```
after voting: 2070295  
after implying threshold: 1496  
after non-max suppression: 72  
[131. 127. 121. 118. 118. 107. 105. 104. 101. 99. 95.]
```

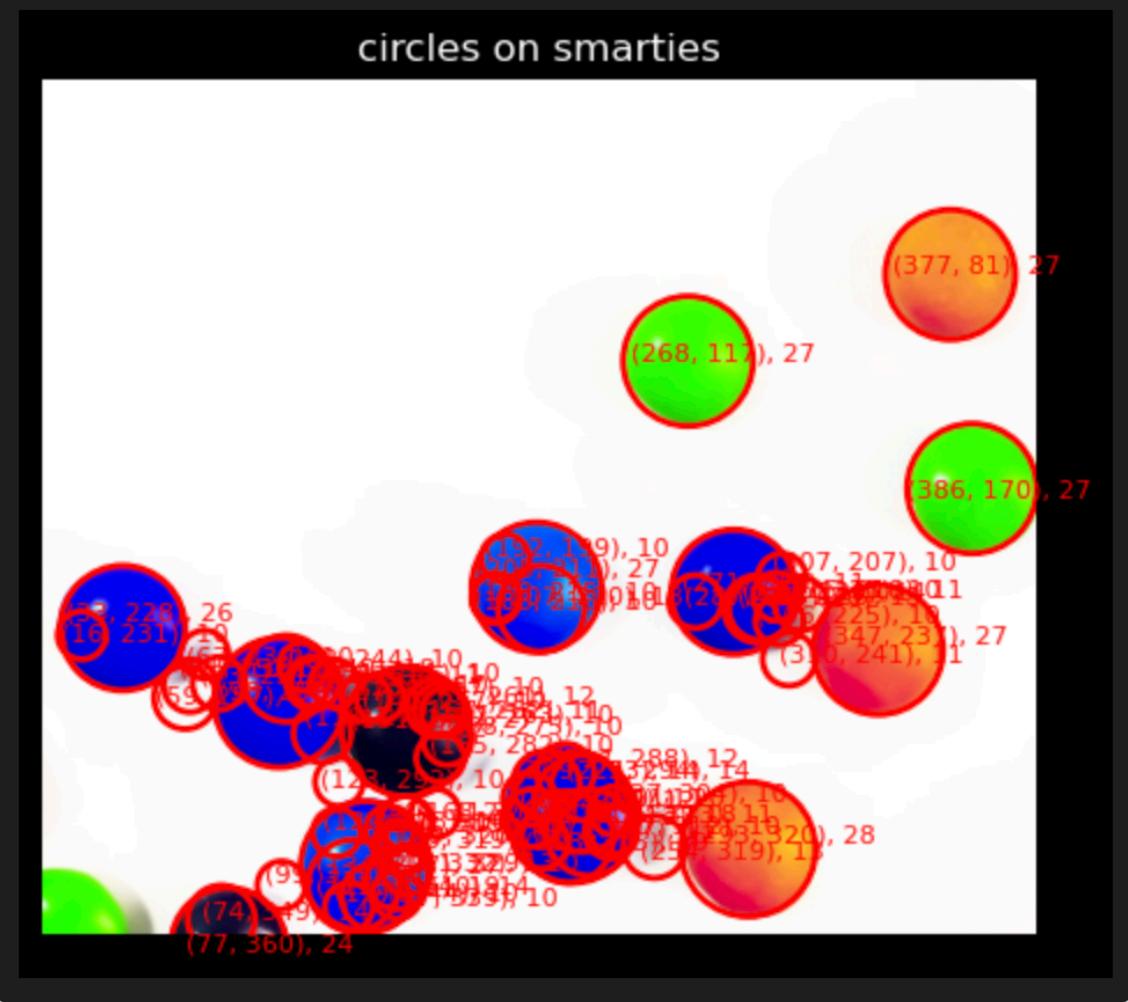


Figure 15 smarties with larger threshold

When keeping the threshold\_ratio as 0.1, and enlarge the threshold\_a = threshold\_b = 30 and threshold\_r = 10, the result is:

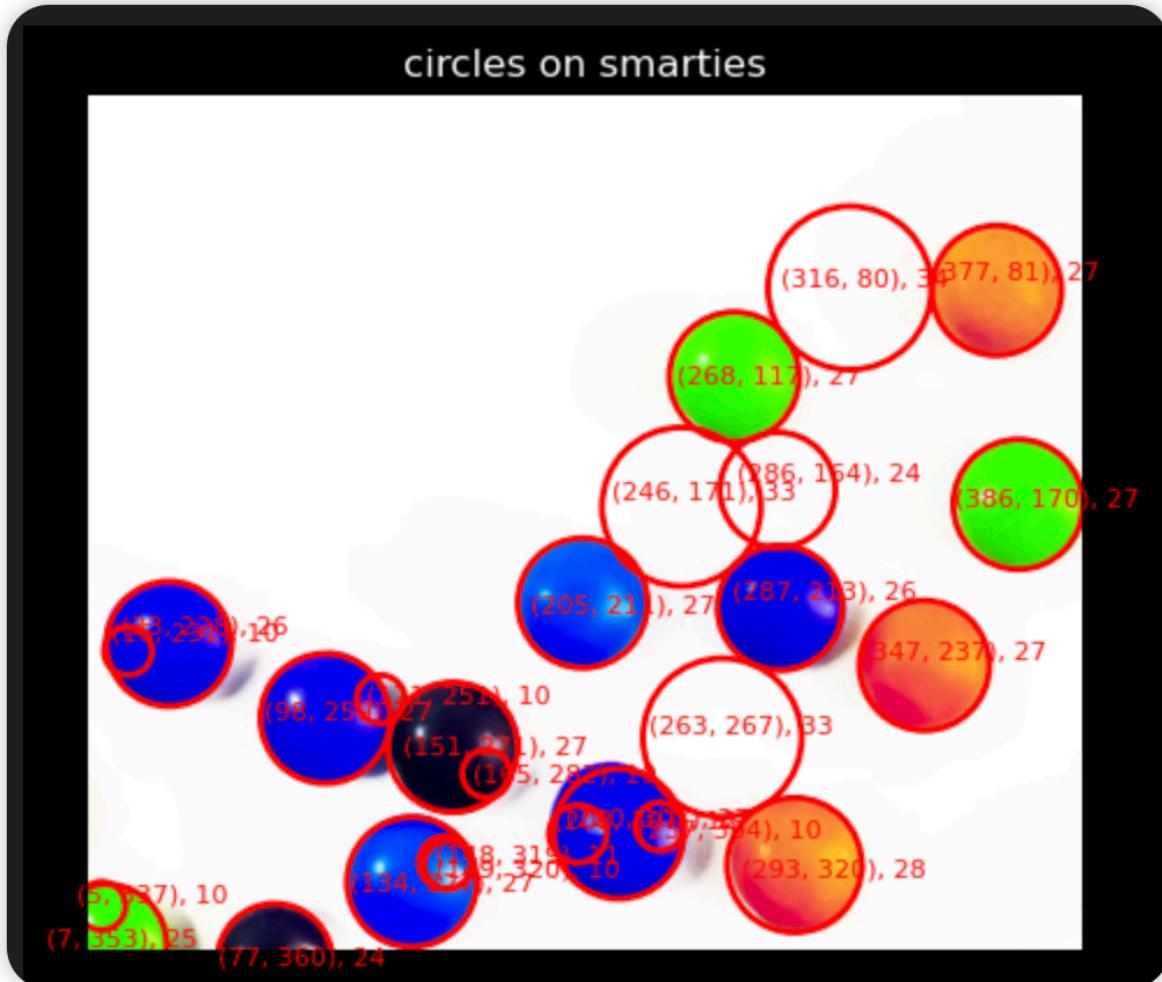


Figure 16 smarties with large thresholds

Fewer FPs are detected on the 12 smarties, while the bottom-left corner circles are still detected.

Keep enlarging the threshold\_r to 50 while keeping others the same, it shows a fairly good result:

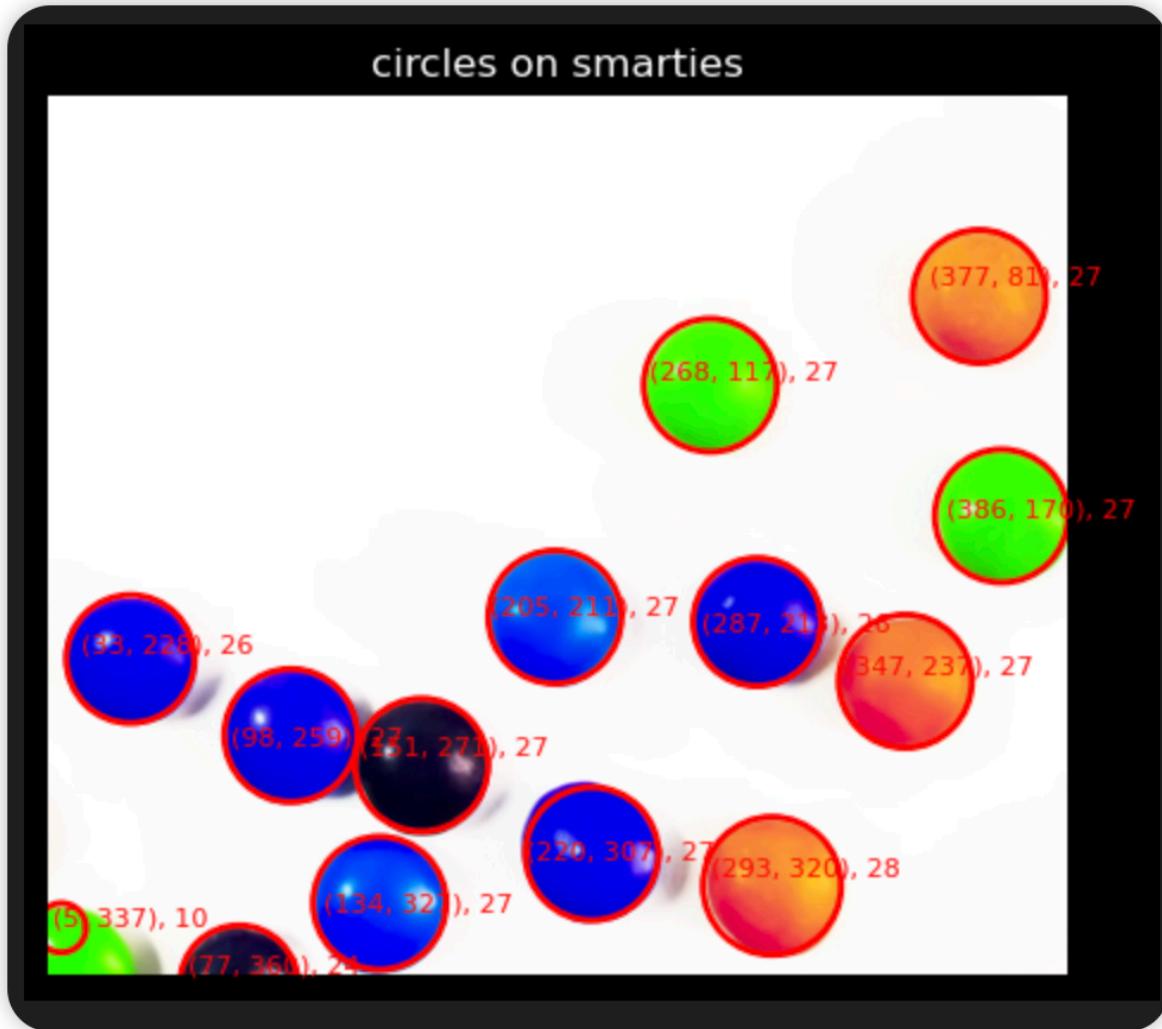


Figure 17 smarties with large thresholds

We can conclude that the bottom-left circles are been filtered by comparing them with the small FP circles overlayed beside them during non-maximum suppression.

From the above analysis and vacuolations, we see that finding a proper set of thresholds is sometimes time-consuming and needs to be tried multiple times to get fairly decent results. The more ground truth the picture has, the more combinations of thresholds we need to try. On the other hand, the complex the pattern is, the more combinations of thresholds we need to try.

## 7. Examine the impacts of different parameters

Part of the analysis is included in the last question. In this question, we focus on other parameters inside the Hough transformation.

For coins, when setting a large threshold\_ratio (0.9) without non-maximum suppression, while setting other increments as 1 unit of pixel and radius range [5, 35], fewer circles are detected:



Figure 18 coins with threshold\_ratio 0.9

After changing threshold\_ratio to 0.4 while keeping others the same, more circles are detected:

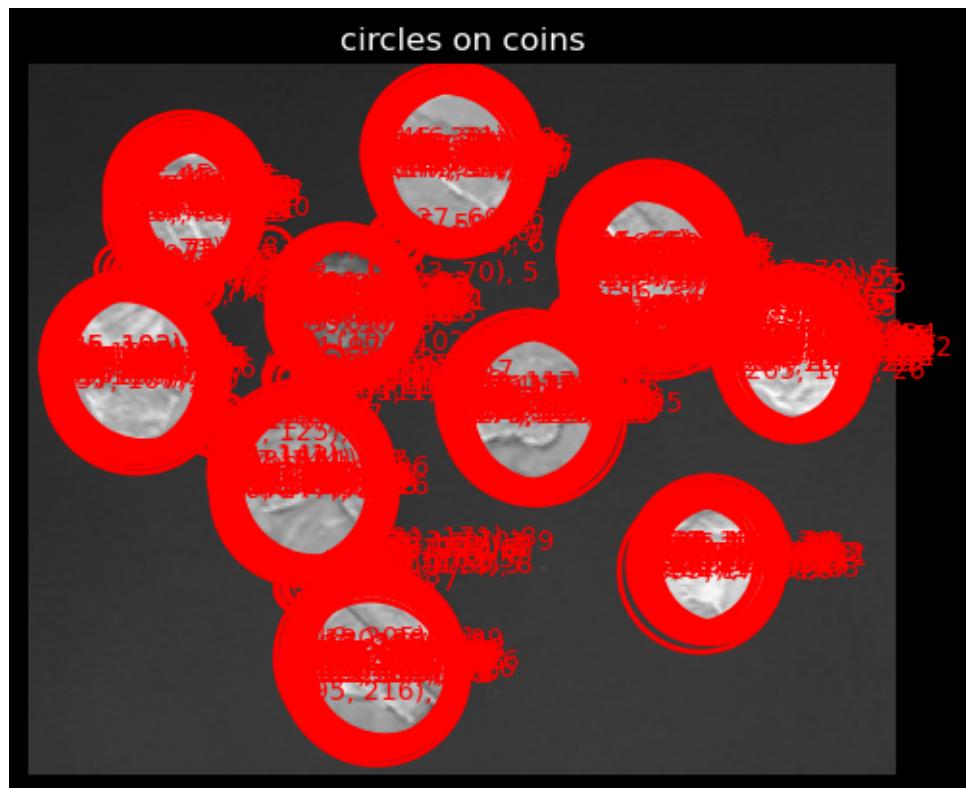


Figure 19 coins with threshold 0.4

We can easily conclude that, when threshold\_ratio is set to a large number, only a few circles are left after applying the threshold. Some TP may be left because of improper threshold\_ratio. Hence it is better to set it to a moderately small/large number to ensure it can select all TPs.

Now set threshold\_ratio to 0.7, and all increments to 1:

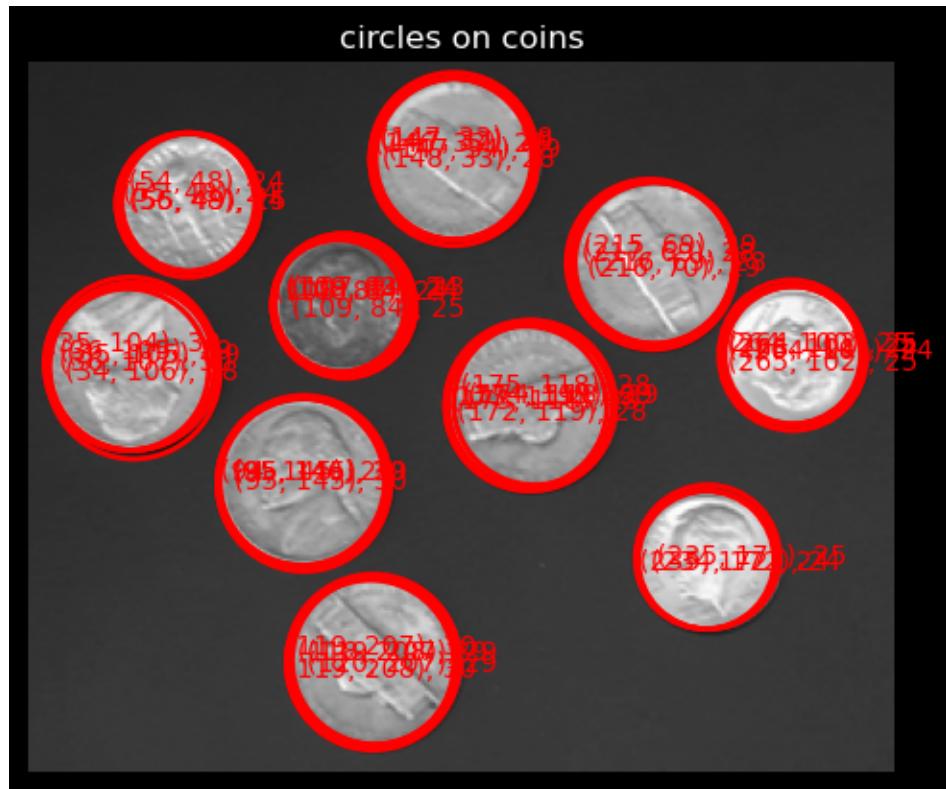


Figure 20 Coins with radius increment 1

When enlarging the radius increment to 3 while keeping others the same:



Figure 21 Coins with radius increment 3

Set radius increment to 10:



Figure 22 Coins with radius increment 10

We can conclude that the radius increment has a strong impact on the results. Different values may generate variant results (i.e., high variance). Hence, it is recommended to set it to a small value to detect all the possible TPs, while keep increasing it for a more efficient computation time.

Similarly, set the circle centre increment to [1, 1]:

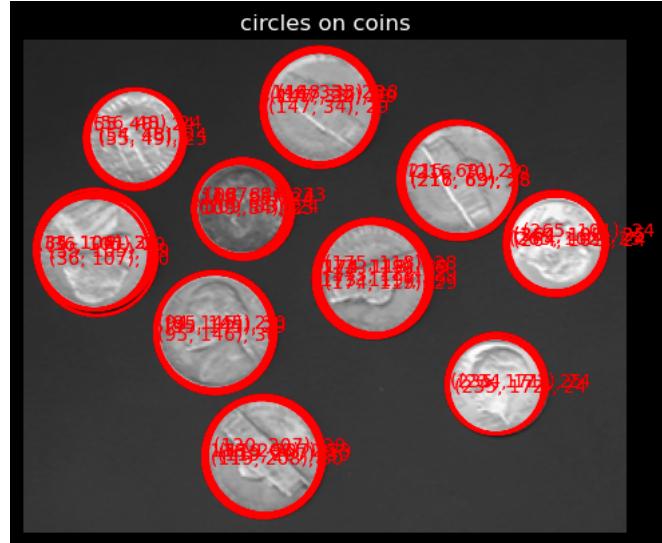


Figure 23 Coins with circle centre increment [1, 1]

Set it to [3, 3]:



Figure 24 Coins with centre increment [3, 3]

Set it to [6, 6]:

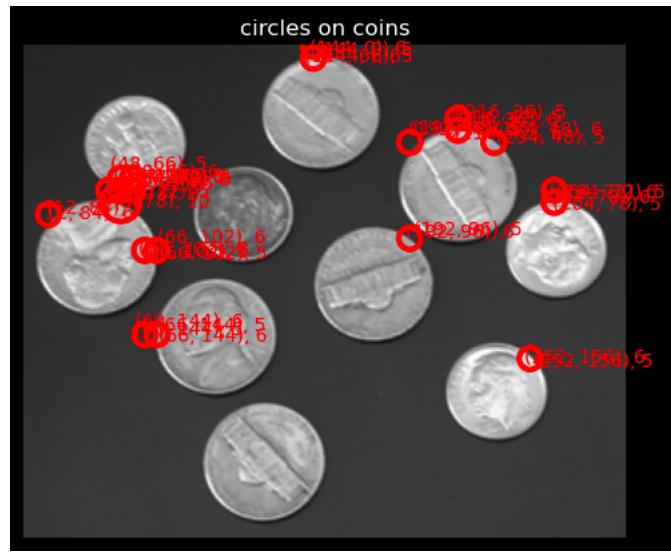


Figure 25 Coins with circle centre increment [6, 6]

We see that fewer circles are detected since the parameter space is getting smaller. It is recommended to set the value to a properly smaller value and increase it gradually to save more computation time.

When setting theta increment to  $5*(2*pi/180)$ :

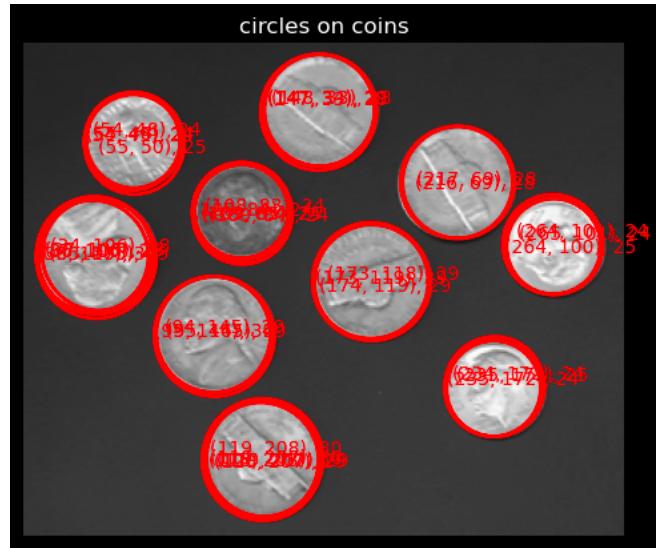
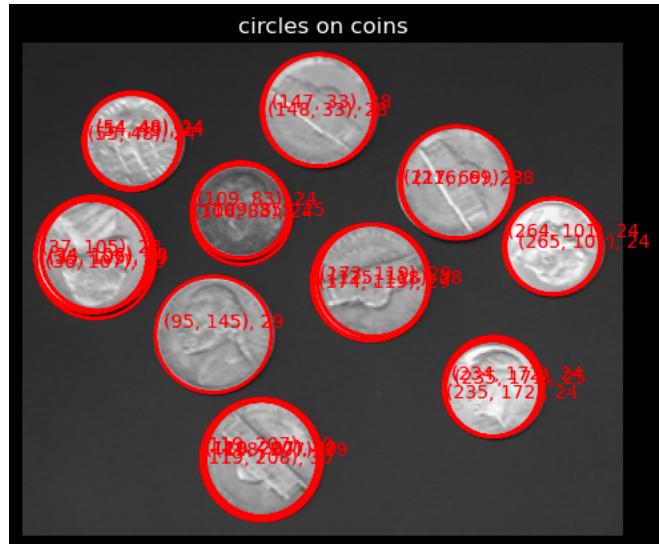


Figure 26 Coins with theta increment to  $5*(2*pi/180)$

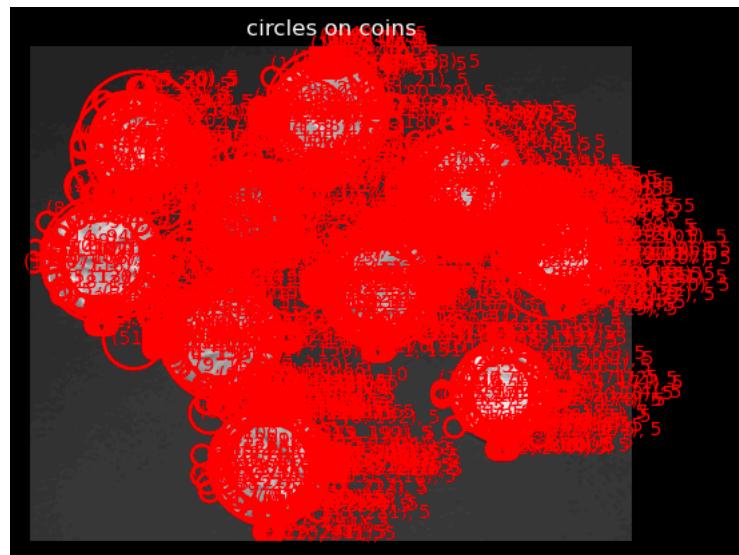
Set it to  $10*(2*pi/180)$ :



*Figure 27 Coins with theta increment to  $10^*(2*pi/180)$*

We can see fewer circles are detected since the parameter space is getting smaller. . It is recommended to set the value to a properly smaller value and increase it gradually to save more computation time.

Now set threshold\_ratio to 0.3, radius\_increment to 5, other increments to 1, and threshold\_a = threshold\_b = threshold\_r = 1:



*Figure 28 Coins with threshold 1*

Set threshold\_a = threshold\_b = threshold\_r = 5:

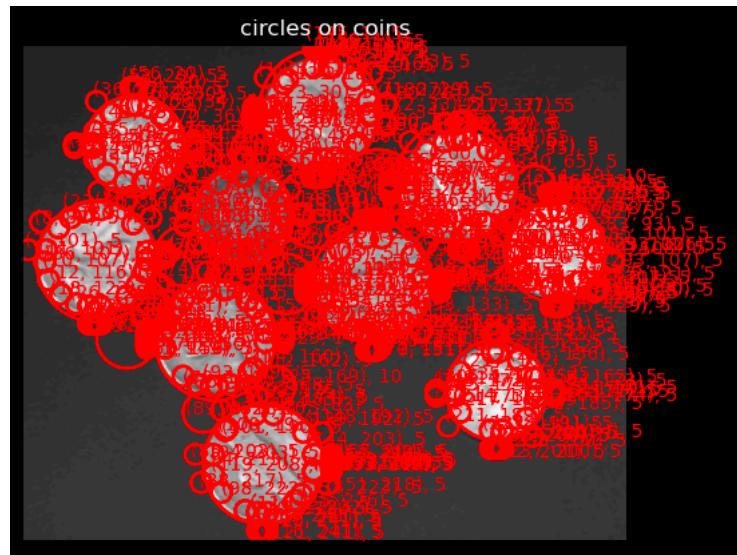


Figure 29 Coins with threshold 5

Set threshold\_a = threshold\_b = threshold\_r = 20:

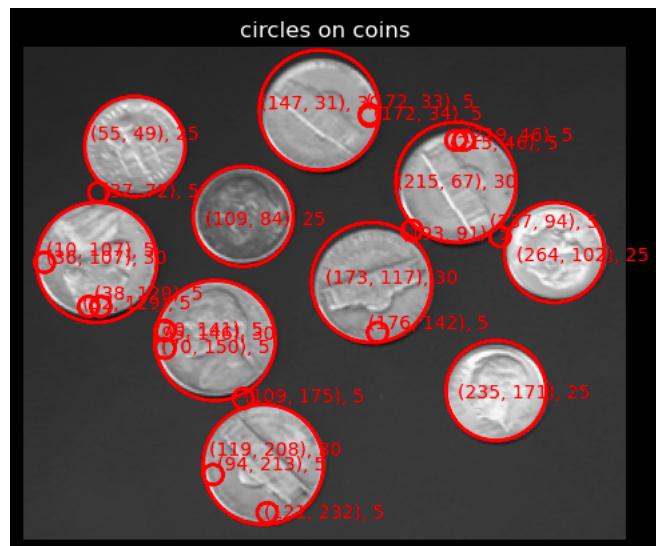


Figure 30 Coins with threshold 20

We can conclude that the larger the threshold, the larger the neighbourhood will be for each non-maximum suppression. It's more recommended to use a large threshold after applying threshold and use non-maximum suppression to filter noises.

## 8. Report the time and memory usage

For coins, when using threshold\_ratio 0.9, radius\_increment 1, circle\_centre increment [1, 1], theta\_increment 2\*pi/180, radius\_range [5, 35], threshold\_a=threshold\_b=threshold\_r 20, the average running time in Jupyter Notebook is 28.32s.

For non-maximum suppression, the time complexity is  $O(a*b*r)$  where  $a$ ,  $b$ ,  $r$  are the shape of the vote matrix, i.e., the parameter space. The space complexity is also  $O(a*b*r)$ .

For the Hough circle transformation function, the time complexity is  $O(H*W*\theta*\text{radius})$ , where  $H$  and  $W$  represent the height and width of the image, and  $\theta$  and  $\text{radius}$  represent each search space. The space complexity is  $O(a*b*r)$ , where  $a$ ,  $b$ ,  $r$  represent each search space.

To optimise the algorithm:

- Instead of voting for every possible parameter combination, employ strategies like edge gradient-based voting or adaptive voting thresholds to prioritise more promising candidate circles.
- Calculation of different combinations of parameters can be accelerated by parallelisation.
- Using sparse sampling of accumulator space can also speed up the voting process while maintaining robustness.
- Using sparse matrices or compressed representations to store accumulator values and reduce memory usage.
- Techniques like thresholding accumulator values or early stopping criteria based on accumulated evidence can terminate voting for non-promoting circles early.

## Task2: Two-view Homography Estimation

### 1. Read RGB images as NumPy arrays



Figure 31 RGB images in Task 2

## 2. Implement the homography function

```
Assignment_3 - task2.ipynb

1 def homography(u1, v1, u2, v2):
2     """Implementation of the DLT algorithm for normalised correspondences.
3
4     Args:
5         u1 (np.ndarray or list): The normalised row-wise coordinates of the first image.
6         v1 (np.ndarray or list): The normalised column-wise coordinates of the first image.
7         u2 (np.ndarray or list): The normalised row-wise coordinates of the second image.
8         v2 (np.ndarray or list): The normalised column-wise coordinates of the second image.
9
10    Returns:
11        H (np.ndarray): The 3*3 homography matrix that wraps (u1, v1) into (u2, v2).
12    """
13    # input correspondences should have the same length
14    if not len(u1) == len(v1) == len(u2) == len(v2):
15        raise ValueError("All input arrays should have the same length")
16    # homography requires at least 4 correspondences
17    if len(u1) < 4:
18        raise ValueError("At least 4 corresponding points are required")
19    # the number of total correspondences
20    n_correspondences = len(u1)
21    # initialise the A matrix where each correspondence contributes two independent rows
22    A = np.zeros((2 * n_correspondences, 9))
23    # construct the A matrix for each correspondence
24    for i in range(n_correspondences):
25        A[2 * i] = [0, 0, 0, -u1[i], -v1[i], -1, v2[i] * u1[i], v2[i] * v1[i], v2[i]]
26        A[2 * i + 1] = [u1[i], v1[i], 1, 0, 0, 0, -u2[i] * u1[i], -u2[i] * v1[i], -u2[i]]
27    # use SVD to decompose the A matrix
28    U, S, V_t = np.linalg.svd(A)
29    # take the rightmost column of V as the length-9 vector h
30    h = V_t.T[:, -1]
31    # fix the norm of h to 1
32    h /= np.linalg.norm(h)
33    # reshape (9,) vector h to a (3, 3) matrix H
34    H = h.reshape(3, 3)
35    # scale the bottom-right element to 1
36    H /= H[-1, -1]
37    return H
```

Figure 32 The homography function based on the described interface

The function takes normalised correspondences from both images and conducts an SVD decomposition on the constructed matrix A from all correspondences.

The returned homography matrix H is constructed from the last column of the decomposed matrix V, scaled its norm to 1. Notice that, according to conventions, H is divided by the bottom-right corner element to ensure its value equals 1.

### 3. Implement the compute\_normalisation\_matrix function and the homography\_w\_normalisation function

```
Assignment_3 - task2.ipynb

1 def compute_normalisation_matrices(img1, img2):
2     """Generate normalisation matrices for input images.
3
4     Args:
5         img1 (np.ndarray): The image to calculate the normalisation matrix T.
6         img2 (np.ndarray): The image to calculate the normalisation matrix T'.
7
8     Returns:
9         T, T' (tuple(np.ndarray, np.ndarray)): The normalisation matrices for img1 and img2.
10    """
11    # the height and width for both images
12    w1, h1 = img1.shape[1], img1.shape[0]
13    w2, h2 = img2.shape[1], img2.shape[0]
14    # calculate the normalisation matrix T and T'
15    T1 = np.linalg.inv(np.array([[w1 + h1, 0, w1 / 2],
16                                [0, w1 + h1, h1 / 2],
17                                [0, 0, 1]]))
18    T2 = np.linalg.inv(np.array([[w2 + h2, 0, w2 / 2],
19                                [0, w2 + h2, h2 / 2],
20                                [0, 0, 1]]))
21    return T1, T2
```

Figure 33 The compute\_normalisation\_matrix for computing normalisation matrices from input pictures

The screenshot above shows the implementation of calculating normalisation matrices T and T' based on slides. The function returns calculated T and T' for later use.

```

Assignment_3 - task2.ipynb

1 def homography_w_normalisation(img1, u1, v1, img2, u2, v2):
2     """Implementation of the DLT algorithm for normalised correspondences.
3
4     Args:
5         img1 (np.ndarray): The first image.
6         u1 (np.ndarray or list): The unnormalised horizontal coordinates of the first image.
7         v1 (np.ndarray or list): The unnormalised vertical coordinates of the first image.
8         img2 (np.ndarray): The second image.
9         u2 (np.ndarray or list): The unnormalised horizontal coordinates of the second image.
10        v2 (np.ndarray or list): The unnormalised vertical coordinates of the second image.
11
12    Returns:
13        H (np.ndarray): The denormalised homography matrix.
14    """
15    # input correspondences should have the same length
16    if not len(u1) == len(v1) == len(u2) == len(v2):
17        raise ValueError("All input arrays should have the same length")
18    # homography requires at least 4 correspondences
19    if len(u1) < 4:
20        raise ValueError("At least 4 corresponding points are required")
21
22    # compute the normalisation matrices
23    T1, T2 = compute_normalisation_matrices(img1, img2)
24    # construct homography indices
25    X1 = np.vstack((u1, v1, np.ones(len(u1))))
26    X2 = np.vstack((u2, v2, np.ones(len(u2))))
27    # normalise the homography indices with calculated normalisation matrices
28    X1_norm = T1 @ X1
29    X2_norm = T2 @ X2
30    # transform from homography coordinates to Cartesian coordinates
31    u1_norm, v1_norm = X1_norm[0] / X1_norm[-1], X1_norm[1] / X1_norm[-1]
32    u2_norm, v2_norm = X2_norm[0] / X2_norm[-1], X2_norm[1] / X2_norm[-1]
33    # apply DLT algorithm to find the normalised homography
34    H_norm = homography(u1_norm, v1_norm, u2_norm, v2_norm)
35    # denormalise the homography matrix H_norm
36    H = np.linalg.inv(T2) @ H_norm @ T1
37    # set H_33 to 1 for normalisation purpose
38    H /= H[-1, -1]
39
40    return H

```

*Figure 34 The homography\_w\_normalisation function to calculate unnormalised homography matrix H*

The function above first calculates the normalisation matrices T1 and T2 by using the previous function compute\_normalisation\_matrix. Then, the unnormalised coordinates u1, v1, u2 and v2 are constructed and transformed into normalised coordinates using T1 and T2.

The homography function is used to calculate homography matrix with normalised coordinates u1\_norm, v1\_norm, u2\_norm and v2\_norm. Then, the homography matrix is denormalised by using T1 and T2 and scaled by using the bottom-right corner entry.

## 4. Use SIFT keypoint detectors and feature descriptors to find correspondences

```
Assignment_3 - task2.ipynb

1 # initialise the SIFT feature detector
2 sift = cv2.SIFT_create()
3
4 mountain1_bgr = cv2.imread('Task2/images/mountain1.jpg')
5 mountain2_bgr = cv2.imread('Task2/images/mountain2.jpg')
6
7 # detect keypoints and compute descriptors
8 kp1, des1 = sift.detectAndCompute(mountain1_gray, None)
9 kp2, des2 = sift.detectAndCompute(mountain2_gray, None)
10
11 # draw the keypoints
12 mountain1_sift = cv2.drawKeypoints(mountain1_gray, kp1, mountain1_bgr, flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
13 mountain2_sift = cv2.drawKeypoints(mountain2_gray, kp2, mountain2_bgr, flags=cv2.DRAW_MATCHES_FLAGS_DRAW_RICH_KEYPOINTS)
14 fig, ax = plt.subplots(1, 2, figsize=(10, 8))
15 ax[0].imshow(mountain1_sift)
16 ax[0].set_axis_off()
17 ax[0].set_title(f"mountain1.jpg\n{n(len(kp1)} keypoints")
18 ax[1].imshow(mountain2_sift)
19 ax[1].set_axis_off()
20 ax[1].set_title(f"mountain2.jpg\n{n(len(kp2)} keypoints")
21 plt.show()
```

Figure 35 Use SIFT to find keypoints and descriptors

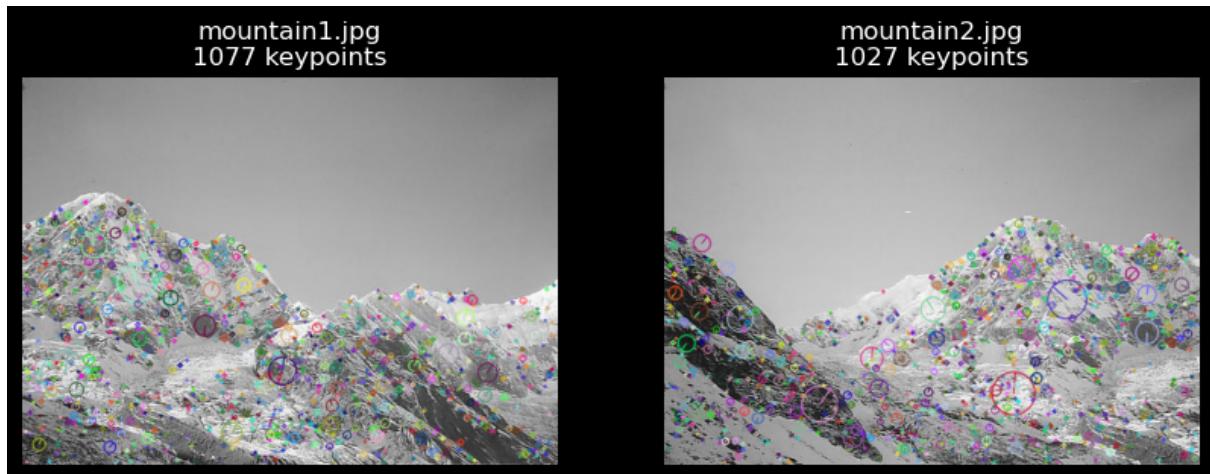


Figure 36 The detected keypoints from both images

First, read the original images in BGR and grayscale formats. Then, use `cv2.SIFT_create()` to create SIFT detectors and find keypoints (`kp1`, `kp2`) and descriptors (`des1`, `des2`) for `mountain1` and `mountain2`. The results show `mountain1` has 1077 keypoints and `mountain2` has 1027 keypoints.

```

Assignment_3 - task2.ipynb

1 # reference: https://docs.opencv.org/4.x/d5/d6f/tutorial_feature_flann_matcher.html
2
3 # find matches by using FLANN
4 matcher = cv2.DescriptorMatcher_create(cv2.DescriptorMatcher_FLANNBASED)
5 knn_matches = matcher.knnMatch(des1, des2, 2)
6
7 # filter matches using the Lowe's ratio test
8 # the lower the ratio, the fewer the correspondences
9 ratio = 0.5
10 all_matches = []
11 for m, n in knn_matches:
12     if m.distance < ratio * n.distance:
13         all_matches.append(m)

```

*Figure 37 Match and select correspondences with a threshold*

The figure above shows how to use a matcher to match keypoints in OpenCV. Notice that we use a ratio of 0.5 as a threshold to select from all the correspondences. The lower the ratio, the fewer the matches.

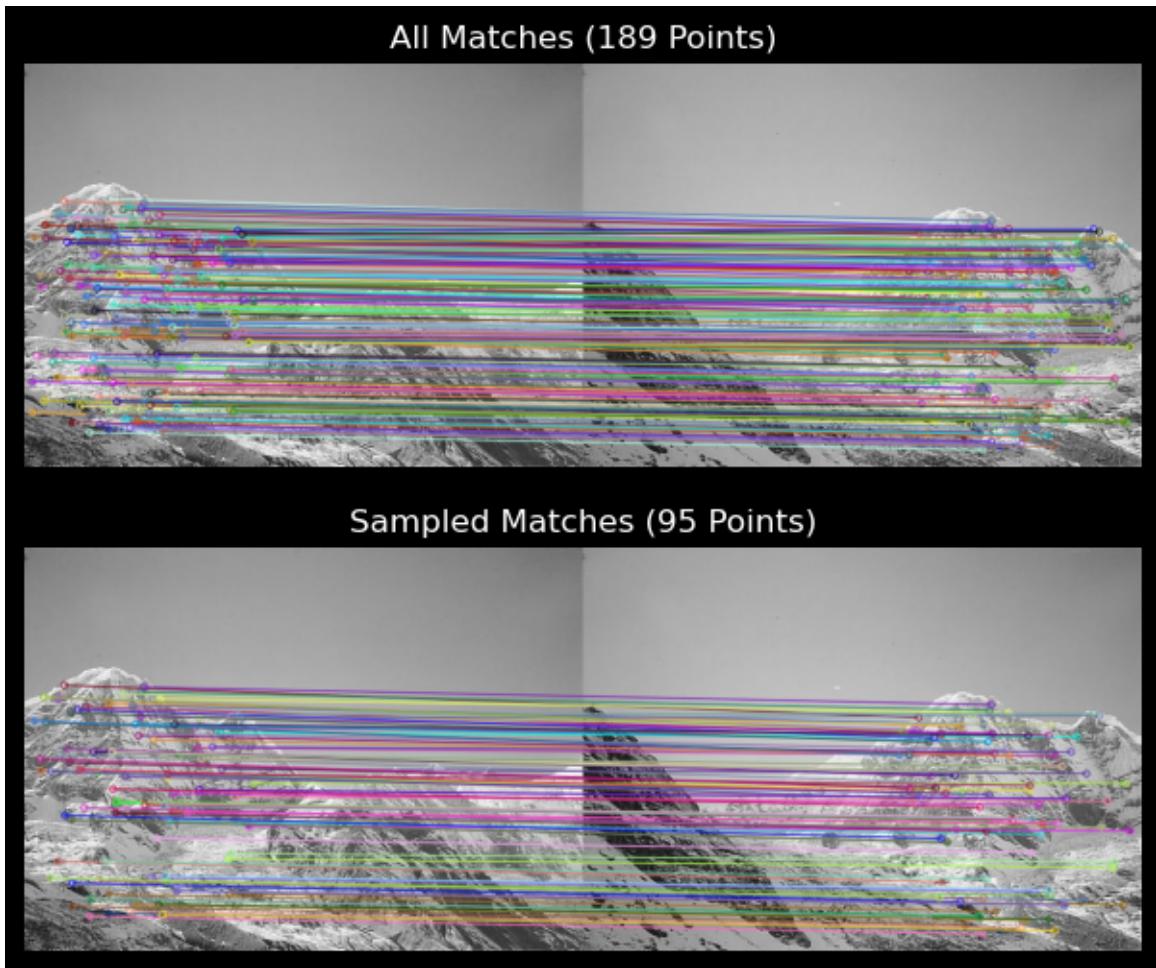
```

Assignment_3 - task2.ipynb

1 # draw all matches
2 img_all_matches = np.empty((max(mountain1_gray.shape[0], mountain2_gray.shape[0]),
3                             mountain1_gray.shape[1] + mountain2_gray.shape[1], 3),
4                             dtype=np.uint8)
5 cv2.drawMatches(mountain1_gray, kp1, mountain2_gray, kp2, all_matches, img_all_matches, flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
6
7 # sample half of the correspondences
8 np.random.seed(42)
9 N = int(np.ceil(len(all_matches) / 2))
10 sample_idx = np.random.choice(len(all_matches), N, replace=False)
11 sample_matches = np.asarray(all_matches)[sample_idx]
12
13 # draw sampled matches
14 img_matches = np.empty((max(mountain1_gray.shape[0], mountain2_gray.shape[0]),
15                         mountain1_gray.shape[1] + mountain2_gray.shape[1], 3),
16                         dtype=np.uint8)
17 cv2.drawMatches(mountain1_gray, kp1, mountain2_gray, kp2, sample_matches, img_matches, flags=cv2.DrawMatchesFlags_NOT_DRAW_SINGLE_POINTS)
18
19 fig, ax = plt.subplots(2, 1, figsize=(8, 6))
20 ax[0].imshow(img_all_matches)
21 ax[0].set_title(f"All Matches ({len(all_matches)} Points)")
22 ax[0].set_axis_off()
23 ax[1].imshow(img_matches)
24 ax[1].set_title(f"Sampled Matches ({len(sample_matches)} Points)")
25 ax[1].set_axis_off()

```

*Figure 38 Draw all the matches and the selected matches by using a seed 42*



*Figure 39 The all matches and the selected matches*

There are 189 matches under the ratio of 0.5, and 95 selected matches according to the requirement.

```

Assignment_3 - task2.ipynb

1 # convert cv2.DMatch to np.array
2 samples = np.array([[s.queryIdx, s.trainIdx, s.distance] for s in sample_matches])
3 # save sampled matches as a .npy file
4 samples_path = 'Task2/sample_matches.npy'
5 np.save(samples_path, samples)
6 print(f"Saved into {samples_path}")
7 # load samples from .npy file
8 print(f"Loading from {samples_path}")
9 samples_load = np.load("Task2/sample_matches.npy")
10 print(samples_load[:5])

```

*Figure 40 Save and load from a "npy" file*

```
Saved into Task2/sample_matches.npy
Loading from Task2/sample_matches.npy
[[ 539.        1016.        140.10353088]
 [ 477.        915.        102.95144653]
 [ 94.         554.        149.14086914]
 [ 84.         599.        99.85489655]
 [ 228.        647.        66.61831665]]
```

*Figure 41 The NumPy array loaded from the "npy" file*

A “.npy” file is saved for the NumPy array object which contains the coordinates of matches from OpenCV. The above picture shows the array information of the first 5 matches. The first column is the query index for mountain1, the second is the target index for mountain2, and the last is the distance of each match.

## 5. Implement homography\_w\_normalisation\_ransac

```
Assignment_3 - task2.ipynb

1 def homography_w_normalisation_ransac(img1, kp1, img2, kp2, corr, n_matches,
2                                         dist_ratio, max_iter=10, seed=42):
3     """ Implementation of using RANSAC to select the best performing H
4     that has the most amount of inliners
5
6     Args:
7         img1 (np.ndarray): The input query image.
8         kp1 (np.ndarray): The keypoints derived from OpenCV in img1.
9         img2 (np.ndarray): The input target image.
10        kp2 (np.ndarray): The keypoints derived from OpenCV in img2.
11        corr (np.ndarray): The matches from img1 and img2, aligned with kp1 and kp2.
12        n_matches (int): The number of matches to sample.
13        dist_ratio (float): The ratio of distance for selecting inliners.
14        max_iter (int): The max iteration in RANSAC.
15        seed (int): The NumPy random seed for sampling in RANSAC.
16
17     Returns:
18         best_H (np.ndarray): The best performing H that has the most amount of inliners.
19     """
20     # set the seed for sampling in RANSAC
21     np.random.seed(seed)
22     # record the maximum number of inliners during RANSAC
23     max_inliers = 0
24     # record the best performed H that has the most amount of inliners
25     best_H = np.zeros((3, 3))
26     for _ in range(max_iter):
27         # sample n matches to estimate H
28         sample_idx = np.random.choice(len(corr), n_matches, replace=False)
29         sample_matches = np.asarray(corr)[sample_idx]
30         # construct the coordinates u1, v1, u2, v2 from matches
31         u1, v1 = np.array([]), np.array([])
32         u2, v2 = np.array([]), np.array([])
33         for m in sample_matches:
34             v1_, u1_ = kp1[m.queryIdx].pt
35             v2_, u2_ = kp2[m.trainIdx].pt
36             u1 = np.append(u1, u1_)
37             v1 = np.append(v1, v1_)
38             u2 = np.append(u2, u2_)
39             v2 = np.append(v2, v2_)
40         # calculate the homography matrix for unnormalised coordinates
41         H = homography_w_normalisation(img1, u1, v1, img2, u2, v2)
42         # calculate the symmetric transfer error for each correspondence
43         X_1 = np.hstack((u1[:, np.newaxis], v1[:, np.newaxis], np.ones((len(u1), 1))))
44         X_2 = np.hstack((u2[:, np.newaxis], v2[:, np.newaxis], np.ones((len(u2), 1))))
45         # calculate the symmetric transfer error for each match
46         sym_trans_err = np.square(X_1.T - np.linalg.inv(H) @ X_2.T).T.sum(axis=1) + \
47                         np.square(X_2.T - H @ X_1.T).T.sum(axis=1)
48         # calculate the number of inliers that are within dist_ratio * max(dist)
49         # record the best H that has the largest number of inliers
50         n_inliers = len(sym_trans_err[sym_trans_err < dist_ratio * np.max(sym_trans_err)])
51         if n_inliers > max_inliers:
52             max_inliers = n_inliers
53             best_H = H
54
55     return best_H
```

Figure 42 The homography\_w\_normalisation\_ransac function to select the best performing H using RANSAC

The screenshot above shows the implementation of using RANSAC to select the best H. In each round, the algorithm samples several matches and uses these matches to calculate a homography matrix H. We use symmetric transfer error and a distance ratio for each match to select inliers. Record and return the best-performing H at the end.

```
Assignment_3 - task2.ipynb

1 n_matches = 15
2 dist_ratio = 0.08
3 max_iter = 30
4 best_H = homography_w_normalisation_ransac(mountain1_gray, kp1, mountain2_gray, kp2, sample_matches,
5                                              n_matches=n_matches, dist_ratio=dist_ratio, max_iter=max_iter)
6 print(f"Each H is calculated from {n_matches} matches\n")
7 Inliers are selected under the distance ratio {dist_ratio}\n
8 In the total {max_iter} iterations"
9 print(f"The best H found by RANSAC:\n{best_H}")
```

Figure 43 Show the best performing H

```
Each H is calculated from 15 matches
Inliers are selected under the distance ratio 0.08
In the total 30 iterations
The best H found by RANSAC:
[[ 8.50960152e-01 -2.66757425e-01  5.21526448e+01]
 [ 1.19510164e-01  6.07816803e-01  2.59509346e+02]
 [-1.33230239e-04 -6.77076891e-04  1.00000000e+00]]
```

Figure 44 The best H found by RANSAC

The result shown above is the best-performing H found by RANSAC. In the total 30 samplings, each H is calculated using 15 sampled matches (set seed to 42), and the liners were selected with a distance (symmetric transfer error) threshold of 0.08.

## 6. Warp and stitch the results

```
Assignment_3 - task2.ipynb

1 def calculate_new_size(img, H):
2     """ Calculate the transformed image size.
3
4     Args:
5         img (np.ndarray): The input image to be trasnformed.
6         H (np.ndarray): The transformation matrix.
7
8     Returns:
9         height, width, shift_u, shift_v (tuple(int, int, int, int)): The new size and shift infromation.
10
11    """
12    corners = np.array([[0, 0, 1],
13                        [0, img.shape[1], 1],
14                        [img.shape[0], 0, 1],
15                        [img.shape[0], img.shape[1], 1]]) # bottom right
16    # transformed coordinates by H
17    corners_trans = (H @ corners.T).T # ((3, 3) @ (4, 3).T).T = (4, 3)
18    u2 = corners_trans[:, 0] / corners_trans[:, 2]
19    v2 = corners_trans[:, 1] / corners_trans[:, 2]
20    # find the new height and width of the transferred image
21    height = np.ceil(u2.max() - u2.min()).astype(int)
22    width = np.ceil(v2.max() - v2.min()).astype(int)
23    # find the shift in both horizontal and vertical directions
24    # to calculate the homography coordinates of the tranformed image
25    shift_u = np.ceil(u2.min() - 0).astype(int)
26    shift_v = np.ceil(v2.min() - 0).astype(int)
27    return height, width, shift_u, shift_v
28
29
30 height, width, shift_u, shift_v = calculate_new_size(mountain1, best_H)
31 print(f"new image height: {height}, width: {width}")
32 print(f"shift vertically: {shift_u}, horizontally shift_v: {shift_v}")
```

Figure 45 The function used to calculate the new size and shift information for inverse wrapping

```
new image height: 522, width: 772
shift vertically: -131, horizontally shift_v: 260
```

For mountain1, the transformed image size has a height of 522 and a width of 772. The shift information is used in the later inverse warping with the bilinear interpolation technique.

```

Assignment_3 - task2.ipynb

1 def inverse_warping(img, H):
2     """Inverse warping for the input image with the transformation matrix.
3
4     Args:
5         img (np.ndarray): The input image to be transformed.
6         H (np.ndarray): The transformation matrix
7
8     Returns:
9         res (np.ndarray): The transformed image.
10    """
11    # calculate the size and shift information for later use
12    new_height, new_width, shift_u, shift_v = calculate_new_size(mountain1, best_H)
13    # initialise the transformed image to return
14    res = np.zeros((new_height, new_width, 3 if img.ndim == 3 else 1))
15    # check if (u, v) is valid in img
16    cond = lambda x, y: (0 <= x < img.shape[0]) and (0 <= y < img.shape[1])
17    # iterate over the pixels in res
18    for u2 in range(new_height):
19        for v2 in range(new_width):
20            # the coordinate in target image (X_2 = best_H @ X_1)
21            X_2 = np.array([u2 + shift_u, v2 + shift_v, 1])
22            # the coordinate in the query image (inv(best_H) @ X2 = X1)
23            X_1 = np.linalg.inv(H) @ X_2
24            # transform the homography coordinates to Cartesian coordinates
25            u1 = X_1[0] / X_1[-1]
26            v1 = X_1[1] / X_1[-1]
27            if cond(u1, v1):
28                if int(u1) == u1 and int(v1) == v1:
29                    res[u2, v2] = img[int(u1), int(v1)]
30                else:
31                    # bilinear interpolation
32                    x1, x2 = int(np.floor(u1)), int(np.ceil(u1))
33                    y1, y2 = int(np.floor(v1)), int(np.ceil(v1))
34                    if cond(x1, y2) and cond(x2, y2) and cond(x1, y1) and cond(x2, y1):
35                        if x1 == x2:
36                            res[u2, v2] = img[x1, y2] * (y - y1) / (y2 - y1) + img[x1, x1] * (y2 - y) / (y2 - y1)
37                        elif y1 == y2:
38                            res[u2, v2] = img[x2, y2] * (x2 - x) / (x2 - x1) + img[x1, y2] * (x - x1) / (x2 - x1)
39                        else:
40                            R1 = img[x1, y2] * ((x2 - u1) / (x2 - x1)) + img[x2, y2] * ((u1 - x1) / (x2 - x1))
41                            R2 = img[x1, y1] * ((x2 - u1) / (x2 - x1)) + img[x2, y1] * ((u1 - x1) / (x2 - x1))
42                            res[u2, v2] = R1 * (v1 - y1) / (y2 - y1) + R2 * (y2 - v1) / (y2 - y1)
43
44    return res.astype('uint8')
45
46 mountain1_trans = inverse_warping(mountain1, best_H)
47
48 # plot the trasnformed result
49 fig, ax = plt.subplots()
50 ax.imshow(mountain1_trans)
51 ax.set_title(f"Trasnformed mountain1\n{mountain1_trans.shape}")
52 ax.set_axis_off()

```

Figure 46 Inverse warping for mountain1



Figure 47 The transformed mountain1

```

Assignment_3 - task2.ipynb

1 # scale the transformed mountain1
2 ratio = 0.89
3 mountain1_resize = cv2.resize(mountain1_trans, (int(mountain1_trans.shape[1] * ratio),
4                                                 int(mountain1_trans.shape[0] * ratio)))
5
6 # stitch images (left: mountain2, right: mountain1)
7 stitched_img = np.zeros((374, mountain2.shape[1] + mountain1_resize.shape[1], 3)).astype('uint8')
8 mountain2_cut = 313
9 mountain1_cut = 70
10 stitched_img[:, :mountain2_cut] = mountain2[:, :mountain2_cut]
11 stitched_img[:, mountain2_cut:(mountain2_cut + (mountain1_resize.shape[1] - mountain1_cut))] = mountain1_resize[mountain1_resize.shape[0] - 374:, mountain1_cut:]
12
13 plt.figure(figsize=(12, 8))
14 plt.title("The Stitched Result\n(stitched_img.shape)")
15 plt.imshow(stitched_img[:, :mountain2_cut + (mountain1_resize.shape[1] - mountain1_cut)])
16 plt.axis('off')
17 plt.show()

```

Figure 48 Stitch two images

The stitching process includes scaling for mountain2. The stitch points from both images are selected based on the visual results from the previous SIFT keypoints and descriptors.

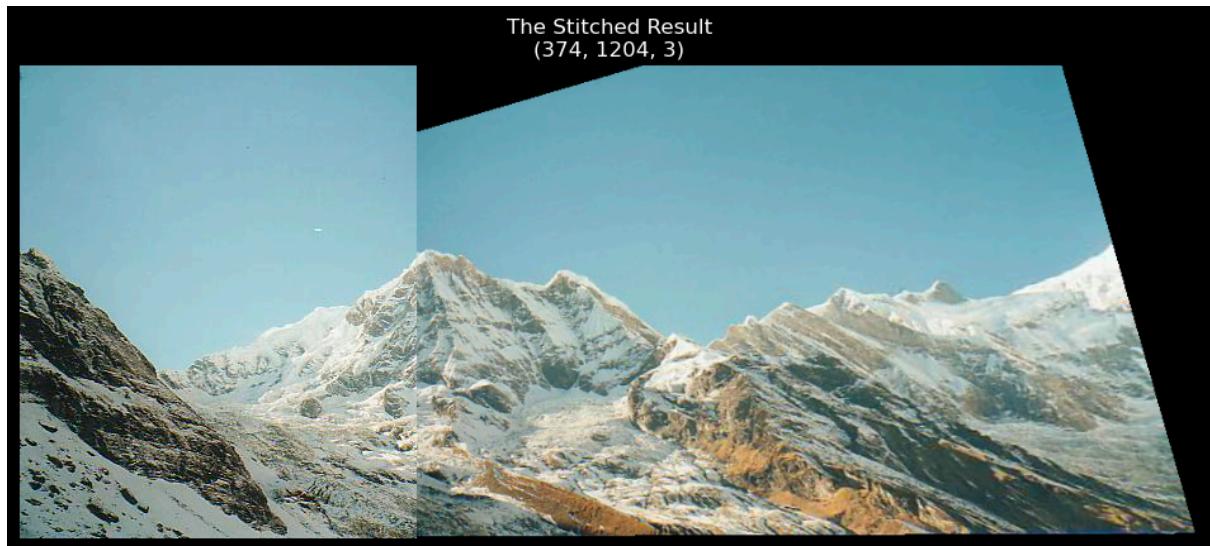


Figure 49 The stitch result (left: mountain2, right: mountain1)

## 7. Discussion

First, we use `ginput()` to select 8 correspondences from both images and save the correspondence coordinates into “`point_1.npy`” for `mountain1` and “`point_2.npy`” for `mountain2`.

```
Assignment_3 - select_points.py

1 import numpy as np
2 import matplotlib.pyplot as plt
3
4 path = 'Task2/images/mountain2.jpg'
5 save_path = 'Task2/mountain2_point.npy'
6 # Load an image
7 mountain = plt.imread(path)
8 # Display the image
9 plt.imshow(mountain)
10 # Use plt.ginput() to select points
11 selected_points = plt.ginput(n=8, timeout=0)
12 # Close the plot
13 plt.close()
14 # Print the selected points
15 print("Selected points:", selected_points)
16 np.save(save_path, selected_points)
```

Figure 50 Select 8 correspondences

The manually selected correspondences are:



*Figure 51 The manually selected correspondences*

Use homography\_w\_normalisation to calculate the homography matrix:

```
Assignment_3 - task2.ipynb
1 # calculate the homography matrix with the manually selected correspondences
2 H_select = homography_w_normalisation(mountain1, points_1[:, 0], points_1[:, 1],
3                                         mountain2, points_2[:, 0], points_2[:, 1])
4 print(H_select)
```

*Figure 52 The homography matrix calculated with the 8 correspondences*

```
[ [ 6.37507672e-01  1.00711146e-01  2.59295683e+02]
[ -2.67214379e-01  8.32356450e-01  5.36026067e+01]
[ -6.17521544e-04 -1.86741367e-04  1.00000000e+00]]
```

*Figure 53 The calculated homography matrix*

Compared with the results (RANSAC) in Task 2.5, some entries have significant differences.

The following plot shows the stitch results by using the manually selected correspondences. Notice that we use the similar stitch location as the previous one for a better comparison:

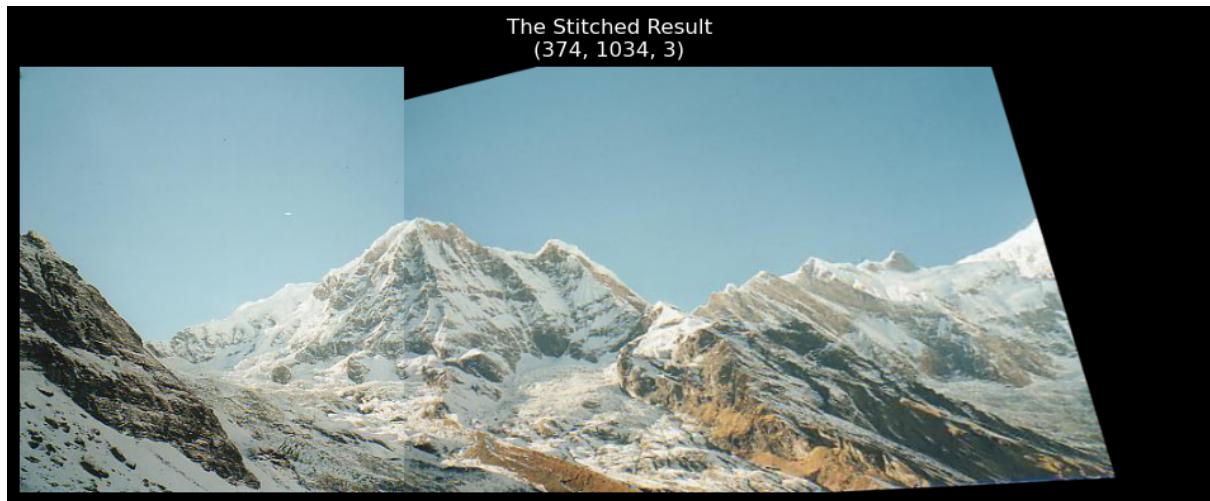


Figure 54 The stitch result from manually selected correspondences

Visually, the new stitch result is more accurate, especially around the stitch line.

#### Theoretical analysis:

- Biased model assumption: The homography estimation tasks normally assume a planar scene and a pinhole camera model. If the scene is not strictly planar, the estimated homography may introduce bias, especially in regions where the scene deviates from planarity.
- Improper optimisation method: Some homography estimation techniques rely on optimisation methods like RANSAC or least squares. Some methods may converge to local minima or be sensitive to outliers, which result in biased estimates if not appropriately handled.
- Parameterisation: The choice of parameterisation for representing the homography matrix can introduce the estimation bias, such as under-parameterisation or over-parameterisation.
- The distribution of correspondence: The distribution of correspondences significantly influences homography estimation. A uniform distribution allows for accurate estimation across the entire scene, while clustering, outliers, sparse distributions, and non-uniform scales can introduce biases and affect the reliability of the estimated homography.

#### Empirically analysis:

- Noise in correspondences: Homography estimation methods are sensitive to noise in correspondences between image points. When there are outliers in the correspondences, the estimated homography tends to be biased towards the outliers.
- Insufficient correspondences: Normally, a homography requires a minimum number of correspondences. Insufficient correspondence may be biased due to the limited information available.

- Lens distortion: Homography estimation typically neglects lens distortions. In real-world scenarios, lens distortions can introduce bias, especially in wide-angle lenses or fisheye cameras.
- Scale ambiguity: Homography estimation alone cannot determine scale, leading to scale ambiguity.
- Inaccurate camera calibration: Homography estimation relies on accurate camera calibration parameters. Errors in camera calibration, such as inaccuracies in focal length or principal point estimation, can propagate into biased homography results.