

# 3D Vision 3

Week 9

Two-view Geometry: Triangulation

Two-view Geometry: Stereo

Two-view Geometry: Optical Flow

# Announcements

- **Survey 2 – Week 8|9** is live on Wattle, please provide a response to improve the course for future years
- **Class Representatives:** another mechanism to provide feedback
  - Arjun Raj (u7526852@anu.edu.au) [COMP4528]
  - Ruxuan Li (u7763307@anu.edu.au) [ENGN4528]
  - Jiabao Han (u7765516@anu.edu.au) [COMP6528]
  - Xingchen Zhang (u7670173@anu.edu.au) [COMP6528]
- **Exam:** 9:00am–12:15pm Saturday 1 June 2024
  - 7-11 Barry Drive, First Floor Left Side

# Weekly Study Plan: Overview

Wk	Starting	Lecture	Lab	Assessment
1	19 Feb	Introduction	X	
2	26 Feb	Low-level Vision 1	1	
3	4 Mar	Low-level Vision 2	1	
		Mid-level Vision 1		
4	11 Mar	Mid-level Vision 2	1	CLab1 report due Friday
		High-level Vision 1		
5	18 Mar	High-level Vision 2	2	
6	25 Mar	High-level Vision 3 <sup>1</sup>	2	
	1 Apr	Teaching break	X	
	8 Apr	Teaching break	X	
7	15 Apr	3D Vision 1	2	CLab2 report due Friday
8	22 Apr	3D Vision 2	3	
9	29 Apr	3D Vision 3	3	
10	6 May	3D Vision 4	3	
		Mid-level Vision 3		
11	13 May	High-level Vision 4	X	CLab3 report due Friday
12	20 May	Course Review	X	

# Weekly Study Plan: Part B

Wk	Starting	Lecture	By
7	15 Apr	3D vision: introduction, camera model, single-view geometry	Dylan
8	22 Apr	3D vision: camera calibration, two-view geometry (homography)	Dylan
9	29 Apr	3D vision: two-view geometry (epipolar geometry, triangulation, stereo)	Dylan
10	6 May	3D vision: multiple-view geometry	Weijian
		Mid-level vision: optical flow, shape-from-X	Dylan
11	13 May	High-level vision: self-supervised learning, detection, segmentation	Dylan
12	20 May	Course review	Dylan

# Outline

1. Two-view Geometry: Triangulation
2. Two-view Geometry: Stereo
3. Two-view Geometry: Optical flow

# Essential & Fundamental Matrices: Summary

- Algebraic representations of epipolar geometry:
  - Projection matrices (given intrinsics + extrinsics)
  - **Essential matrix** (given intrinsics):  $\mathbf{E} = [\mathbf{t}]_{\times} \mathbf{R}$
  - **Fundamental matrix**:  $\mathbf{F} = \mathbf{K}'^{-T} [\mathbf{t}]_{\times} \mathbf{R} \mathbf{K}^{-1} = [\mathbf{e}']_{\times} \mathbf{K}' \mathbf{R} \mathbf{K}^{-1}$  [HZ p.244]
    - Using identity  $[\mathbf{v}]_{\times} \mathbf{M} = \mathbf{M}^{-T} [\mathbf{M}^{-1} \mathbf{v}]_{\times}$  and  $\mathbf{e}' = \mathbf{K}' \mathbf{t}$  (image of camera centre 1)
- Epipolar constraint for corresponding points  $\{\mathbf{x}, \mathbf{x}'\}$ :
$$\mathbf{x}'^T \mathbf{F} \mathbf{x} = 0$$
  - $\mathbf{F}$  is rank 2 and is known only up to scale  $\rightarrow$  7 DoF
  - What is  $\mathbf{F} \mathbf{x}$ ?
- **Estimation**: DLT (again!); assemble a matrix  $\mathbf{A}$ , compute the SVD, enforce rank 2 (with another SVD); or use a nonlinear solver

# The Fundamental Matrix F: Recovering P & P'

- Since we can set the world coordinate system arbitrarily, set

$$\begin{aligned}P &= [I \mid 0] \\P' &= [[e']]_x F \mid e']\end{aligned}$$

where epipole can be found by solving

$$F^T e' = 0$$

via SVD/eigendecomposition [HZ p. 256]

# Triangulation

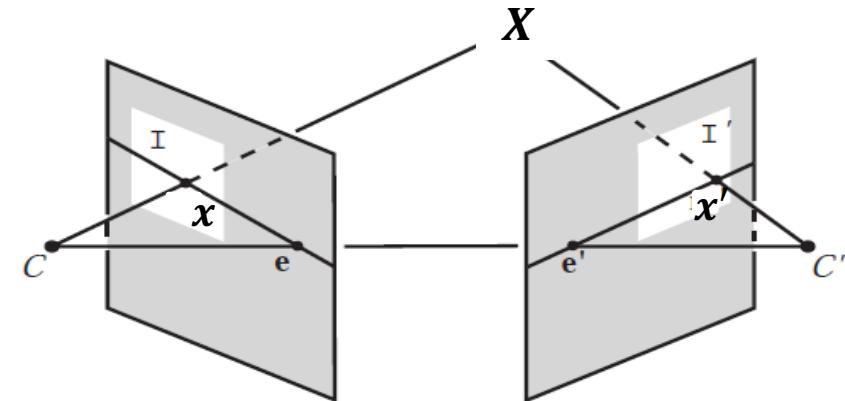
Two-view Geometry

# Triangulation

- **Problem:** Given image points  $x$  and  $x'$  in correspondence across two or more images (taken from calibrated cameras with known intrinsics and extrinsics), compute the 3D location  $X$

# Triangulation: DLT Linear Solution

- Generally, rays  $\overrightarrow{Cx}$  and  $\overrightarrow{C'x'}$  will not intersect exactly
- Can solve via the **DLT algorithm** (again!), finding a **least squares solution** to a system of equations
- Further reading:
  - HZ pp. 312–313



$$\begin{aligned} \cancel{x \times (P X) = 0 \text{ and } x' \times (P' X) = 0} \\ \therefore AX = \begin{bmatrix} u\mathbf{p}_3^T - \mathbf{p}_1^T \\ v\mathbf{p}_3^T - \mathbf{p}_2^T \\ u'\mathbf{p}'_3^T - \mathbf{p}'_1^T \\ v'\mathbf{p}'_3^T - \mathbf{p}'_2^T \end{bmatrix} \begin{bmatrix} x \\ y \\ z \\ 1 \end{bmatrix} = 0 \end{aligned}$$

# Triangulation: DLT Linear Solution

Given  $\mathbf{P}, \mathbf{P}', \mathbf{x}, \mathbf{x}'$

1. Precondition points and projection matrices

2. Create matrix  $\mathbf{A}$

3. SVD:  $\mathbf{A} = \mathbf{U}\Sigma\mathbf{V}^T$

4.  $\mathbf{X} = \mathbf{V}_{:, -1}$  (last column of  $\mathbf{V}$ )

5. Then refine with respect to a geometric error

$$\mathbf{x} = \begin{bmatrix} u \\ v \\ 1 \end{bmatrix} \quad \mathbf{x}' = \begin{bmatrix} u' \\ v' \\ 1 \end{bmatrix}$$

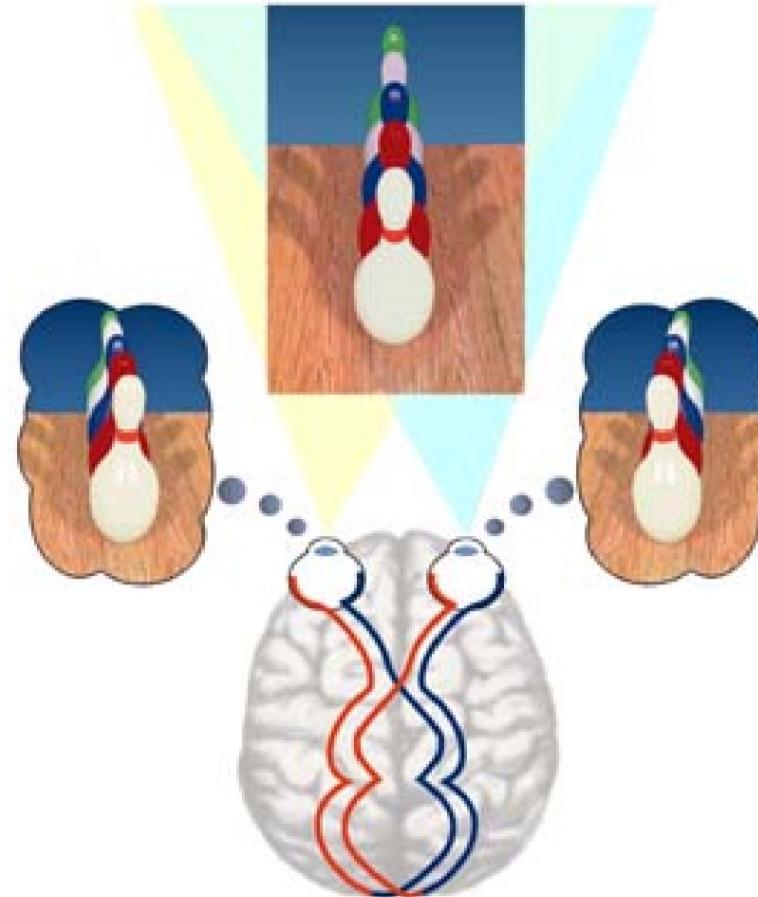
$$\mathbf{P} = \begin{bmatrix} \mathbf{p}_1^T \\ \mathbf{p}_2^T \\ \mathbf{p}_3^T \end{bmatrix} \quad \mathbf{P}' = \begin{bmatrix} \mathbf{p}'_1^T \\ \mathbf{p}'_2^T \\ \mathbf{p}'_3^T \end{bmatrix}$$

$$\mathbf{A} = \begin{bmatrix} u\mathbf{p}_3^T - \mathbf{p}_1^T \\ v\mathbf{p}_3^T - \mathbf{p}_2^T \\ u'\mathbf{p}'_3^T - \mathbf{p}'_1^T \\ v'\mathbf{p}'_3^T - \mathbf{p}'_2^T \end{bmatrix}$$

# Stereo Vision

Two-view Geometry

# Binocular Stereo Vision



# 3D Perception Using Stereo Vision

- Humans do it well from a *single* image and very well from *stereo* (binocular) images
- Mechanism is not well understood:
  - Biological components: some understanding
  - Algorithm: little understanding
- An intermediate goal of computer vision was to mimic the functionality of the biological system
  - *Not* to mimic its mechanics

# Stereo Vision

- Stereo matching computes **depth information** from two images
- Depth information can be used to...
  - Differentiate objects from the background
  - Differentiate objects from one another
  - Expose camouflaged objects
  - Navigate in environment avoiding obstacles

# Stereo Vision: Two Sub-Problems

## **Correspondence Problem**

- The problem of measuring the disparity of each point in the two eye (camera) projections

## **Interpretation Problem**

- The use of disparity information to recover the orientation and distance of surfaces in the scene

# Stereo Vision: Algorithmic Steps

- **Basic steps** to be performed in any **stereo imaging system**:
  - Image Acquisition
  - Camera Modelling
  - Feature Extraction
  - Image Matching
  - Depth Interpolation

# Stereo Image Acquisition

- Capturing two images with a very specific camera geometry

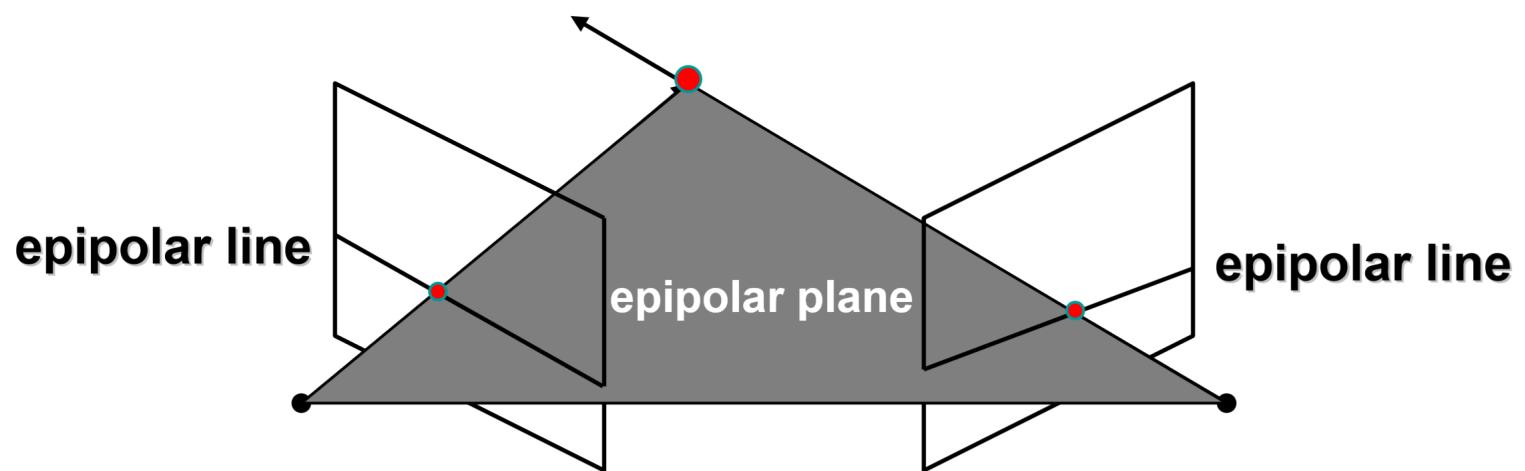


# Camera Modelling

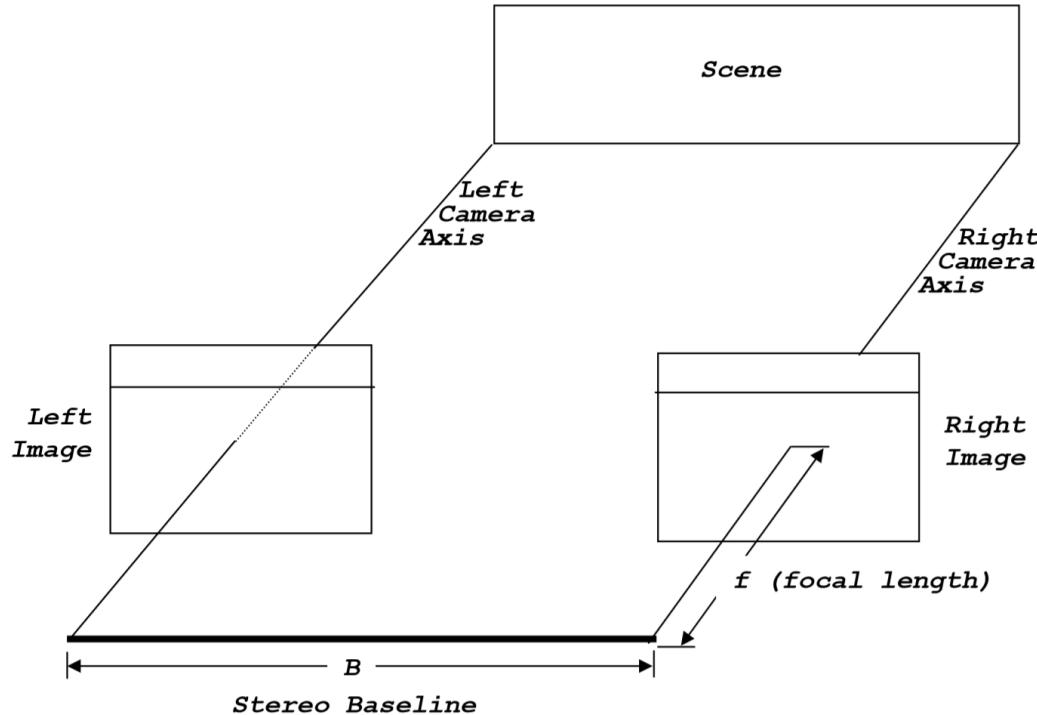
- For accurate depth results the camera parameters must be known
  - Intrinsic parameters: focal length, pixel skew/shape, distortion, etc.
  - Extrinsic parameters: relative rotation and translation



# General Epipolar Constraint



# Simple (Ideal) Stereo Geometry



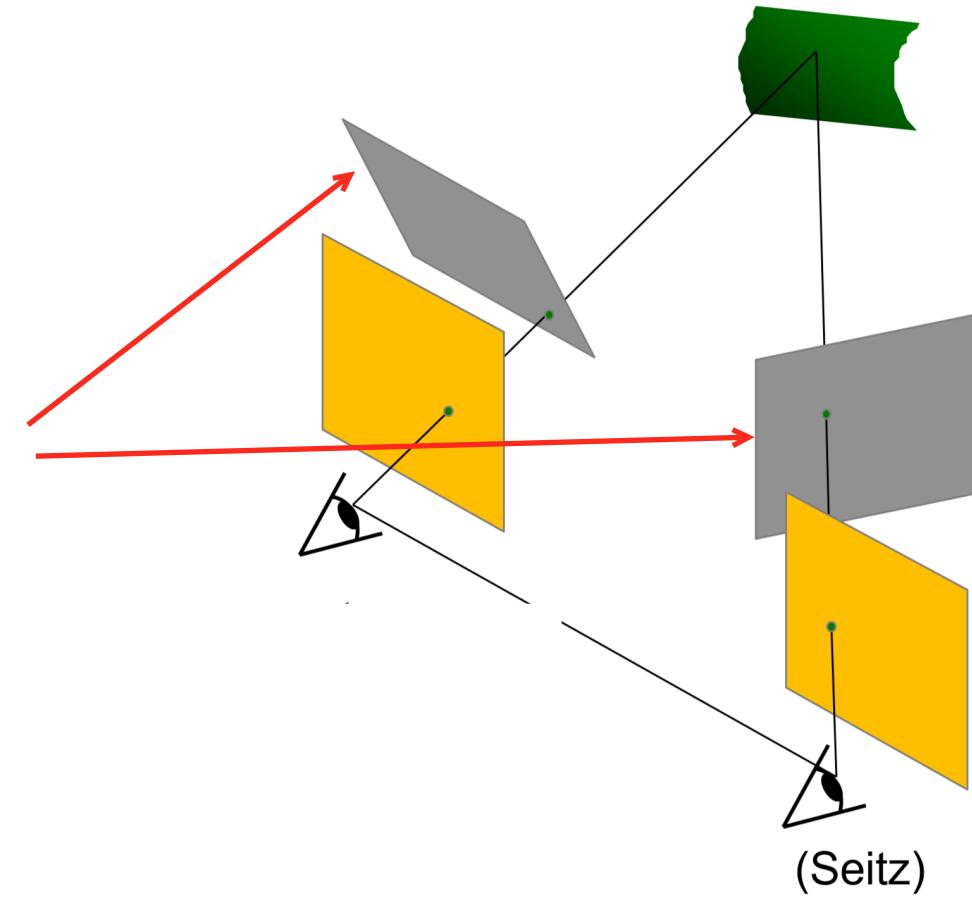
- Two slightly different images
- Extrinsics: identical rotations, translation has an offset in the X direction

# Simplest Case: Rectified Stereo Images

- Assumptions:
  - Image planes of cameras are parallel
  - Principal axes are at same height
  - Identical focal lengths
- Then the epipolar lines fall along the horizontal scan lines of the images
- We will assume images have been rectified so that epipolar lines correspond to scan lines:
  - Simplifies algorithms
  - Improves efficiency

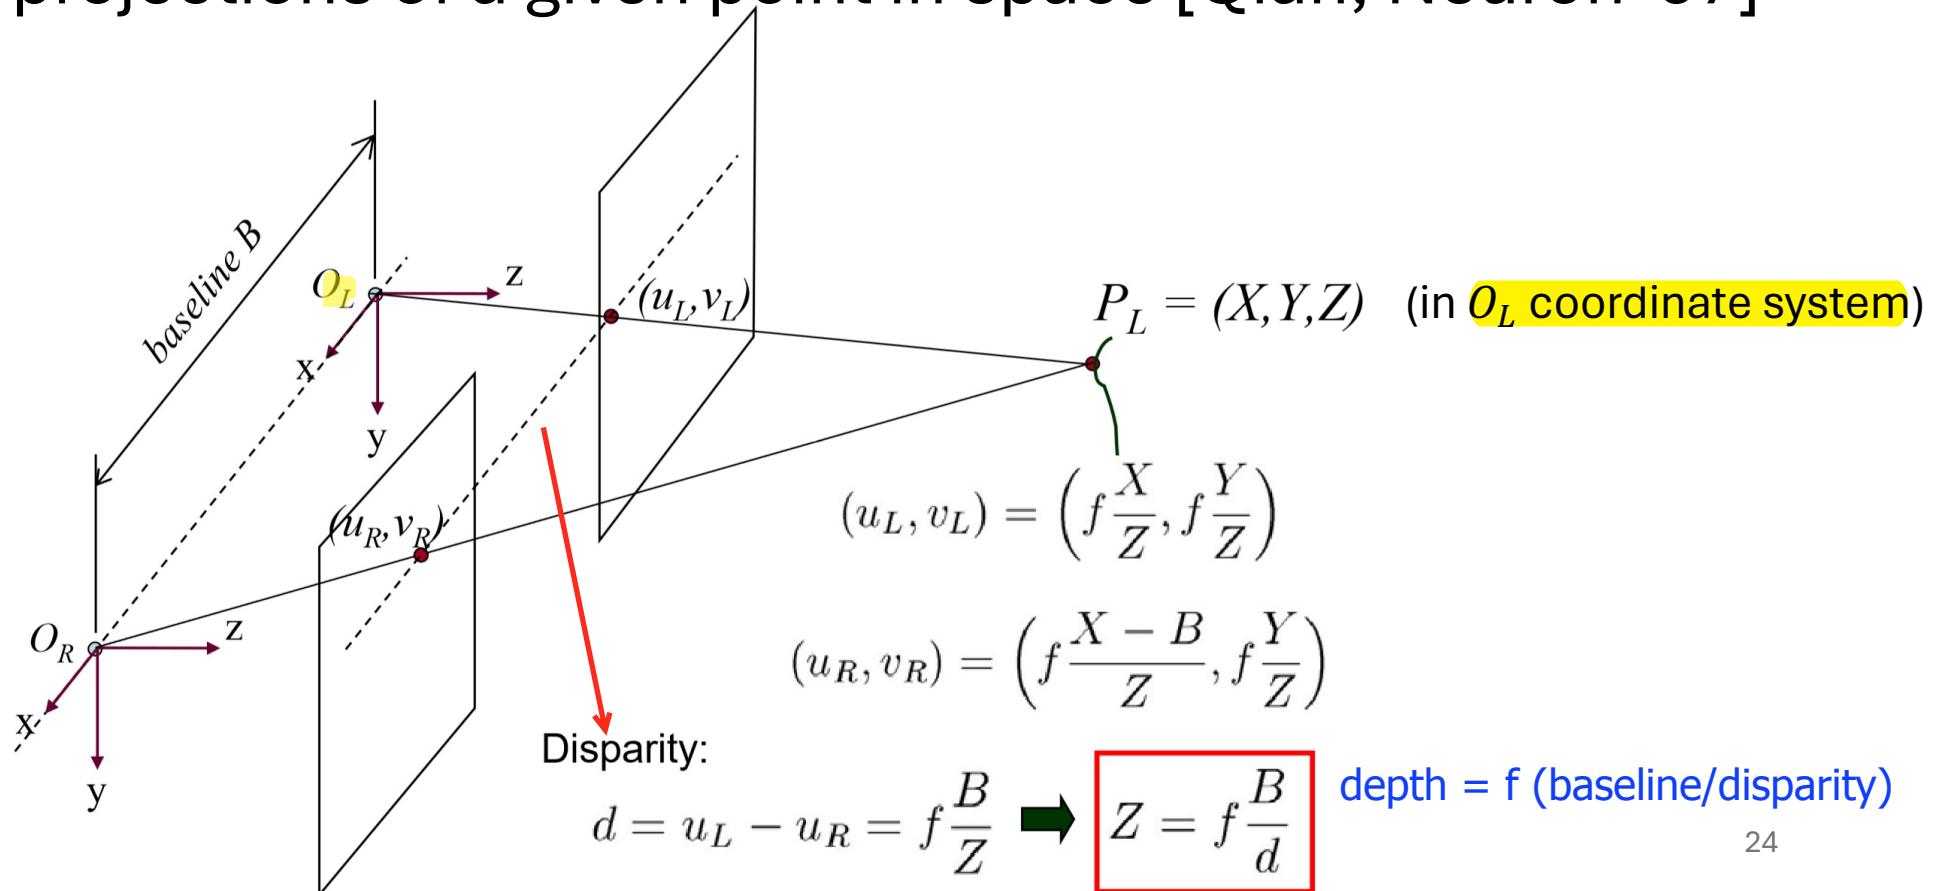
# Simplest Case: Rectified Stereo Images

- We can always achieve this ideal geometry with stereo-rectification
- **Image reprojection:**
  - Reproject image planes onto a common plane parallel to the line between optical centres using suitable planar homographies
  - E.g., using homography estimation
  - Not covered in detail here



# Disparity-to-Depth Computation

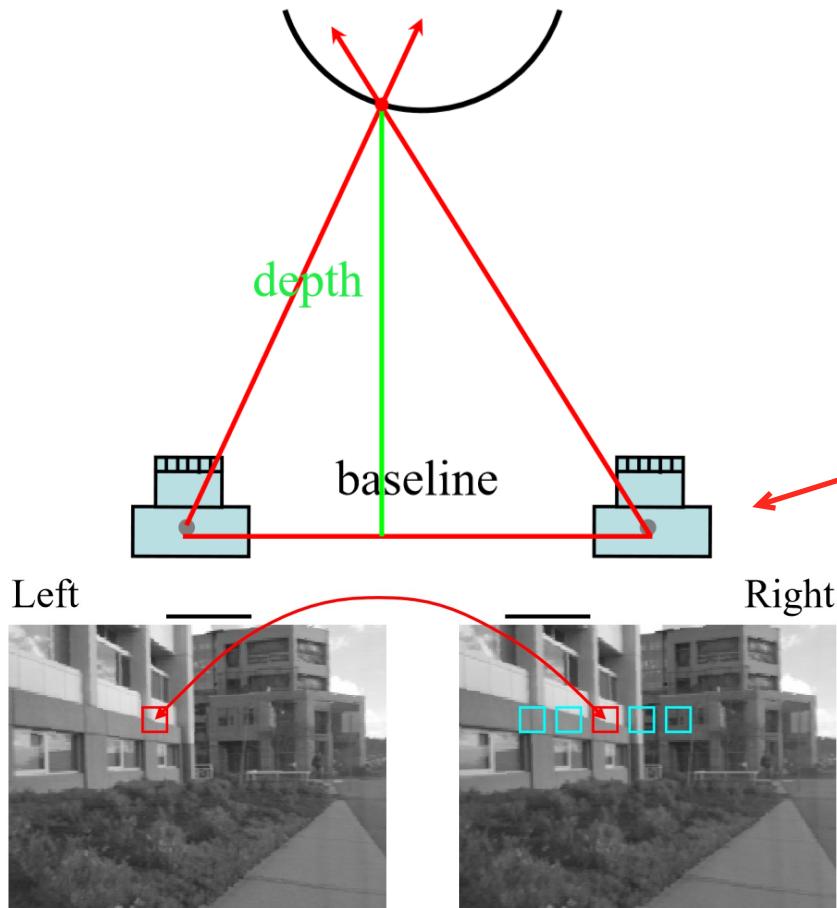
- **Disparity** (human vision): the positional difference between the two retinal projections of a given point in space [Qian, Neuron '97]



# The Correspondence Problem

- What should we match?
  - Objects?
  - Edges?
  - Pixels?
  - Super-pixels (set of pixels)?
  - ...
  - DINO features?

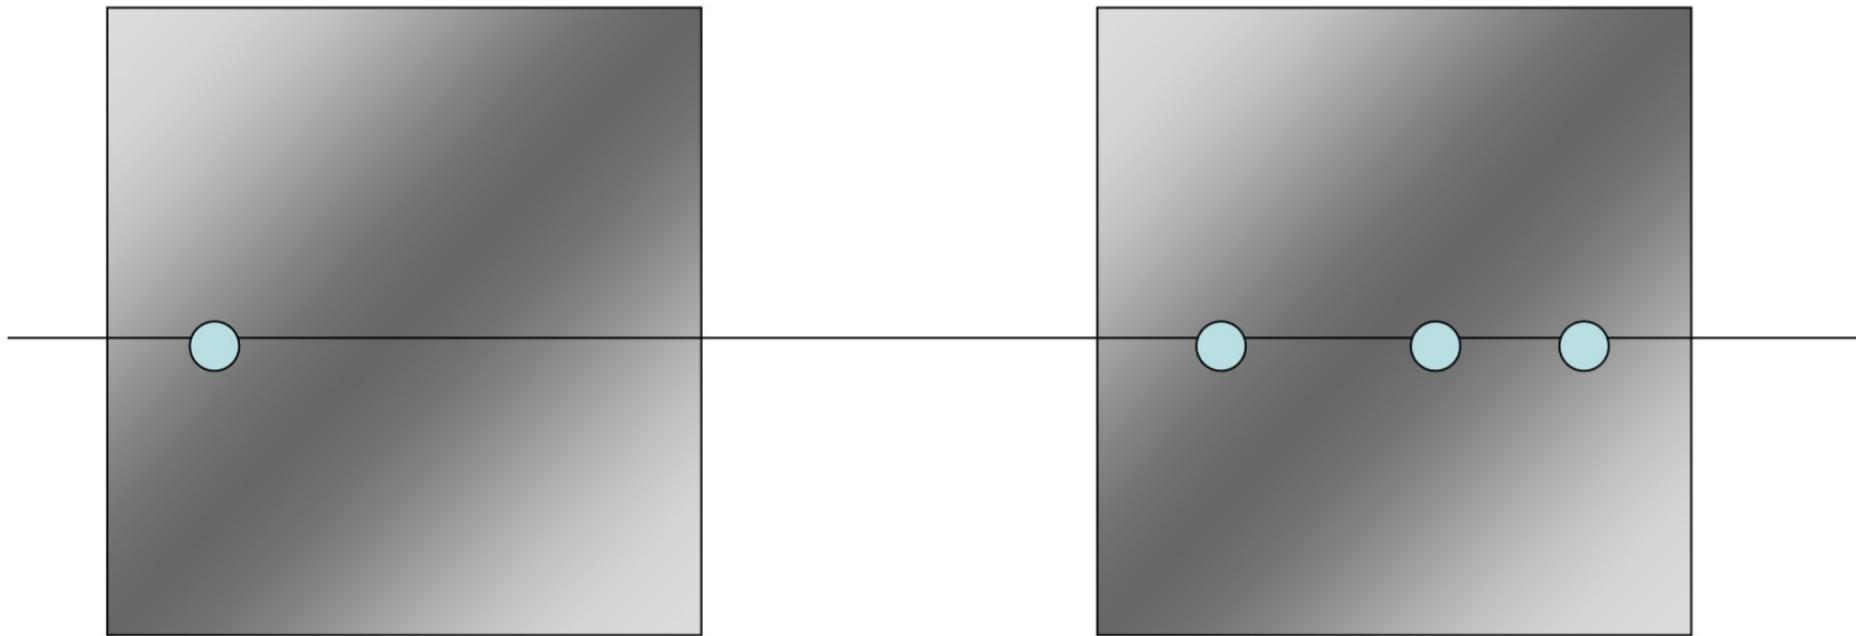
# The Correspondence Problem



- Triangulate on two images of the same point to **recover depth**
  - Feature matching across views
  - Calibrated cameras
- Matching windows across **horizontal scan lines**

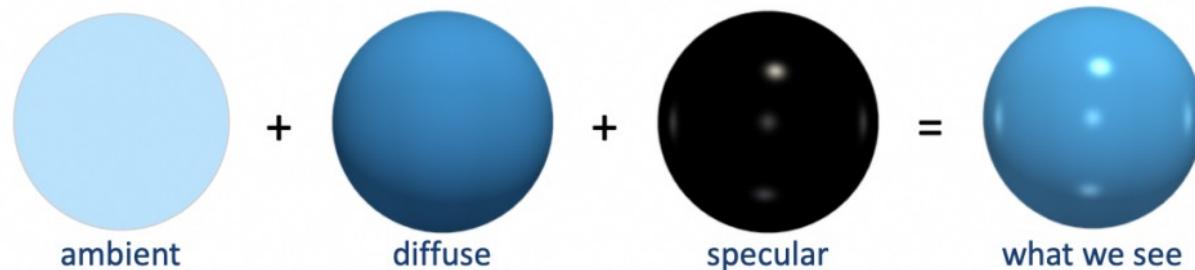
# The Correspondence Problem: The Epipolar Constraint

- The epipolar constraint removes some ambiguity

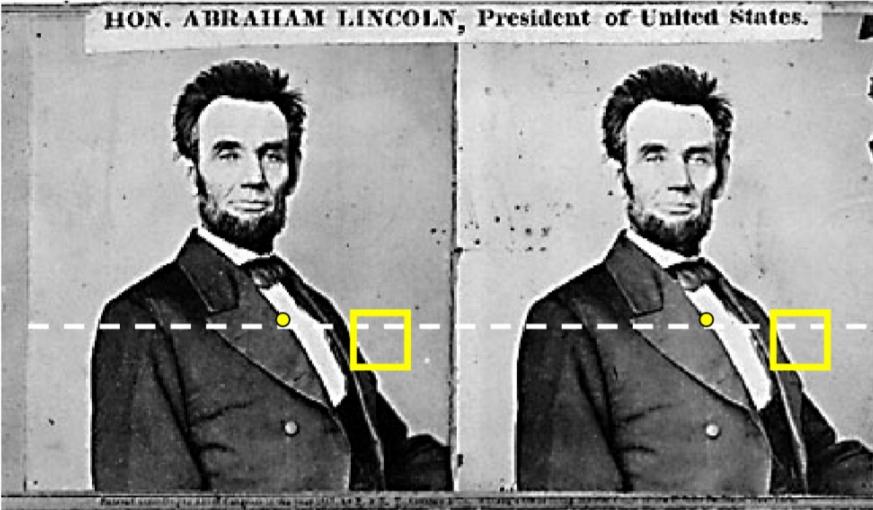


# The Correspondence Problem: The Colour Constancy “Constraint”

- Same world point has same intensity (or colour) in both images
  - True for Lambertian surfaces
    - A Lambertian surface has a brightness that is independent of viewing angle
  - Violations:
    - Noise
    - Specularity
    - Non-Lambertian materials
    - Pixels that contain multiple surfaces
    - Occlusions

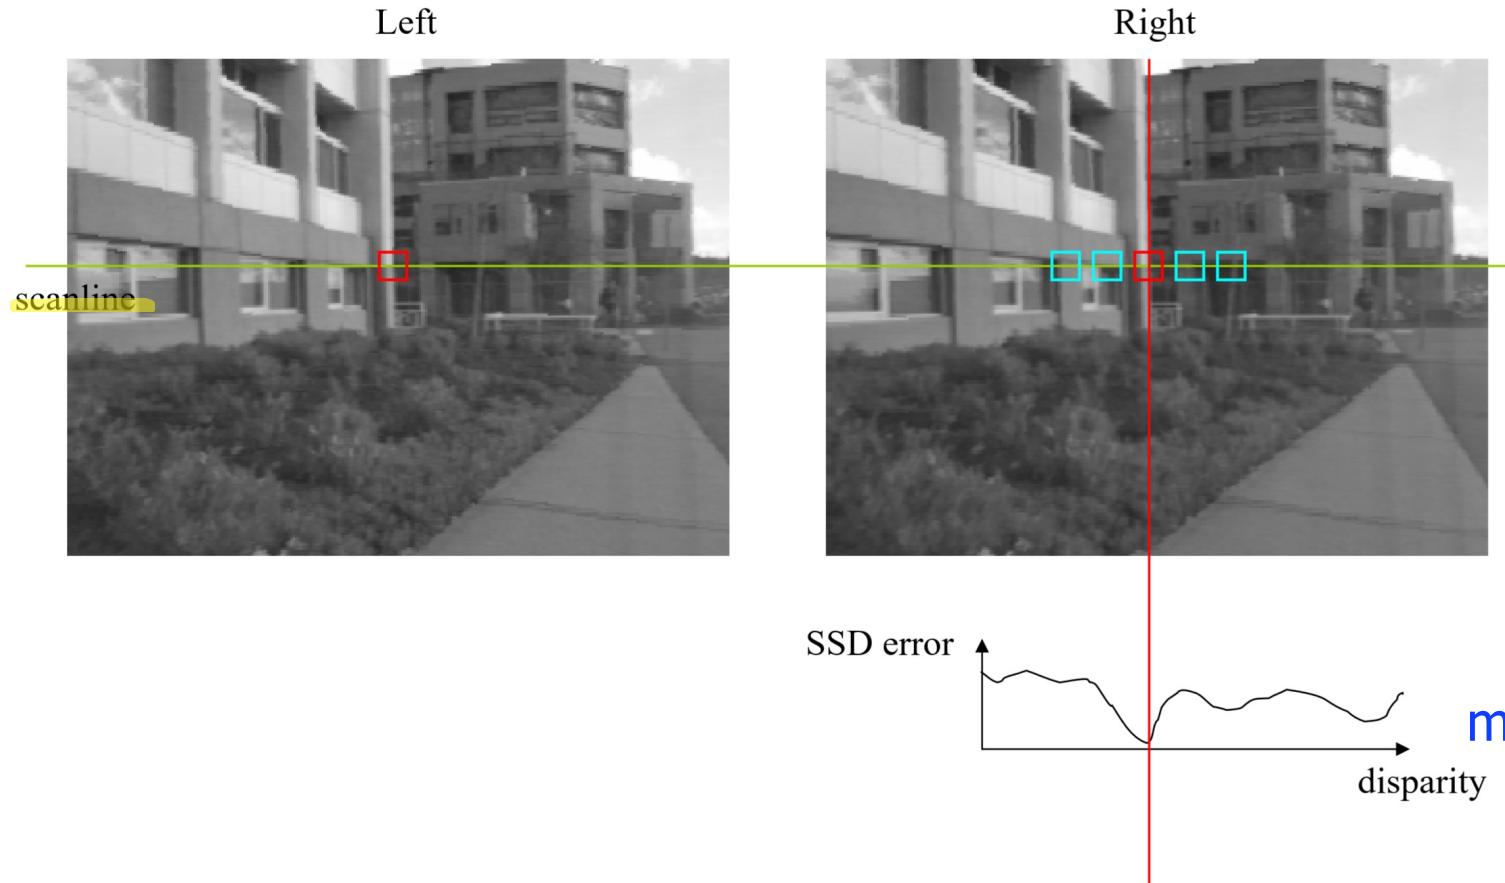


# The Correspondence Problem: Pixel Matching

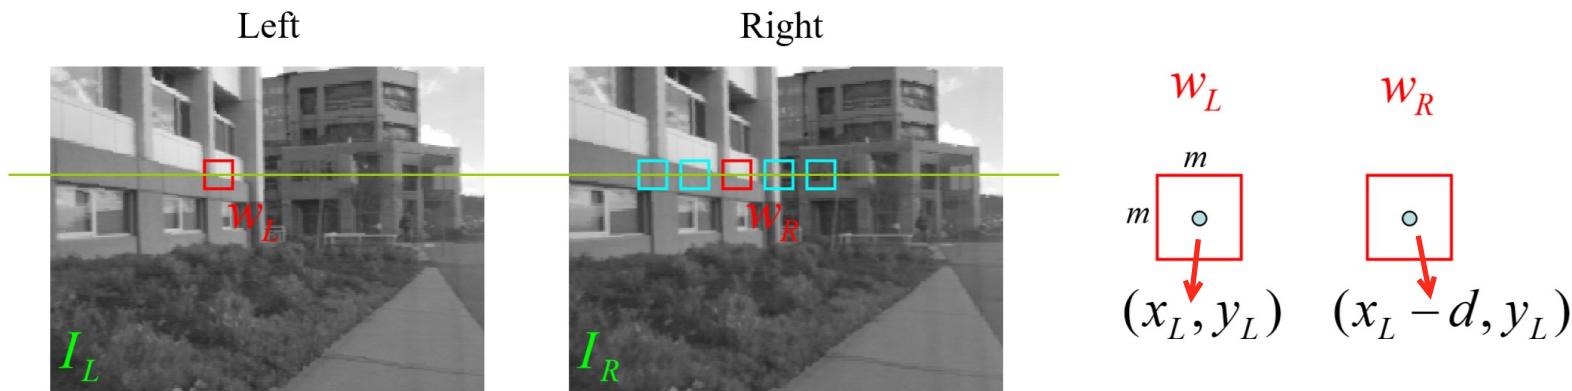


- For each epipolar line:
  - For each pixel in the left image:
    - Compare with every pixel on same epipolar line in right image
    - Pick pixel with minimum matching cost
  - Still too much ambiguity: match windows instead

# The Correspondence Problem: Correspondence Using Correlation



# Sum of Squared (Pixel) Differences



$w_L$  and  $w_R$  are corresponding  $m$  by  $m$  windows of pixels.

We define the window function :

$$W_m(x, y) = \{u, v \mid x - \frac{m}{2} \leq u \leq x + \frac{m}{2}, y - \frac{m}{2} \leq v \leq y + \frac{m}{2}\}$$

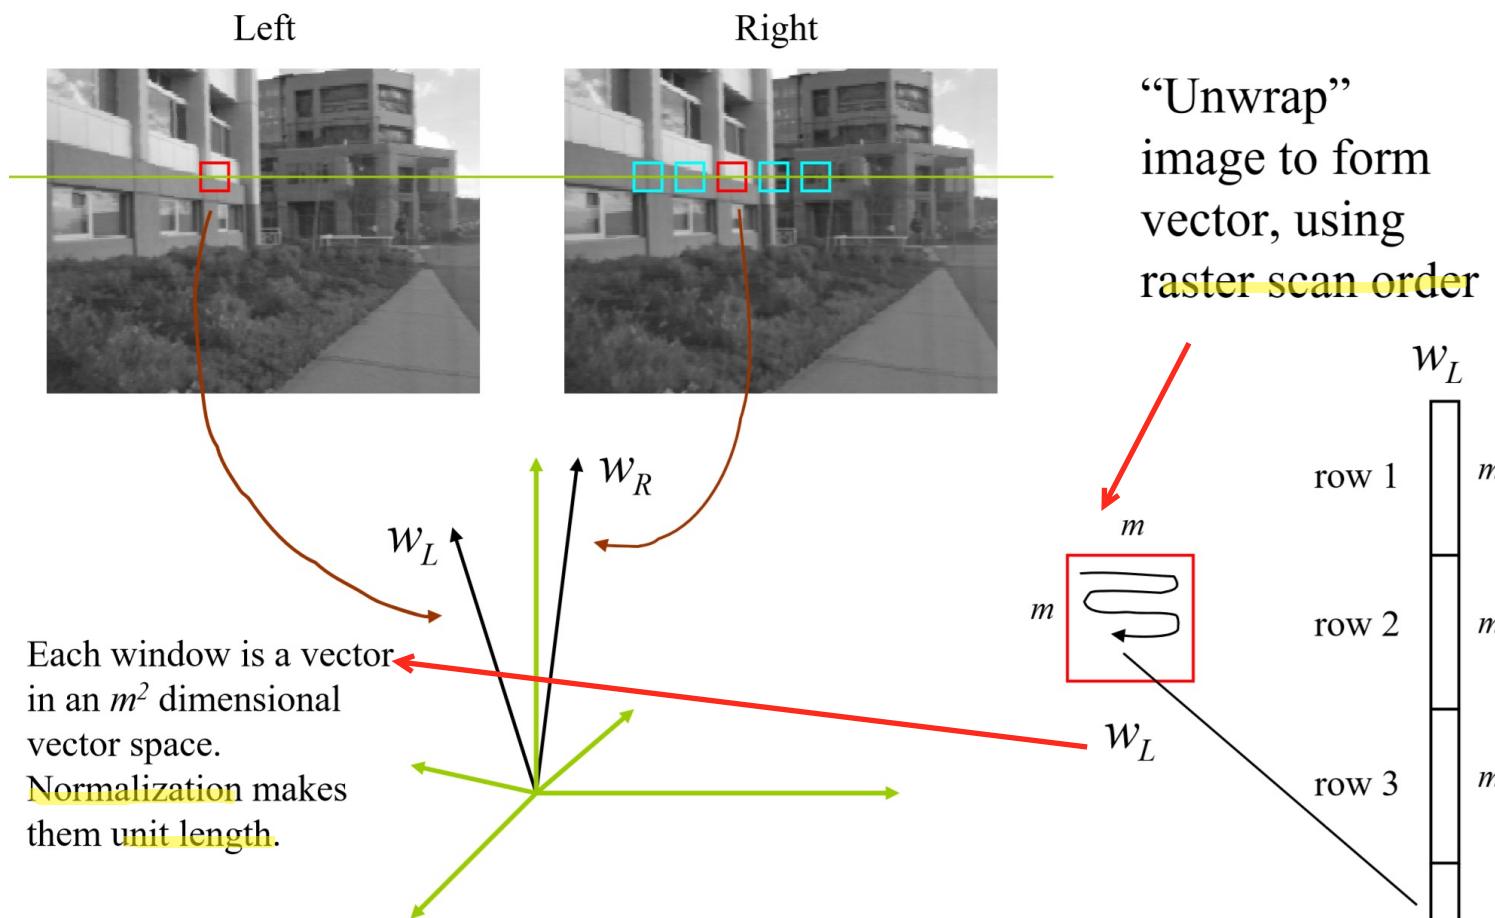
The SSD cost measures the intensity difference as a function of disparity :

$$C_r(x, y, d) = \sum_{(u, v) \in W_m(x, y)} [I_L(u, v) - I_R(u - d, v)]^2$$

d: the disparity

select the min disparity,  
then calculate the depth of the matched window  
by using depth = f (baseline/disparity)

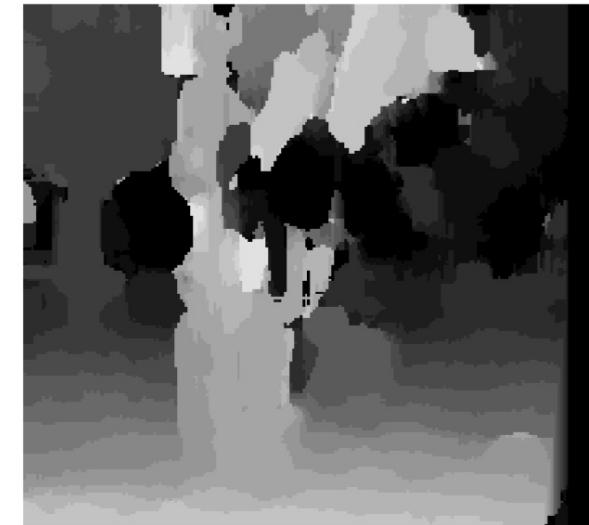
# Image Patches as Vectors



# Effect of Window Size



$W = 3$



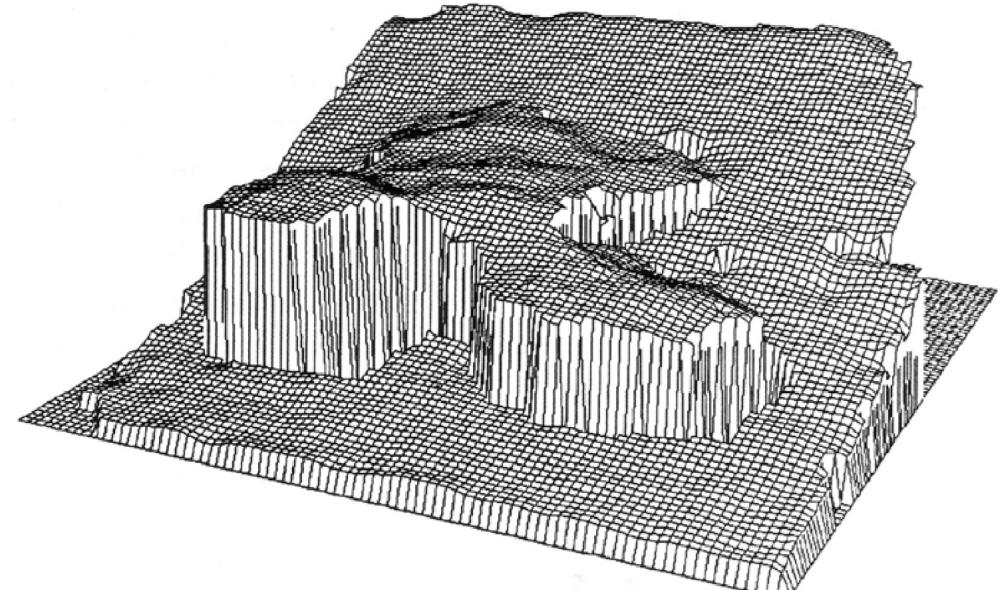
$W = 20$

- Improvement: use an adaptive window size (try multiple sizes and select best match)

# Example Stereo Pair

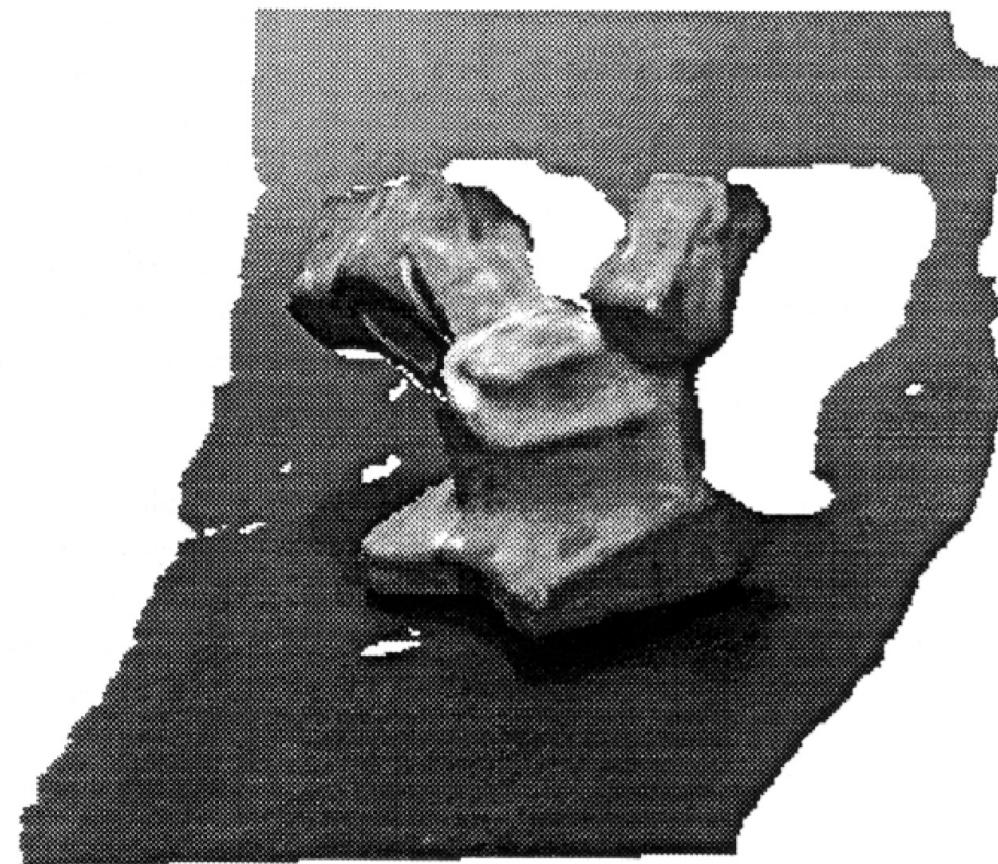


Left Camera



- Obtained **depth map** in 3D

# Novel View Synthesis

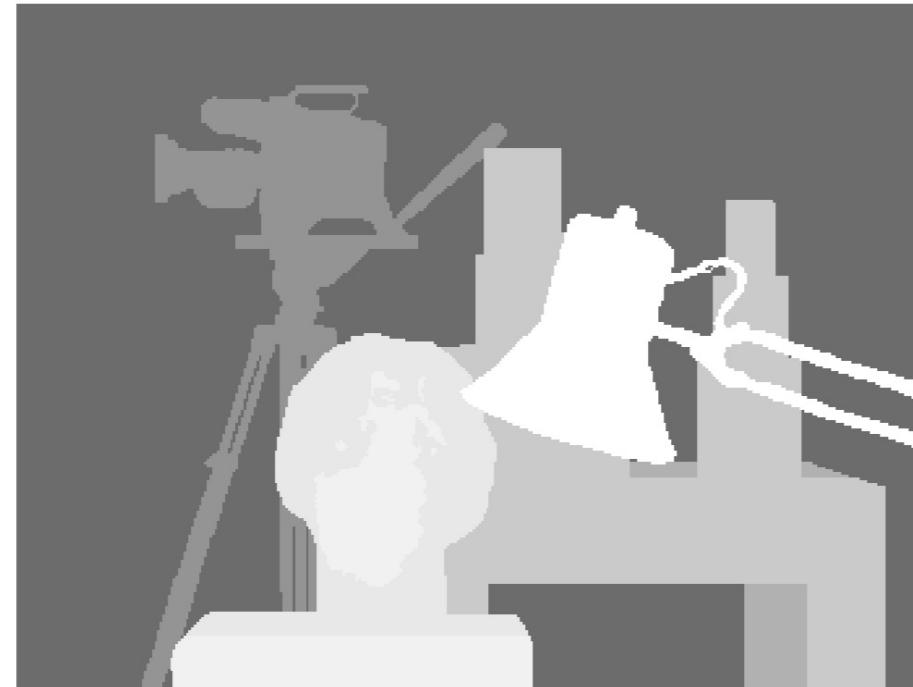


# A Taxonomy of Stereo Algorithms

- D. Scharstein and R. Szeliski, “A Taxonomy and Evaluation of Dense Two-Frame Stereo Correspondence Algorithms”, International Journal of Computer Vision, 47 (2002), pp. 7-42

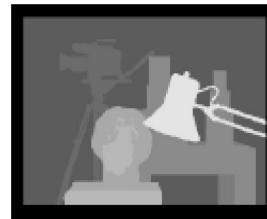


Scene



Ground truth

# A Taxonomy of Stereo Algorithms



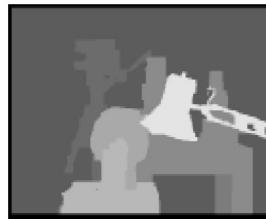
True disparities



19 – Belief propagation



11 – GC + occlusions



20 – Layered stereo



10 – Graph cuts



\*4 – Graph cuts



13 – Genetic algorithm



6 – Max flow



12 – Compact windows



9 – Cooperative alg.



15 – Stochastic diffusion



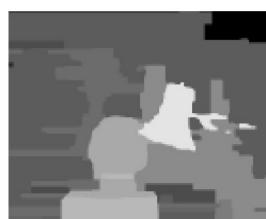
\*2 – Dynamic prgr.



14 – Realtime SAD



\*3 – Scanline opt.



7 – Pixel-to-pixel stereo

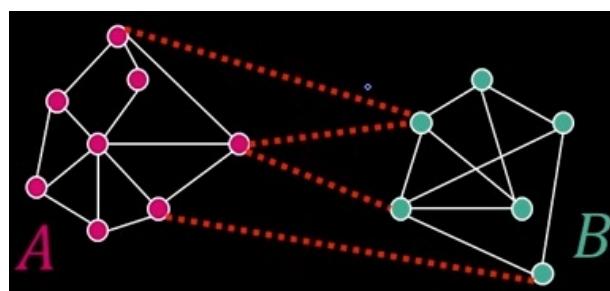
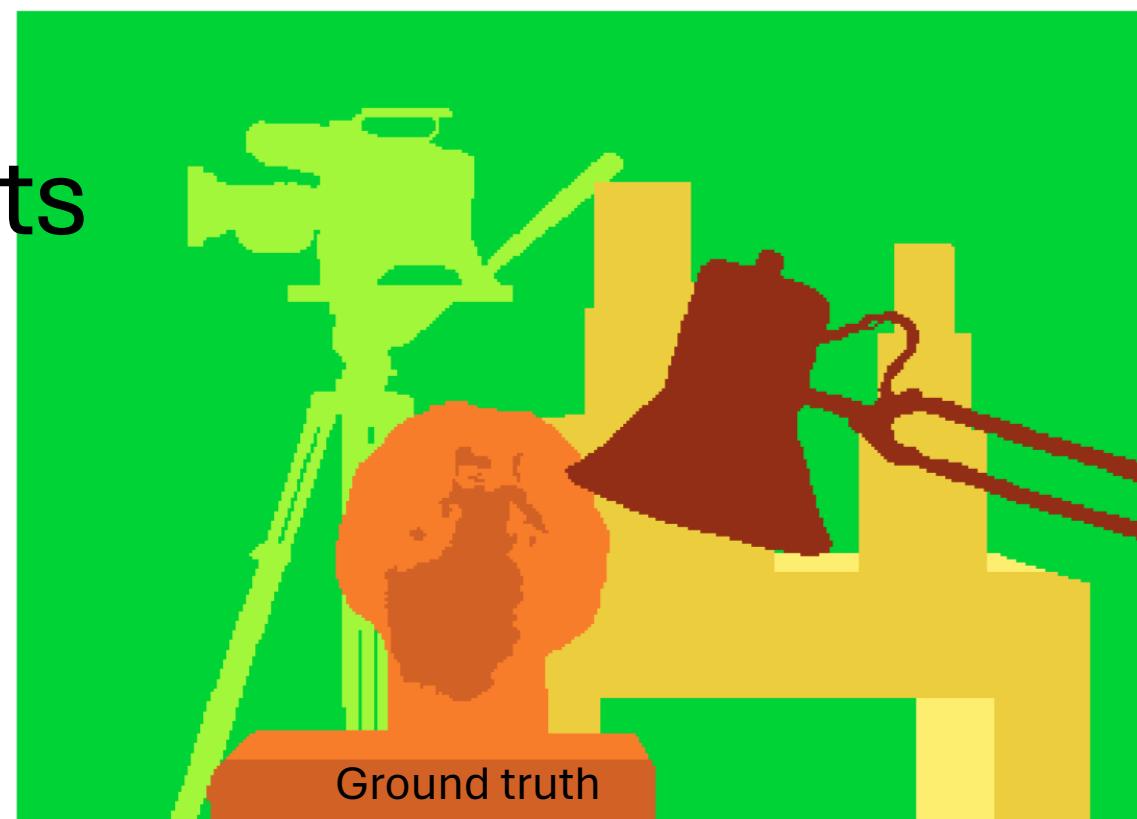
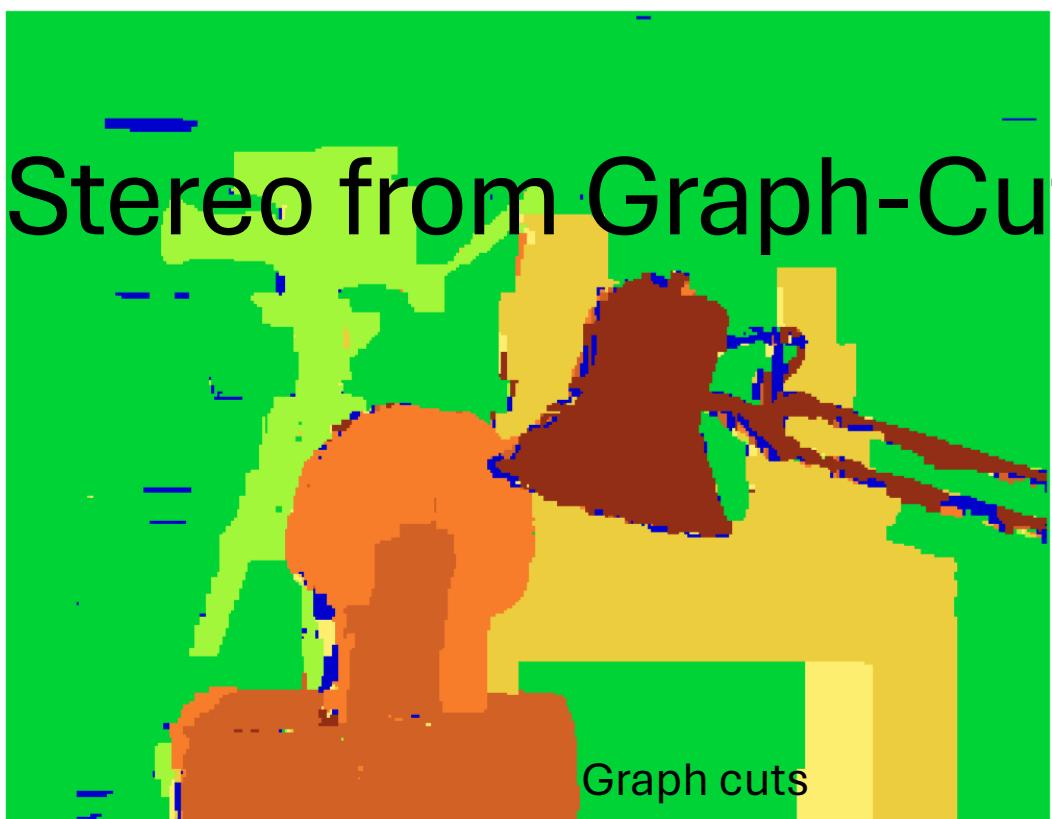


\*1 – SSD+MF  
Scharstein and Szeliski

# Stereo from Segmentation



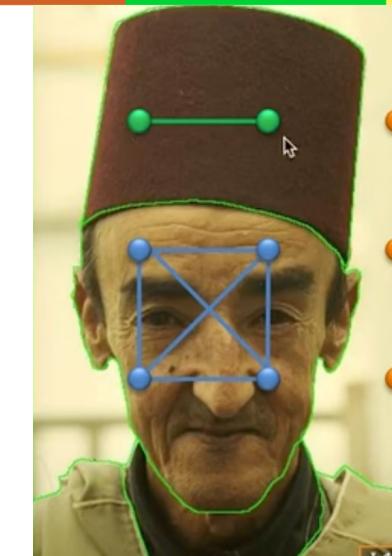
# Stereo from Graph-Cuts



An image is represented by a graph

$$Cut(A, B) = \sum_{p \in A, q \in B} w_{pq}$$

Minimise the cost of a cut,  
also called **min-cut**  
Each subgraph is an image  
segment



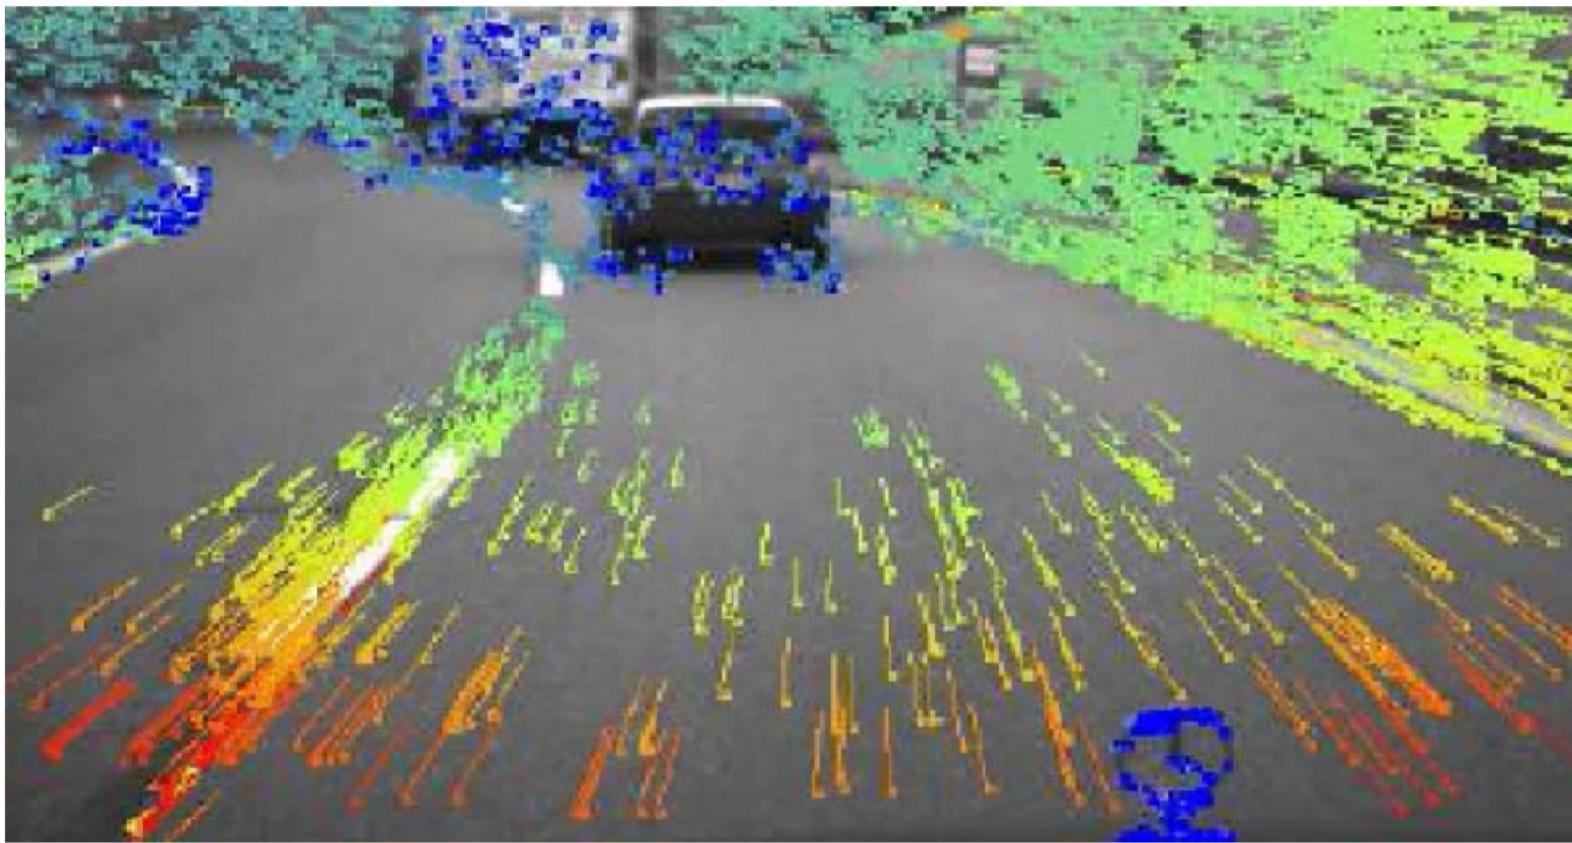
# Summary of Stereo Vision

- Stereo matching + triangulation
  - The correspondence problem and the interpretation problem
- Constraints:
  - Geometry: epipolar constraint
  - Photometric: brightness constancy “constraint”

# Optical Flow

Two-view Geometry

# Optical Flow



Video: Juan Adarve (ANU) Optical Flow

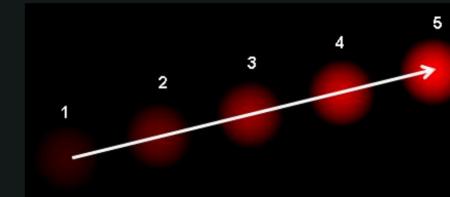
# Optical Flow

- Introduction: structure from continuous infinitesimal motion
- Lucas–Kanade optical flow algorithm (KLT tracker)
- Horn–Schunck optical flow algorithm

# What is Optical Flow?

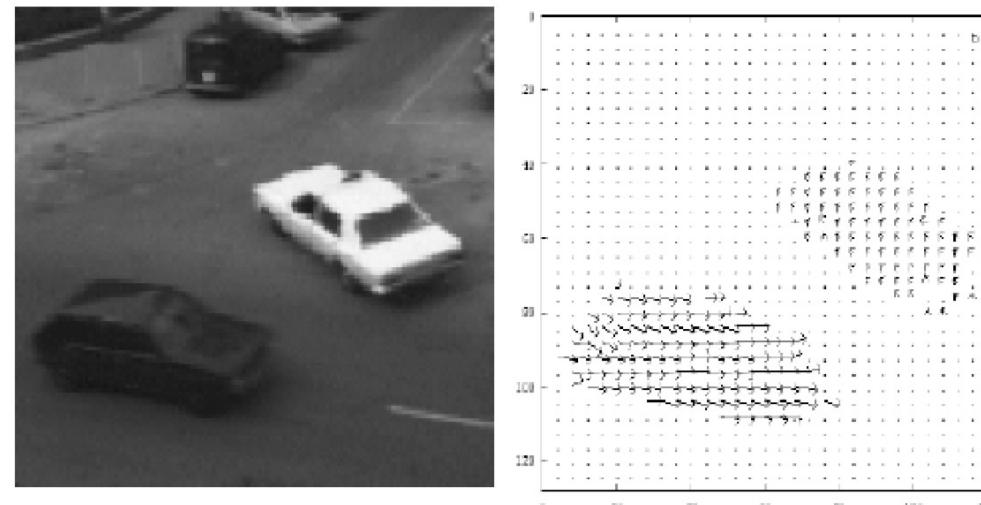
## Optical Flow

Optical flow is the pattern of apparent motion of image objects between two consecutive frames caused by the movement of object or camera. It is 2D vector field where each vector is a displacement vector showing the movement of points from first frame to second. Consider the image below (Image Courtesy: Wikipedia article on Optical Flow).



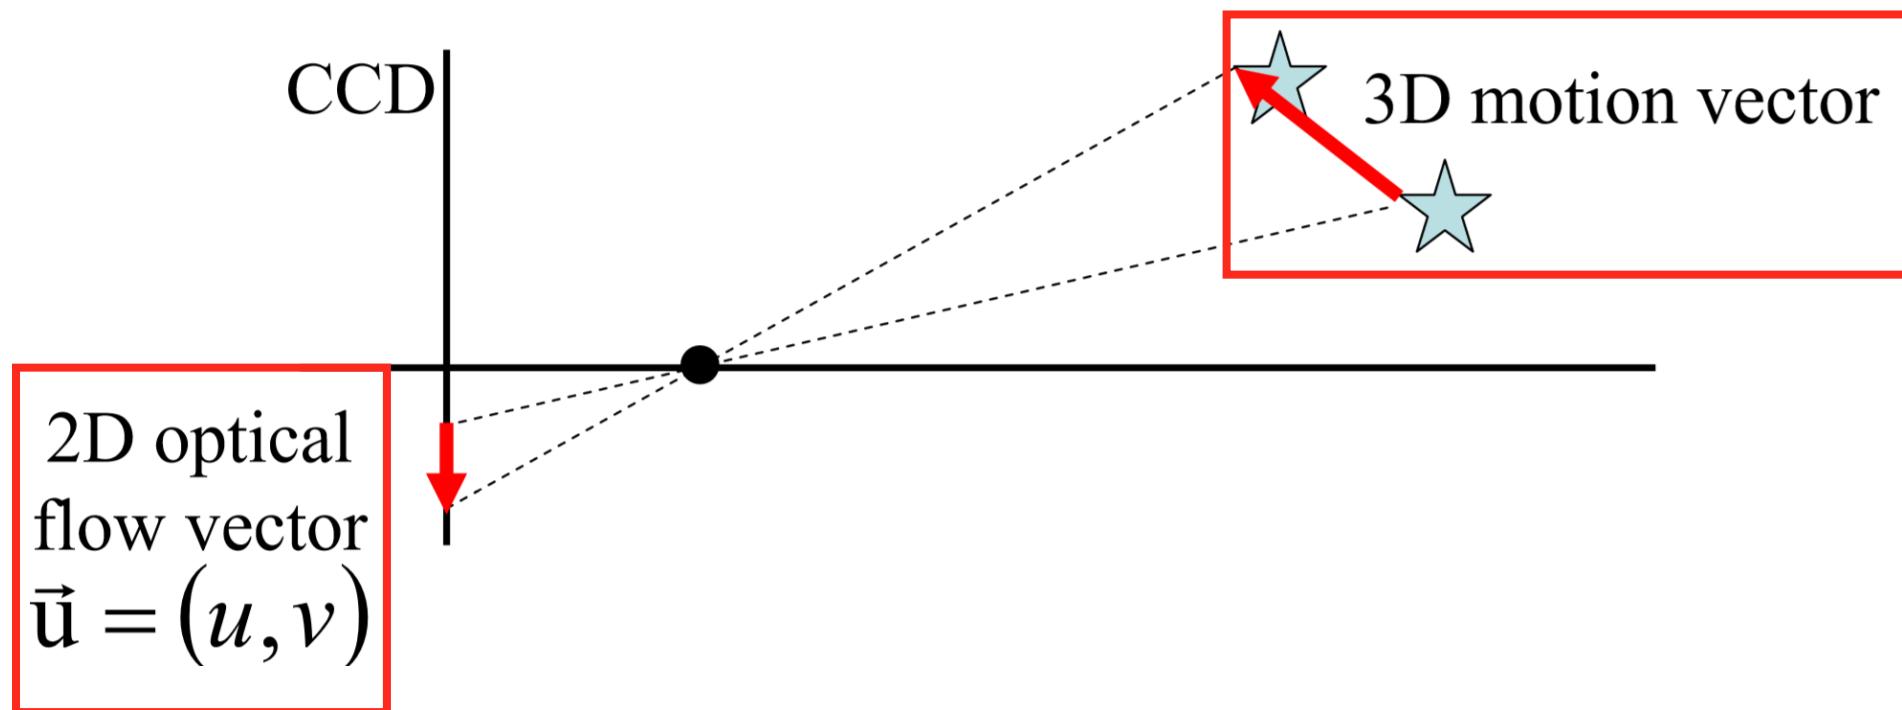
OPTICAL FLOW = apparent motion of  
brightness patterns

Ideally, the optical flow is the projection of the  
three-dimensional velocity vectors on the image

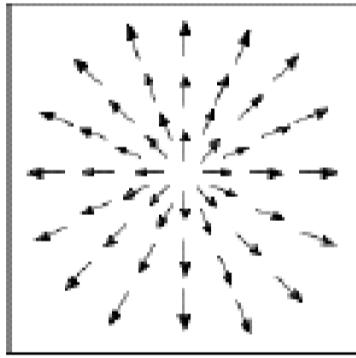


# Motion Field vs Optical Flow

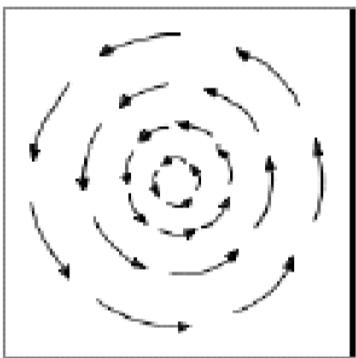
- Motion field: real world 3D motion
- Ideal optical flow: projection of the motion field onto the 2D image



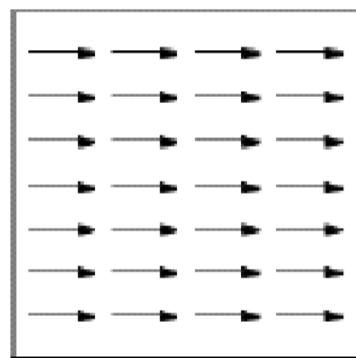
# Some Simple Motion Fields



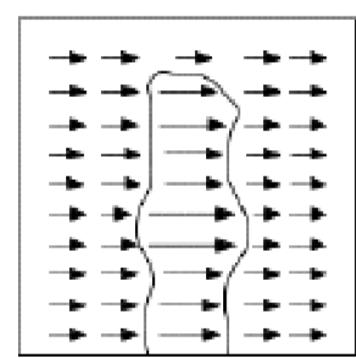
Forward  
motion



Rotation

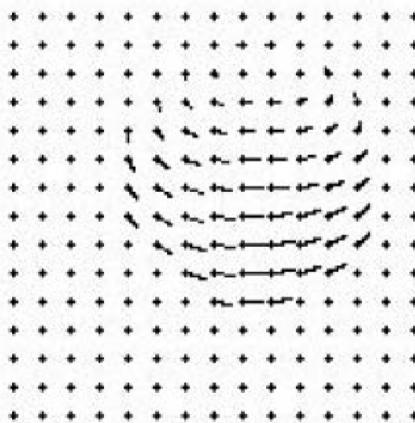
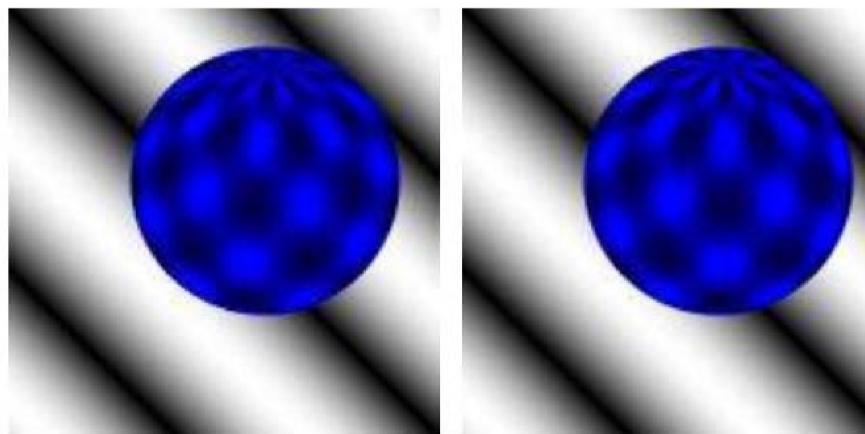


Horizontal  
translation

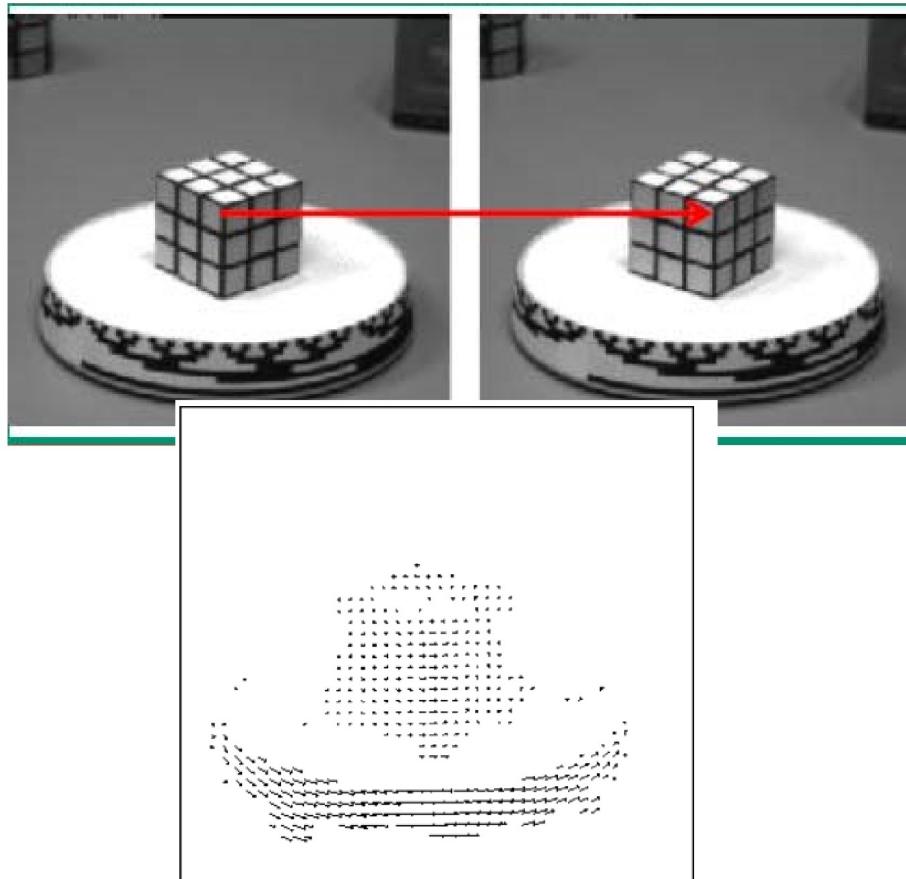


Closer  
objects  
appear to  
move faster!!

# Examples



# Examples



# Examples



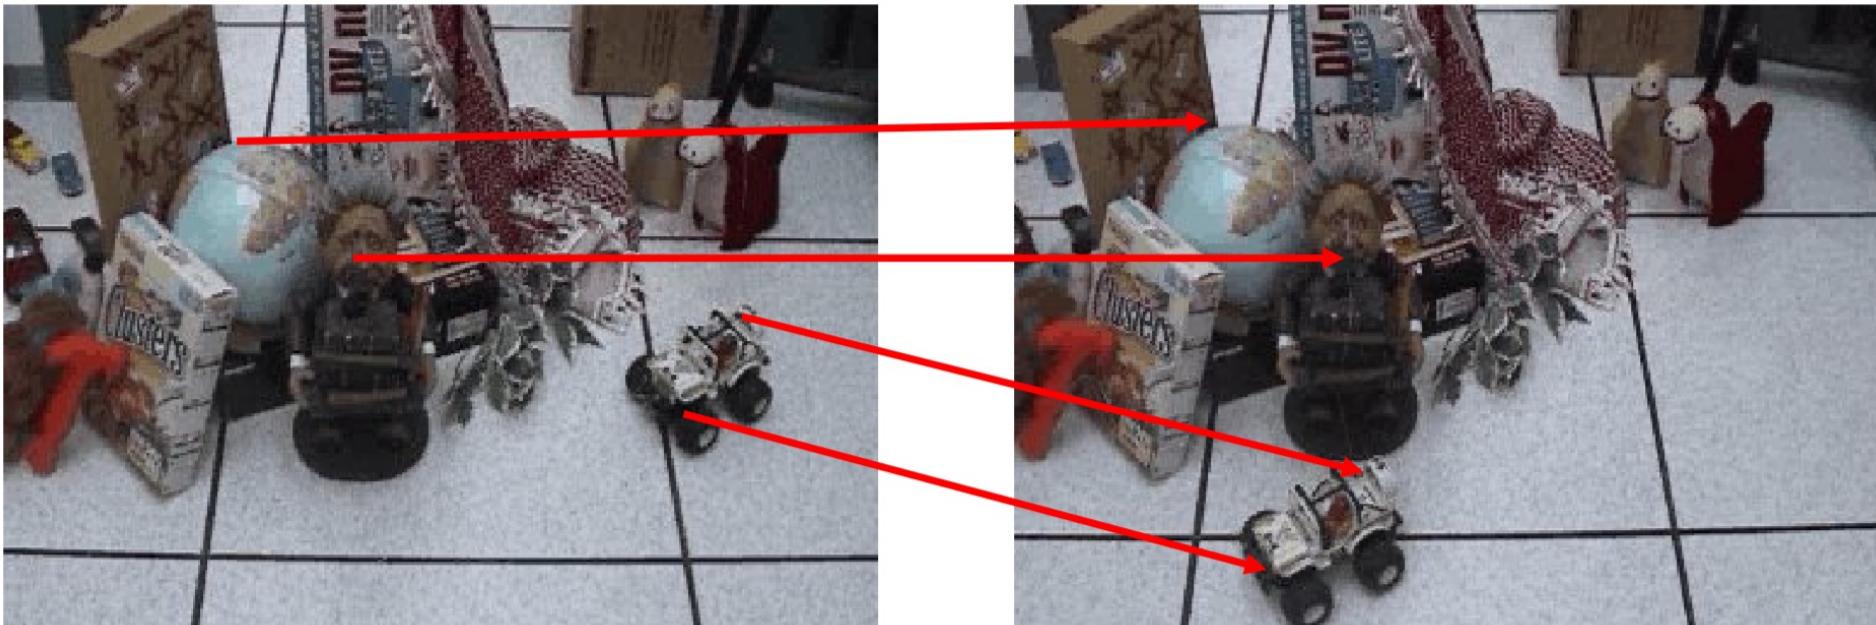
# Optical flow: Applications

- From the optical flow field various parameters can be measured
  - Object shape
  - Object segmentation
  - Camera motion
  - Multiple object motions

# How to Compute Optical Flow?

# Optical Flow Estimation Problem

- Where does each pixel in image 1 go to in image 2?



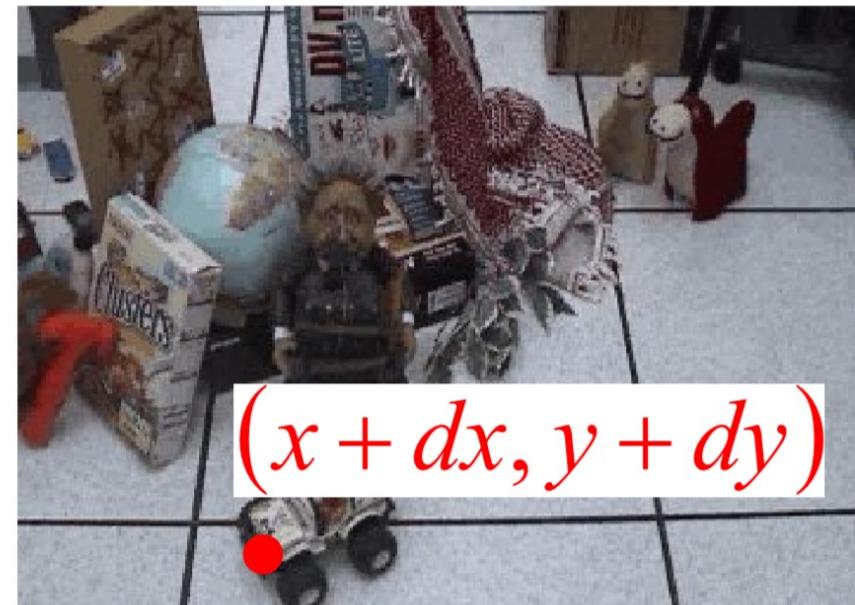
# Estimating Optical Flow

- Assume that the image brightness of the corresponding point is constant over two frames

Time =  $t$



Time =  $t+dt$



$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

# Brightness Constancy Equation

$$I(x, y, t) = I(x + dx, y + dy, t + dt)$$

First order Taylor Expansion

$$= I(x, y, t) + \boxed{\frac{\partial I}{\partial x} dx + \frac{\partial I}{\partial y} dy + \frac{\partial I}{\partial t} dt}$$

Simplify notations:

$$I_x dx + I_y dy + I_t dt = 0$$

Divide by  $dt$  and denote:

$$u = \frac{dx}{dt} \quad v = \frac{dy}{dt}$$

$$I_x u + I_y v = -I_t$$

**Problem I:** One equation, two unknowns

What is  $I_t$ ? The time derivative of the image at  $(x, y)$

shorthand:  $I_x = \frac{\partial I}{\partial x}$

The x-component of the gradient vector.

# Brightness constancy Equation

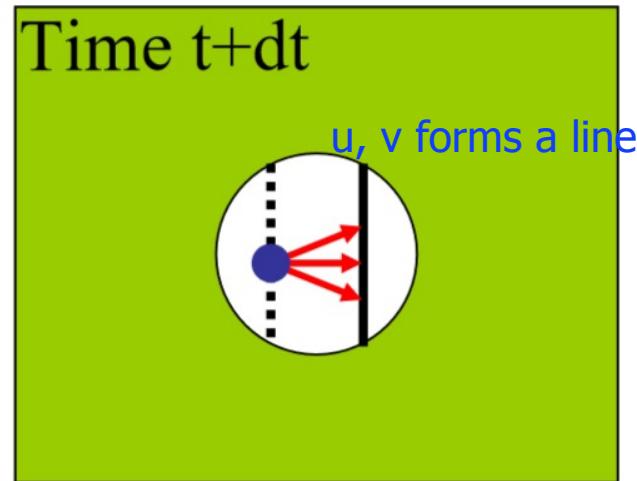
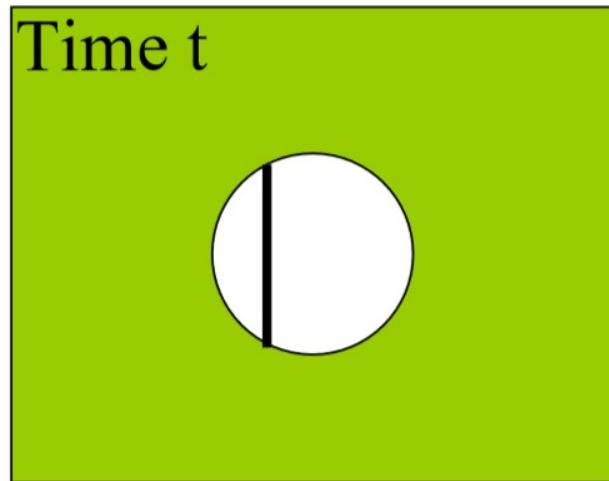
$$I_x u + I_y v = -I_t$$



Note, this assumes a smooth gradient over the area of motion (small scale motion relative to visual texture)

# The Aperture Problem

- For points on a line of fixed intensity we can only recover the normal flow (normal to max gradient)



?

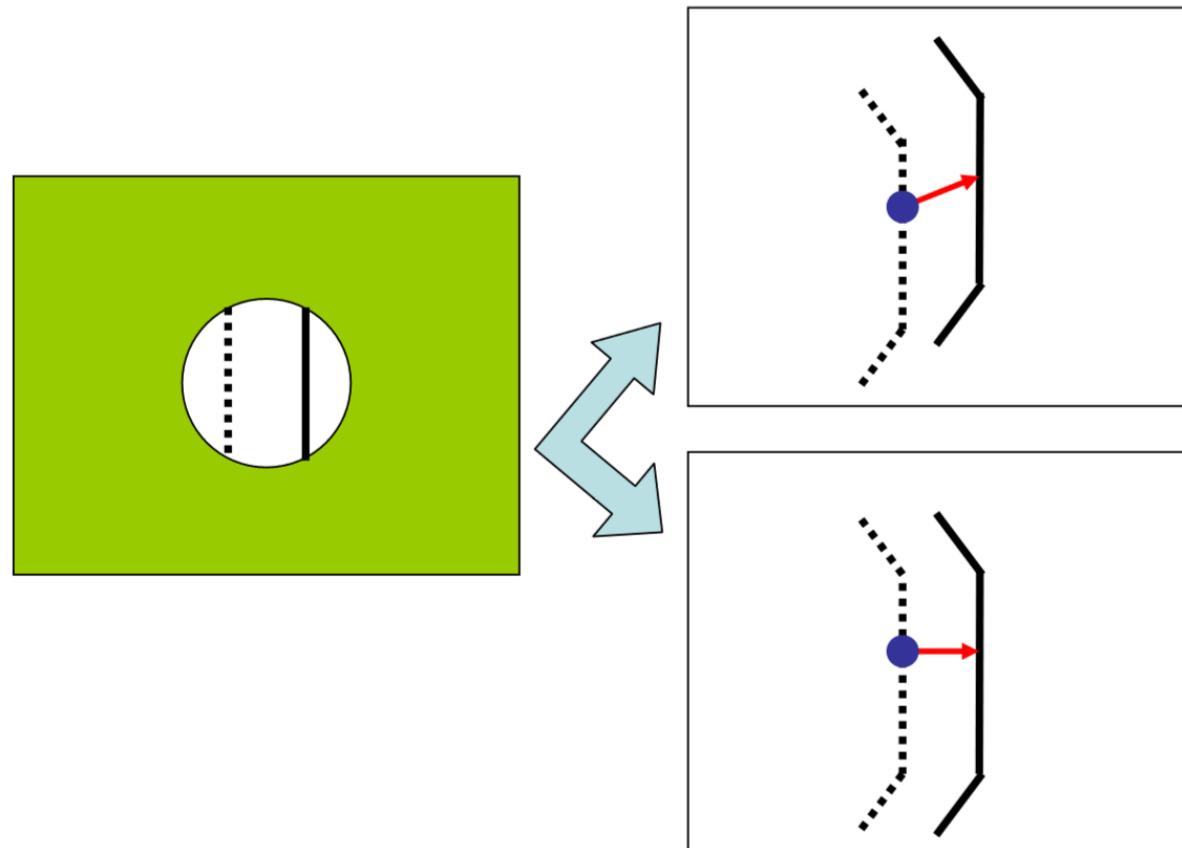


Where did the blue point move to?

We need additional constraints

# The Aperture Problem: a Remedy

- Use local information: sometimes enlarging the aperture can help



# Next Week

- Multiple-view geometry: Structure-from-Motion
- Multiple-view geometry: Radiance fields (fancy triangulation)
- Multiple-view geometry: Learning-based pose-free reconstruction
- Mid-level vision: Optical flow (continued)
- Mid-level vision / single-view geometry: Shape-from-X